

## Outgoing email - SMTP Retrieving email - POP3, IMAP

```
S: 220 mrl.its.yale.edu
C: HELO cyndra.yale.edu
S: 250 Hello cyndra.cs.yale.edu, pleased to meet you
C: MAIL FROM: <spoof@cs.yale.edu>
S: 250 spoof@cs.yale.edu... Sender ok
C: RCPT TO: <cry@yale.edu>
S: 250 cry@yale.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Date: Wed, 23 Jan 2008 11:20:27 -0500 (EST)
C: From: "Y. R. Yang" <cry@cs.yale.edu>
C: To: "Y. R. Yang" <cry@cs.yale.edu>
C: Subject: This is subject
C:
C: This is the message body!
C: Please don't spoof!
C:
C:
S: 250 Message accepted for delivery
C: QUIT
S: 221 mrl.its.yale.edu closing connection
```

### Pros of email design

- separate protocols for different functions
- email retrieval (e.g., POP3, IMAP)
- email transmission (SMTP)
- simple/basic requests to implement basic control; finegrain control through ASCII header and message body
- make the protocol easy to read/debug/extend (analogy with end-to-end layered design?)
- status code in response makes message easy to parse

### Email security

Sender Policy Frame (SPF) - check in channel

DomainKeys Identified Mail (DKIM) - check in sender and receiver

Frequency division multiplexing(FDM)

Time division multiplexing(TDM)

Code division multiplexing(CDM)

Queue theory

Client requests arrival rate lambda/second

Service rate: each call takes on average 1/mu second

Circuit switching

$$\#(\text{transitions } k \rightarrow k+1) = \#(\text{transitions } k+1 \rightarrow k)$$

$$p_k \lambda = p_{k+1} (k+1) \mu$$

$$p_{k+1} = \frac{\lambda}{k+1} p_k = \left(\frac{\lambda}{\mu}\right)^{k+1} p_0$$

$$p_0 = \frac{1}{1 + \frac{\lambda}{\mu} + \frac{1}{2!} \left(\frac{\lambda}{\mu}\right)^2 + \dots + \frac{1}{N!} \left(\frac{\lambda}{\mu}\right)^N}$$

### Packet switching

at equilibrium (time reversibility) in one unit time:

$$\#(\text{transitions } k \rightarrow k+1) = \#(\text{transitions } k+1 \rightarrow k)$$

$$p_k \lambda = p_{k+1} \mu$$

$$p_{k+1} = \frac{\lambda}{\mu} p_k = \left(\frac{\lambda}{\mu}\right)^{k+1} p_0 = \rho^{k+1} p_0$$

$$\rho = \frac{\lambda}{\mu}$$

$$p_0 = 1 - \rho$$

$$\rho = \frac{\lambda}{\mu}$$

$$S = \frac{1}{\mu}$$

$$\text{average queueing delay: } w = S \frac{\rho}{1 - \rho}$$

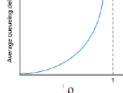
$$\text{queueing + trans} = S \frac{\rho}{1 - \rho} + S = S \frac{1}{1 - \rho}$$

Assume:  
R = link bandwidth (bps)  
L = packet length (bits)  
S = L / R  
a = average packet arrival rate (pkts/sec)

$$\text{utilization: } \rho = \frac{a}{1/S} = aS$$

$$w = S \frac{\rho}{1 - \rho}$$

- $\rho \sim 0$ : average queueing delay small
- $\rho \rightarrow 1$ : delay becomes large
- $\rho > 1$ : more "work" arriving than can be serviced, average delay infinite!

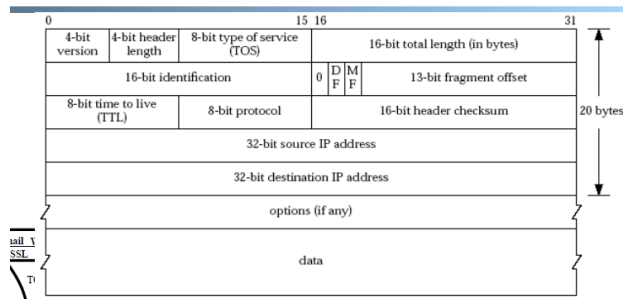


### pros of packet switching:

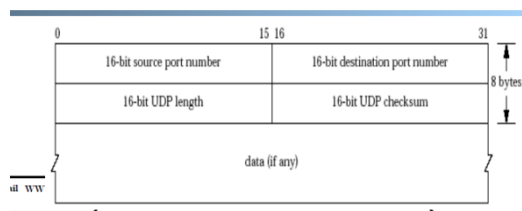
Statistical multiplexing, more efficient bandwidth usage.

Cons: Congestion, packet header overhead, per packet processing overhead.

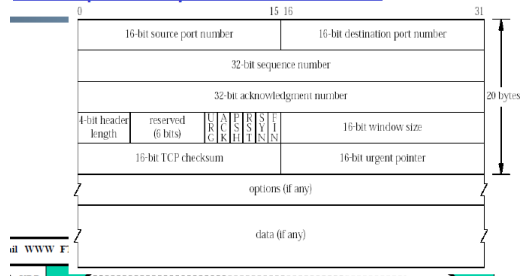
## Network Layer: IPv4 Header



## Transport Layer: UDP Header



## Transport Layer: TCP Header



Key questions to ask about a C-S application

Extensible? Scalable? Handle failures(robust)? Security?

DNS: Port 53, over UDP

**DNS protocol** : typically over UDP (can use TCP);  
*query* and *reply* messages, both with the **same message format**

Identification	Flags	
Number of questions	Number of answer RRs	12 bytes
Number of authority RRs	Number of additional RRs	
Questions (variable number of questions)		Name, type fields for a query
Answers (variable number of resource records)		RRs in response to query
Authority (variable number of resource records)		Records for authoritative servers
Additional information (variable number of resource records)		Additional "helpful" info that may be used

### Pros of DNS design

Hierarchical delegation avoids central control, improving manageability and scalability

- Redundant servers improve robustness
- Caching reduces workload and improves robustness

Cons:

Domain names may not be the best way to name other resources, e.g. files

□ Simple query model makes it hard to implement advanced query

□ Relatively static resource types make it hard to introduce new services or handle mobility

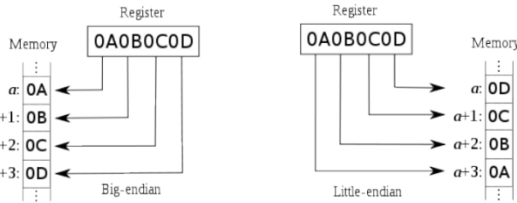
□ Although theoretically you can update the values of the records, it is rarely enabled

□ Early binding (separation of DNS query from application query) does not work well in mobile, dynamic environments

TCP\_UDP

Network endian: big-endian

Intel x86 - little-endian

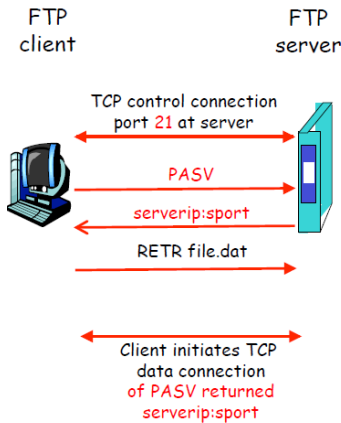


## FTP

Use telnet for the control channel

use nc (NetCat) to receive/send data

FTP PASV: Server Specifies Data Port, Client Initiates Connection



## HTTP

Http is stateless, comparing to FTP

Speed up http protocol

Reduce the number of objects fetched [Browser cache] - if-modified-since

Response 304, not modified

Reduce data volume [Compression of data]

Reduce the latency to the server to fetch the content [Proxy cache]

Forward proxy server - behave like client

Reverse proxy server - behave like server

Keep consistency

pull

Web caches periodically pull the web server to see if a document is modified

push

whenever a server gives a copy of a web page to a web cache, they sign a lease with an expiration time; if the web page is modified before the lease, the server notifies the cache

Increase concurrency [Multiple TCP connections]

Remove the extra RTTs to fetch an object [Persistent

HTTP, aka HTTP/1.1]

Asynchronous fetch (multiple streams) using a single TCP [HTTP/2]

Server push [HTTP/2]

Header compression [HTTP/2]

## Background: Little's Law (1961)



- For any system with no or (low) loss.
- Assume
  - mean arrival rate  $\lambda$ , mean time  $R$  at system, and mean number  $Q$  of requests at system
- Then relationship between  $Q$ ,  $\lambda$ , and  $R$ :

$$Q = \lambda R$$

## Asyn server

In addition to management threads, a system may still need multiple threads for performance (why?)

-FSM code can never block, but page faults, file io, garbage collection may still force blocking

-CPU may become the bottleneck and there maybe multiple cores supporting multiple threads (typically 2 n threads)

## Operational Quantities

- $T$ : observation interval
- $B_i$ : busy time of device  $i$
- $i = 0$  denotes system
- $A_i$ : # arrivals to device  $i$
- $C_i$ : # completions at device  $i$

$$\text{arrival rate } \lambda_i = \frac{A_i}{T}$$

$$\text{Throughput } X_i = \frac{C_i}{T}$$

$$\text{Utilization } U_i = \frac{B_i}{T}$$

$$\text{Mean service time } S_i = \frac{B_i}{C_i}$$

$$\text{Utilization } U_i = \frac{B_i}{T} = \frac{C_i}{T} * \frac{B_i}{C_i} = X_i S_i$$

- Assume flow balanced (arrival=throughput), Little's Law:

$$Q = \lambda R = X S$$

- Assume PASTA (Poisson arrival--memory-less arrival--sees time average), a new request sees  $Q$  ahead of it, and FIFO

$$R = S + Q S = S + X R S$$

- According to utilization law,  $U = X S$

$$R = S + U R \implies R = \frac{S}{1-U}$$

## Why Multiple Servers?

Scalability

Scaling beyond single server capability

There is a fundamental limit on what a single server can process (CPU/bw/disk throughput)

store (disk/memory)

Scaling beyond geographical location capability

There is a limit on the speed of light

Network detour and delay further increase the delay

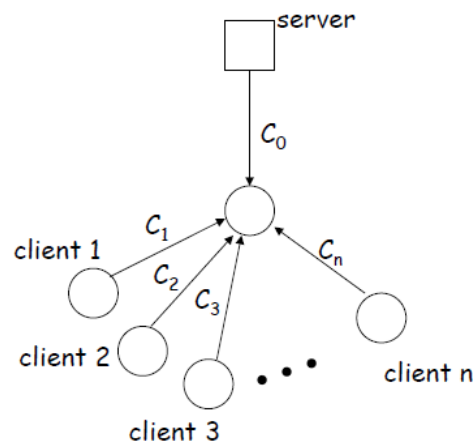
LB/NAT

Pro: Only one public IP address is needed for the load balancer; real servers can use private IP Addresses; Real servers need no change and are not aware of load balancing

cons: The load balancer must be on the critical path and hence may become the bottleneck due to load to rewrite request and response packets. Typically, rewriting responses has more load because there are more response packets

So we configure real servers and load balancer to use same ip address.

Configure address resolution protocol



$$R = \min\{C_0, (C_0 + \sum C_i)/n\}$$

Commands

Telnet

Traceroute

Dig <type> <domain> Type: MX, A, TXT, NS, CNAME