
Network Applications: DNS Details; UDP Network App Programming

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

2/3/2016

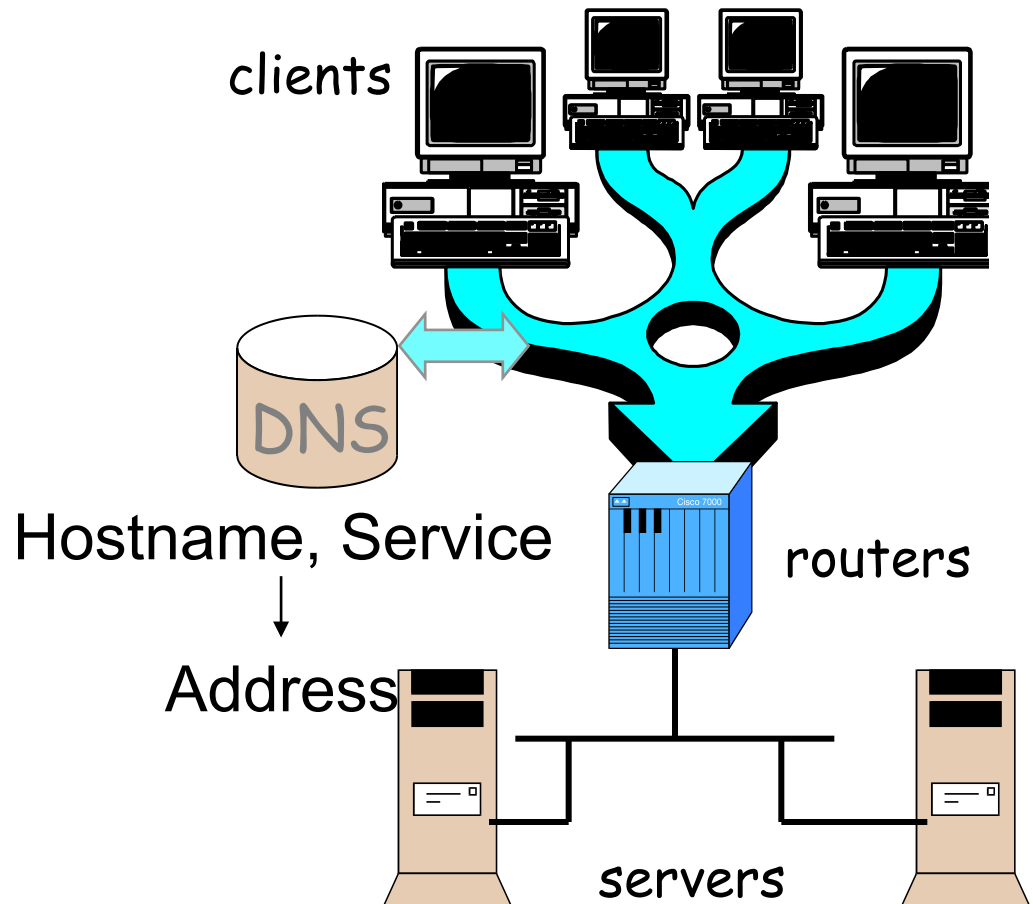
Outline

- Admin and recap
- ❑ Network app programming

Recap: Domain Name System (DNS)

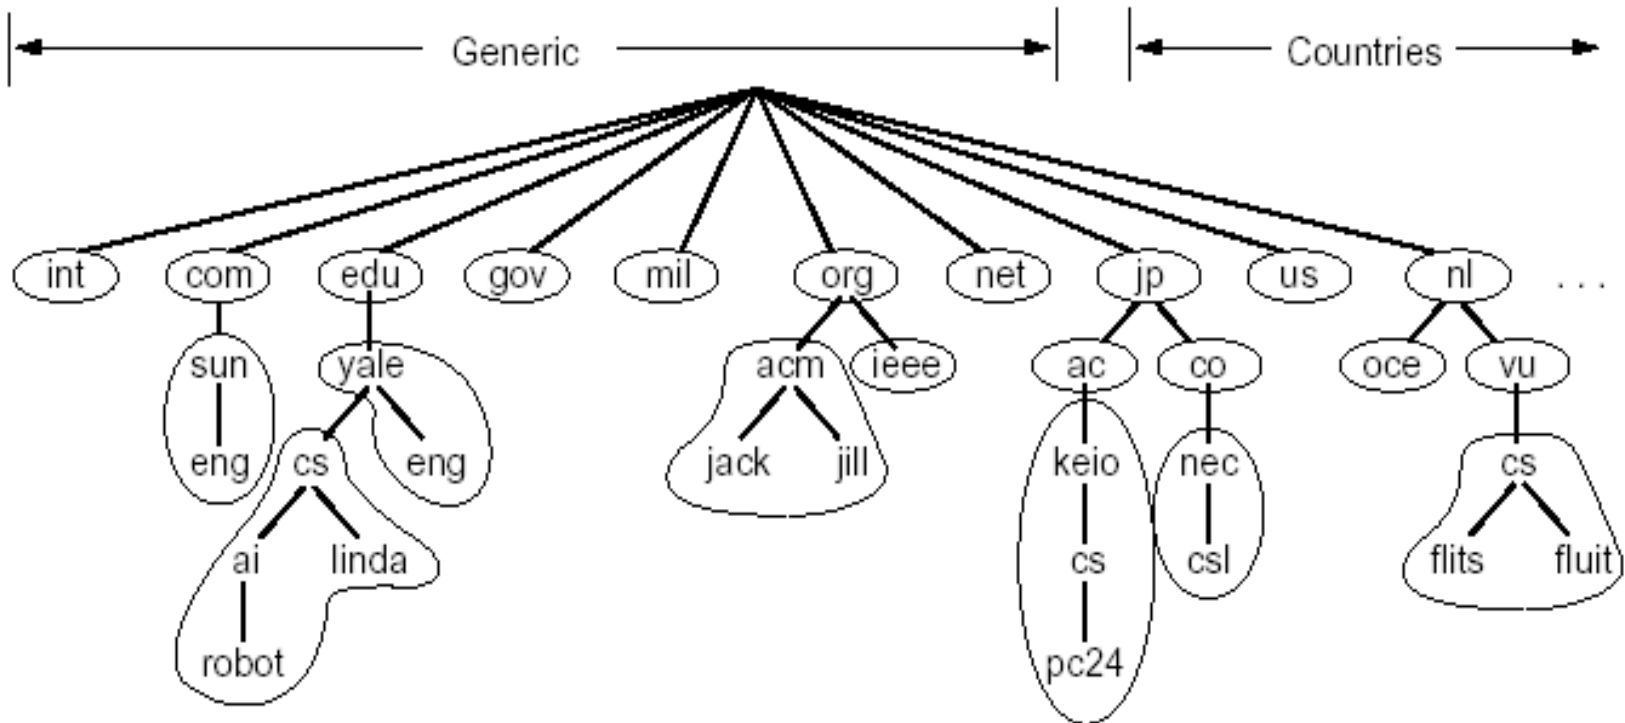
□ Function

- map between (domain name, service) to value, e.g.,
 - (www.cs.yale.edu, Addr)
→ 128.36.229.30
 - (cs.yale.edu, Email)
→ netra.cs.yale.edu



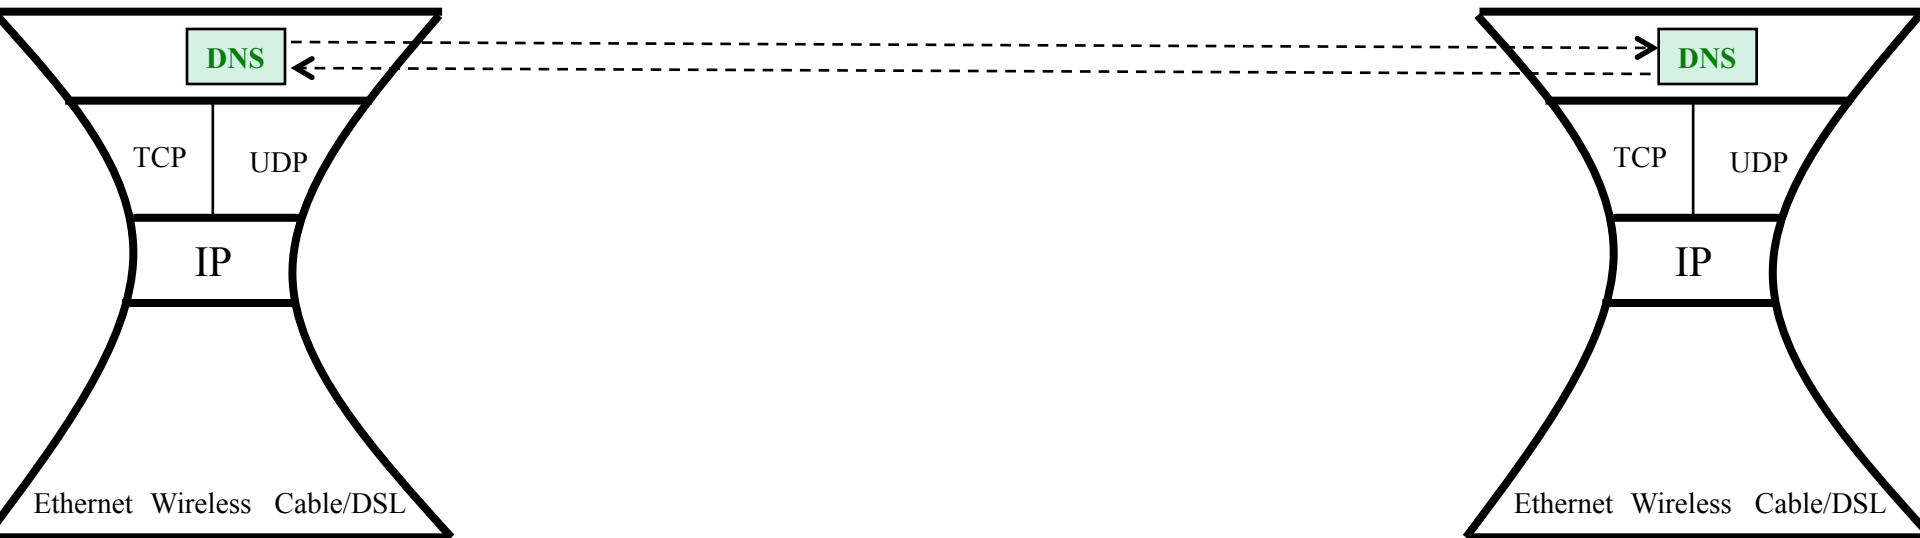
Recap: DNS

- Key design features of DNS
 - Hierarchical domain name space allowing delegation
 - Recursive or iterative queries



DNS Message Format?

Basic encoding decisions: UDP/TCP,
how to encode domain name, how to
encode answers...



Observing DNS Messages

❑ Issue DNS query using the command dig:

- force iterated query to see the trace:

```
%dig +trace www.cnn.com
```

- see the manual for more details

❑ Capture the messages

- DNS server is at port 53

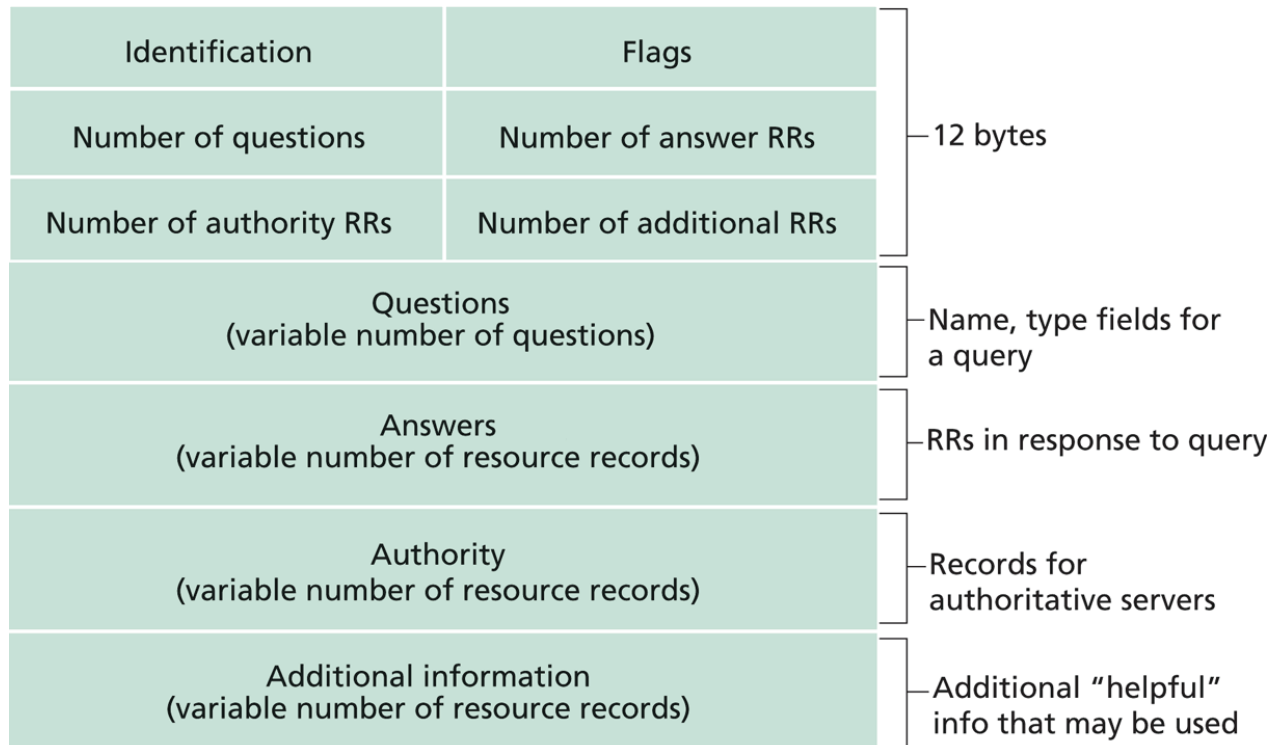
- Display and clear DNS cache

- <https://support.apple.com/en-us/HT202516> (e.g., MAC)

- Try to load the dns-capture file from class Schedule page, if you do not want live capture

DNS Protocol, Messages

DNS protocol : typically over UDP (can use TCP);
query and *reply* messages, both with the *same*
message format



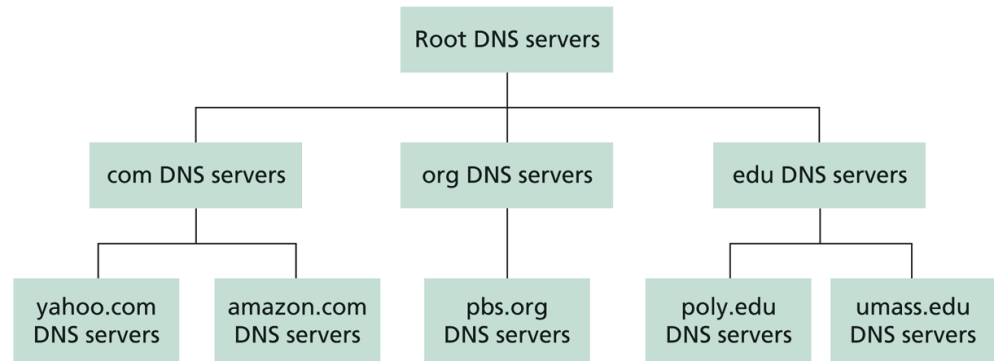
DNS Details

- ❑ Header (Sec. 4.1.1 of <https://www.ietf.org/rfc/rfc1035.txt>)
- ❑ Encoding of questions (Sec. 4.1.2):
 - [Label-length label-chars]
- ❑ Encoding of answers (Sec. 4.1.3)
 - Pointer format (<http://www.iana.org/assignments/dns-parameters/dns-parameters.xhtml>)
- ❑ See example DNS packets

Evaluation of DNS

Key questions to ask about a C-S application

- **extensible?**
- **scalable?**
- **robust?**
- **security?**



What DNS did Right?

- ❑ Hierarchical delegation avoids central control, improving manageability and scalability
- ❑ Redundant servers improve robustness
 - see <http://www.internetnews.com/dev-news/article.php/1486981> for DDoS attack on root servers in Oct. 2002 (9 of the 13 root servers were crippled, but only slowed the network)
- ❑ Caching reduces workload and improves robustness

Problems of DNS

- ❑ Domain names may not be the best way to name other resources, e.g. files
- ❑ Simple query model makes it hard to implement advanced query
- ❑ Relatively static resource types make it hard to introduce new services or handle mobility
- ❑ Although theoretically you can update the values of the records, it is rarely enabled
- ❑ Early binding (separation of DNS query from application query) does not work well in mobile, dynamic environments
 - e.g., load balancing, locate the nearest printer

Outline

- ❑ Recap
- ❑ Network app programming

Socket Programming

Socket API

- ❑ introduced in BSD4.1 UNIX, 1981
- ❑ Two types of sockets
 - Connectionless (UDP)
 - connection-oriented (TCP)

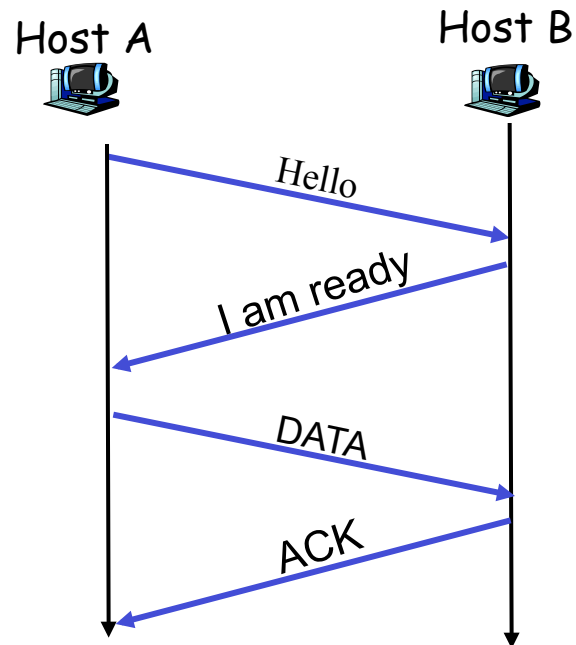
socket

an interface (a “door”) into which one application process can **both send and receive** messages to/from another (remote or local) application process

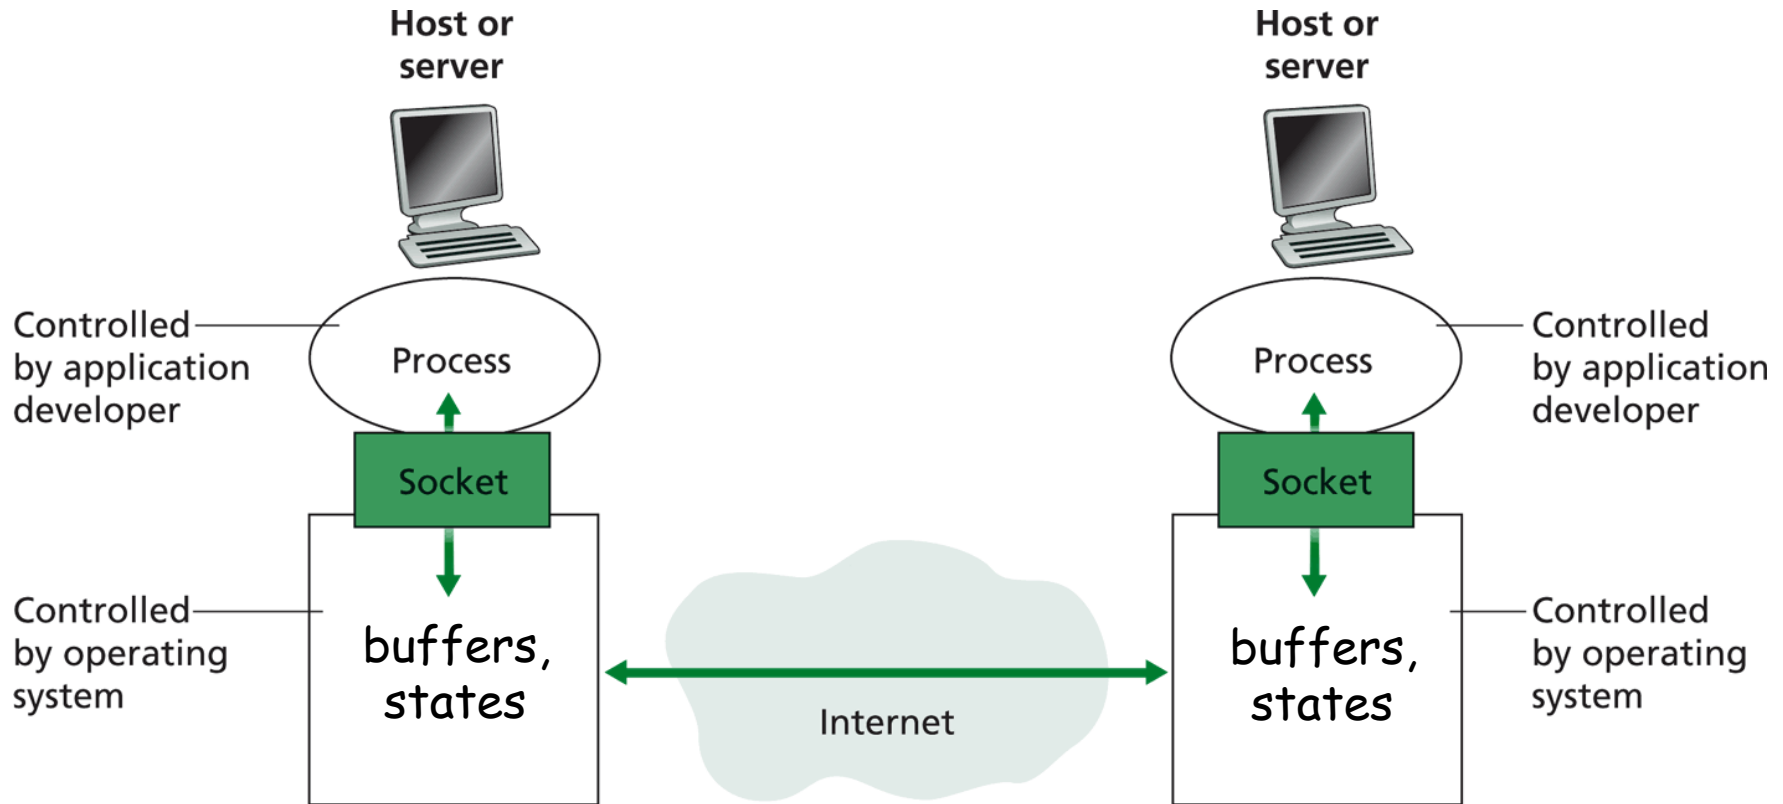
Services Provided by Transport

- User data protocol (UDP)
 - multiplexing/demultiplexing

- Transmission control protocol (TCP)
 - multiplexing/demultiplexing
 - reliable data transfer
 - rate control: flow control and congestion control



Big Picture: Socket



Outline

- ❑ Recap
- ❑ Basic network application programming
 - Overview
 - UDP (Datagram Socket)

DatagramSocket (Java) (Basic)

- ❑ **DatagramSocket()**
constructs a datagram socket and binds it to any available port on the local host
- ❑ **DatagramSocket(int lport)**
constructs a datagram socket and binds it to the specified port on the local host machine.
- ❑ **DatagramPacket(byte[] buf, int length)**
constructs a DatagramPacket for receiving packets of length length.
- ❑ **DatagramPacket(byte[] buf, int length, InetAddress address, int port)**
constructs a datagram packet for sending packets of length length to the specified port number on the specified host.
- ❑ **receive(DatagramPacket p)**
receives a datagram packet from this socket.
- ❑ **send(DatagramPacket p)**
sends a datagram packet from this socket.
- ❑ **close()**
closes this datagram socket.

Connectionless UDP: Big Picture (Java version)

Server (running on `serv`)

create socket,
port=`x`, for
incoming request:
`serverSocket =`
`DatagramSocket(x)`

read request from
`serverSocket`

generate reply, create
datagram using client
host address, port number

write reply to
`serverSocket`

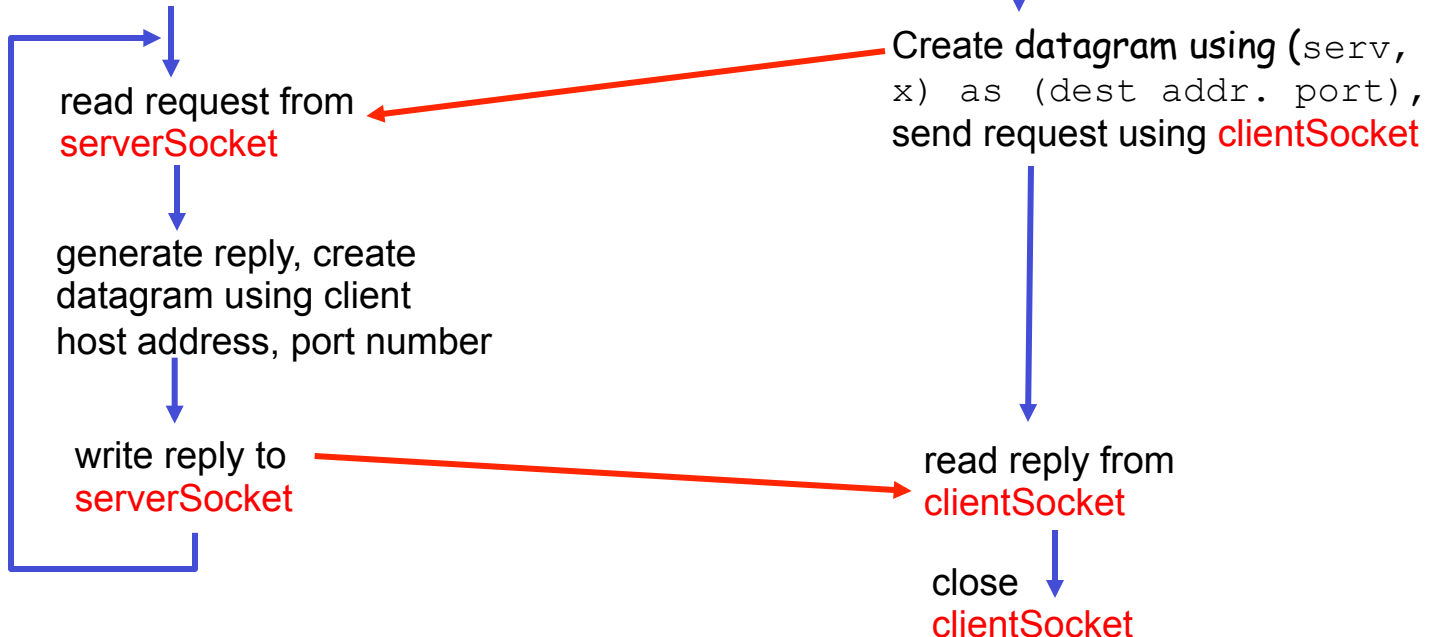
Client

create socket,
`clientSocket =`
`DatagramSocket()`

Create datagram using (`serv`,
`x`) as (dest addr. port),
send request using `clientSocket`

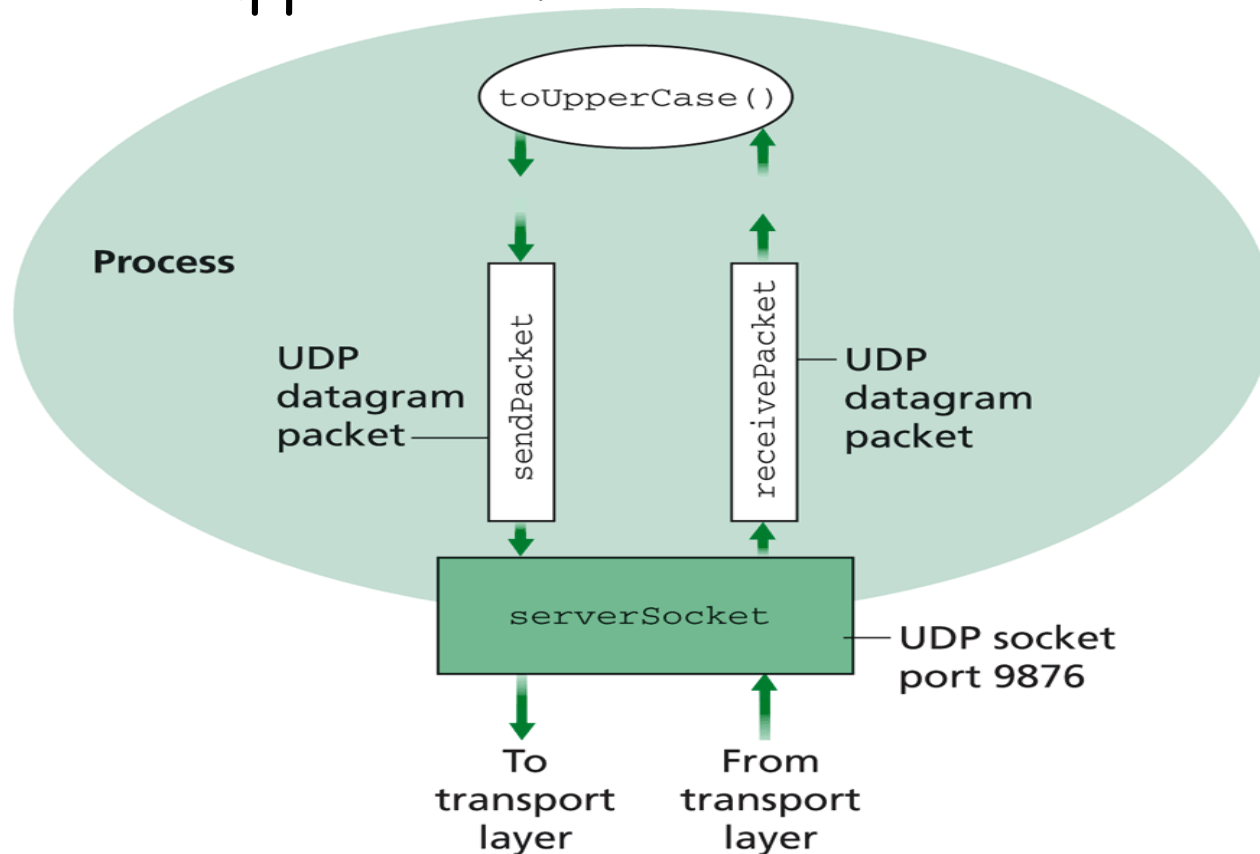
read reply from
`clientSocket`

close
`clientSocket`



Example: UDPServer.java

- A simple UDP server which changes any received sentence to upper case.



Java Server (UDP): Create Socket

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

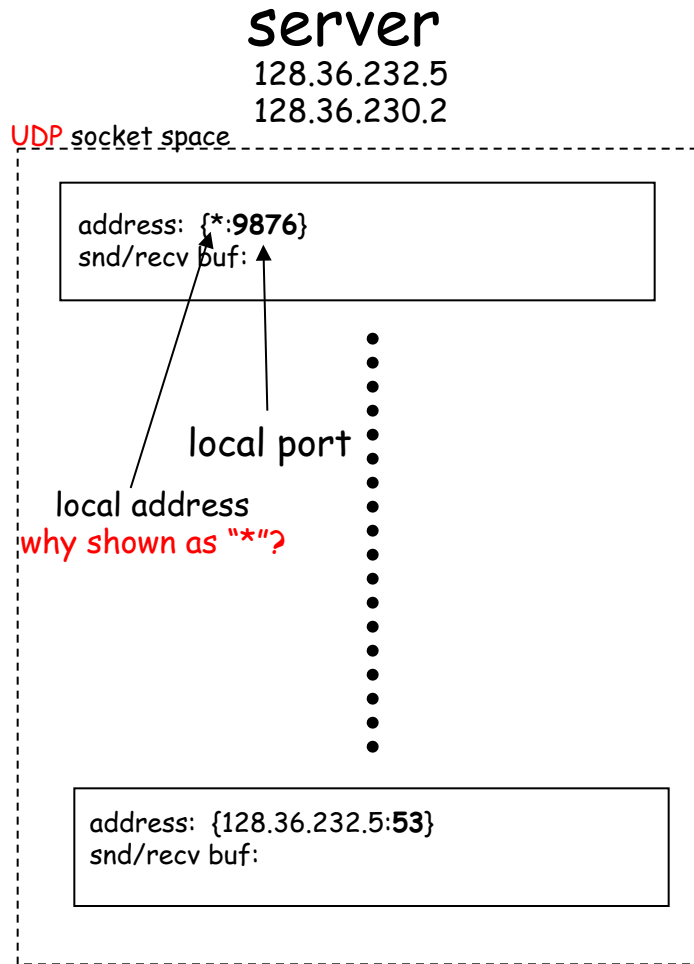
Create
datagram socket
bind at port 9876



```
DatagramSocket serverSocket = new DatagramSocket(9876);
```

Check socket state:
%netstat -a -p udp -n

System State after the Call



"*" indicates that the socket binds to **all** IP addresses of the machine:

% `ifconfig -a`

Binding to Specific IP Addresses

server

Public address: 128.36.59.2

Local address: 127.0.0.1

UDP socket space

address: {127.0.0.1:9876}
snd/rcv buf:

address: {128.36.59.2:9876}
snd/rcv buf:

address: {*:6789}
snd/rcv buf:

•
•
•
•
•

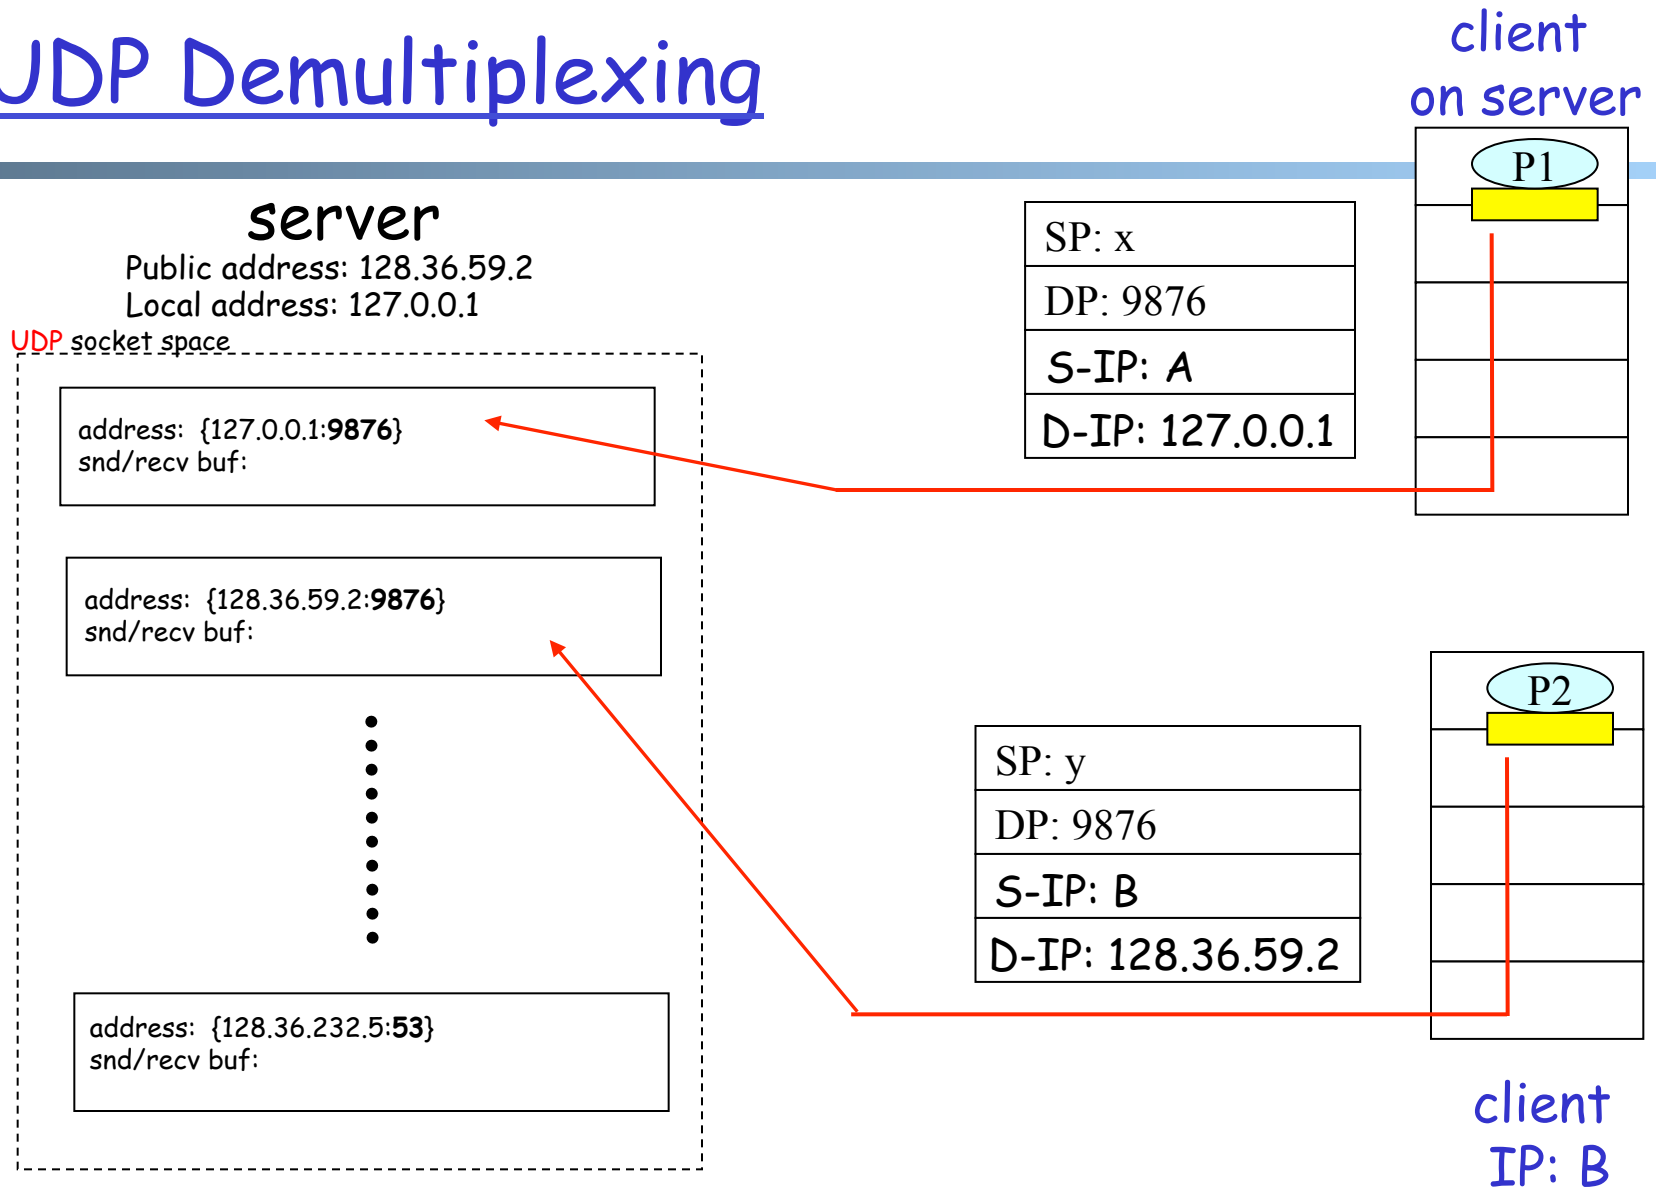
address: {128.36.232.5:53}
snd/rcv buf:

```
InetAddress sIP1 =  
    InetAddress.getByName("localhost");  
DatagramSocket ssock1 = new  
    DatagramSocket(9876, sIP1);
```

```
InetAddress sIP2 =  
    InetAddress.getByName("128.36.59.2");  
DatagramSocket ssock2 = new  
    DatagramSocket(9876, sIP2);
```

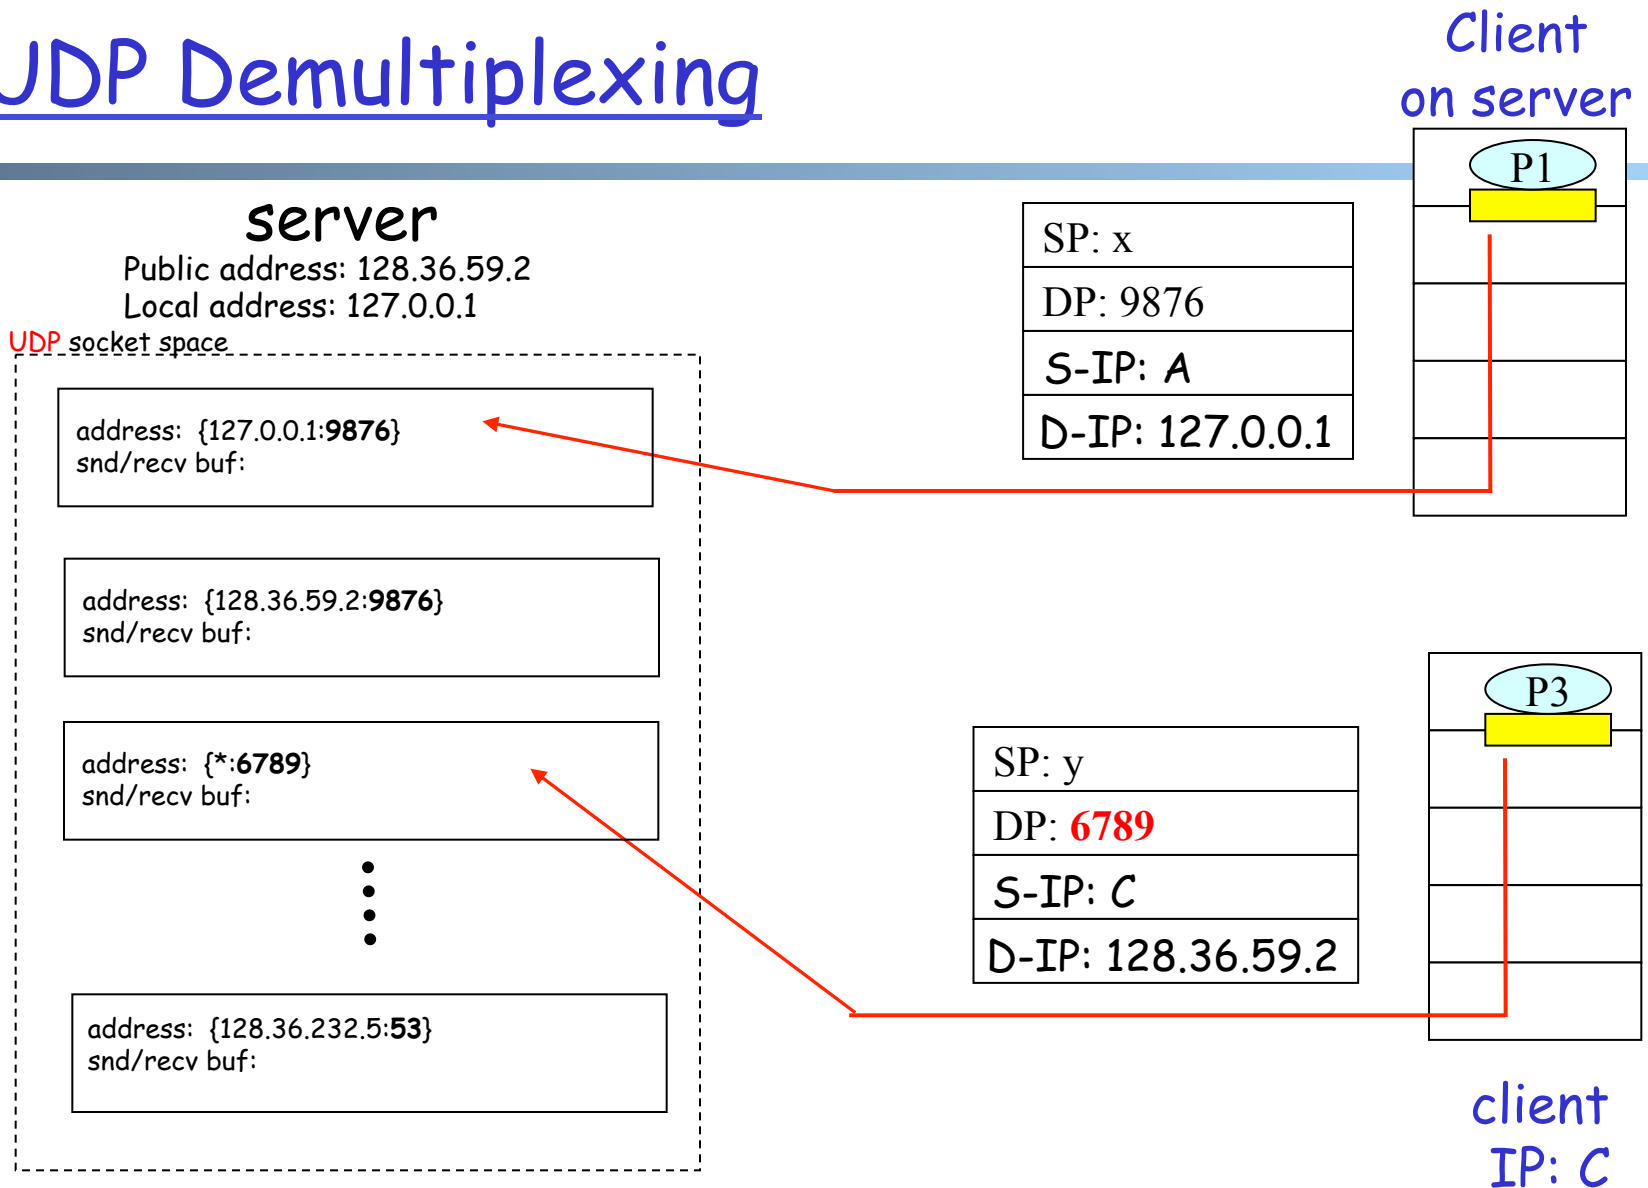
```
DatagramSocket serverSocket = new  
    DatagramSocket(6789);
```

UDP Demultiplexing



UDP demultiplexing is based on matching (dst address, dst port)

UDP Demultiplexing



UDP demultiplexing is based on matching (dst address, dst port)

Per Socket State

- ❑ Each Datagram socket has a set of states:
 - local address
 - send buffer size
 - receive buffer size
 - timeout
 - traffic class

See <http://download.java.net/jdk7/archive/b123/docs/api/java/net/DatagramSocket.html>

Example: socket state after clients sent msgs to the server

Java Server (UDP): Receiving

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception  
    {
```

```
        DatagramSocket serverSocket = new DatagramSocket(9876);
```

```
        byte[] receiveData = new byte[1024];  
        byte[] sendData = null;
```

```
        while(true)  
        {
```

Create space for
received datagram



```
            DatagramPacket receivePacket =  
                new DatagramPacket(receiveData, receiveData.length);
```

Receive
datagram



```
            serverSocket.receive(receivePacket);
```

DatagramPacket

❑ Receiving

- **DatagramPacket(byte[] buf, int length)**
constructs a DatagramPacket for receiving packets of length length.
- **DatagramPacket(byte[] buf, int offset, int length)**
constructs a DatagramPacket for receiving packets starting at offset, length length.

❑ Sending

- **DatagramPacket(byte[] buf, int length, InetAddress address, int port)**
constructs a datagram packet for sending packets of length length to the specified port number on the specified host.
- **DatagramPacket(byte[] buf, int offset, int length, InetAddress address, int port)**

Java Server (UDP): Processing

```
import java.io.*;  
import java.net.*;
```

```
class UDPServer {  
    public static void main(String args[]) throws Exception {
```

```
        ...
```

```
        // process data
```

```
        String sentence = new String(receivePacket.getData(),  
                                     0, receivePacket.getLength());
```

```
        String capitalizedSentence = sentence.toUpperCase();  
        sendData = capitalizedSentence.getBytes();
```

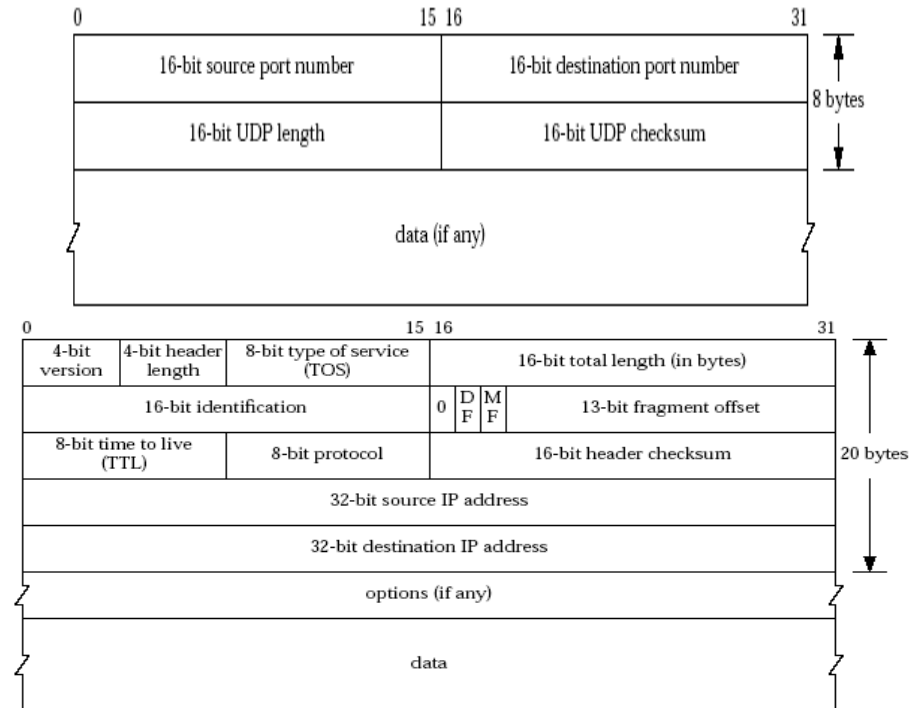
getData() returns a pointer to an underlying buffer array;
for efficiency, don't assume receive() will reset the rest of the array

getLength() returns how much data is valid.

Java Server (UDP): Response

❑ Java DatagramPacket:

- getAddress() / getPort() returns the **source** address/port



Java server (UDP): Reply

Get IP addr
port #, of
sender

```
InetAddress IPAddr = receivePacket.getAddress();  
int port = receivePacket.getPort();
```

Create datagram
to send to client

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length,  
        IPAddr, port);
```

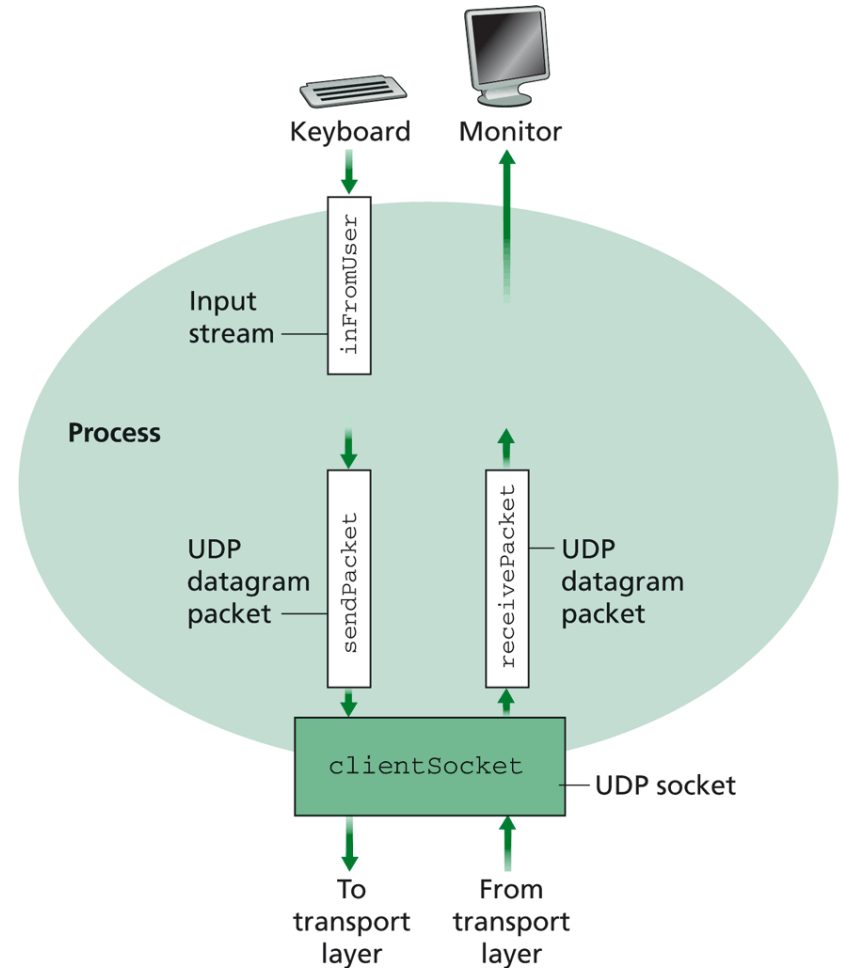
Write out
datagram
to socket

```
serverSocket.send(sendPacket);  
}
```

End of while loop,
loop back and wait for
another datagram

Example: UDPClient.java

- A simple UDP client which reads input from keyboard, sends the input to server, and reads the reply back from the server.



Example: Java client (UDP)

```
import java.io.*;  
import java.net.*;
```

```
class UDPClient {  
    public static void main(String args[]) throws Exception  
    {
```

Create
input stream

```
        → BufferedReader inFromUser =  
            new BufferedReader(new InputStreamReader(System.in));  
        String sentence = inFromUser.readLine();  
        byte[] sendData = sentence.getBytes();
```

Create
client socket

```
        → DatagramSocket clientSocket = new DatagramSocket();
```

Translate
hostname to IP
address using DNS

```
        → InetAddress sIPAddress = InetAddress.getByName("servname");
```


Example: Java client (UDP), cont.

Create datagram
with data-to-send,
length, IP addr, port

Send datagram
to server

Read datagram
from server

```
DatagramPacket sendPacket =  
    new DatagramPacket(sendData, sendData.length, sIPAddress, 9876);  
  
clientSocket.send(sendPacket);  
  
byte[] receiveData = new byte[1024];  
DatagramPacket receivePacket =  
    new DatagramPacket(receiveData, receiveData.length);  
  
clientSocket.receive(receivePacket);  
  
String modifiedSentence =  
    new String(receivePacket.getData());  
  
System.out.println("FROM SERVER:" + modifiedSentence);  
clientSocket.close();  
}  
}
```

Demo

%mac: java UDPServer

% netstat to see buffer

%cicada: java UDPClient <server>

% wireshark to capture traffic

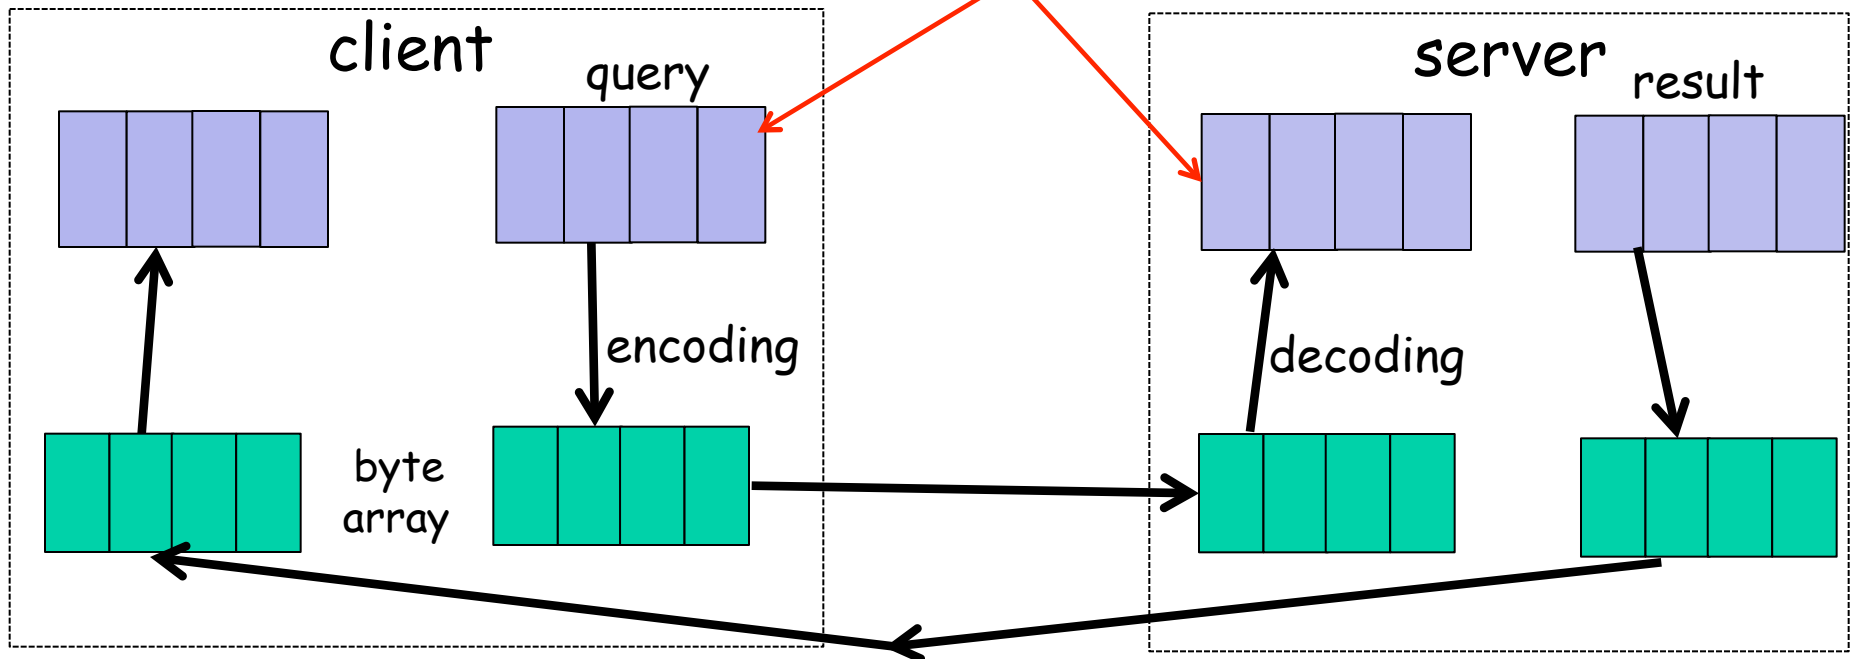
Discussion on Example Code

- ❑ A simple upper-case UDP echo service is among the simplest network service.
- ❑ Are there any problems with the program?

Data Encoding/Decoding

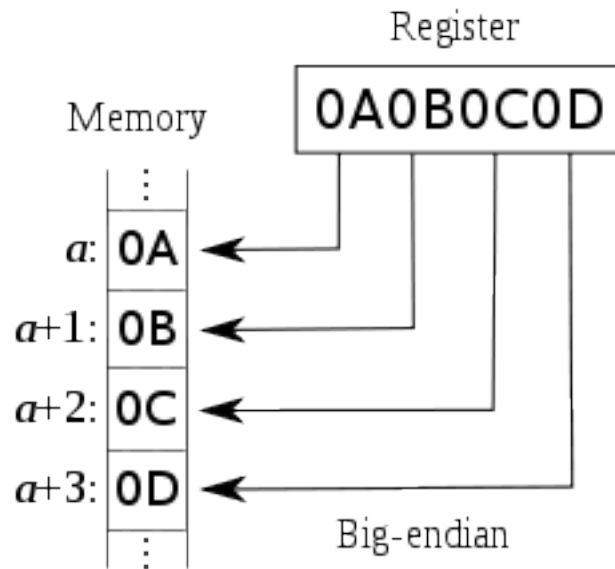
- Pay attention to encoding/decoding of data:
transport layer handles only a sequence of bytes

if not careful, query sent !=
query received (how?)

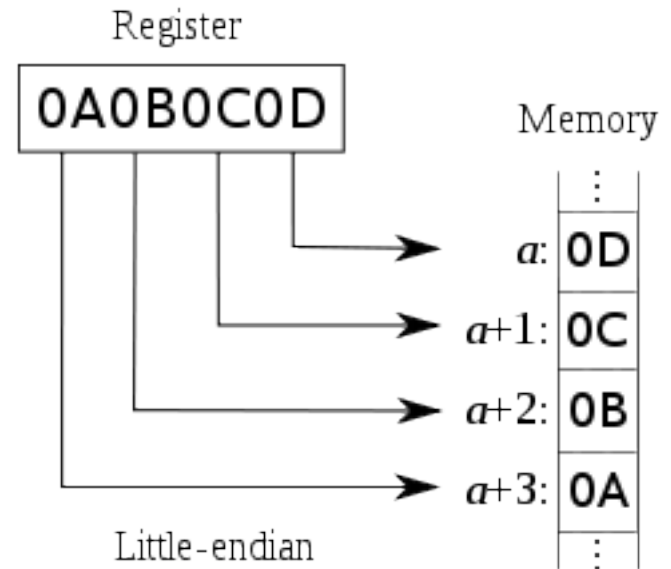


Example: Endianness of Numbers

□ `int var = 0x0A0B0C0D`



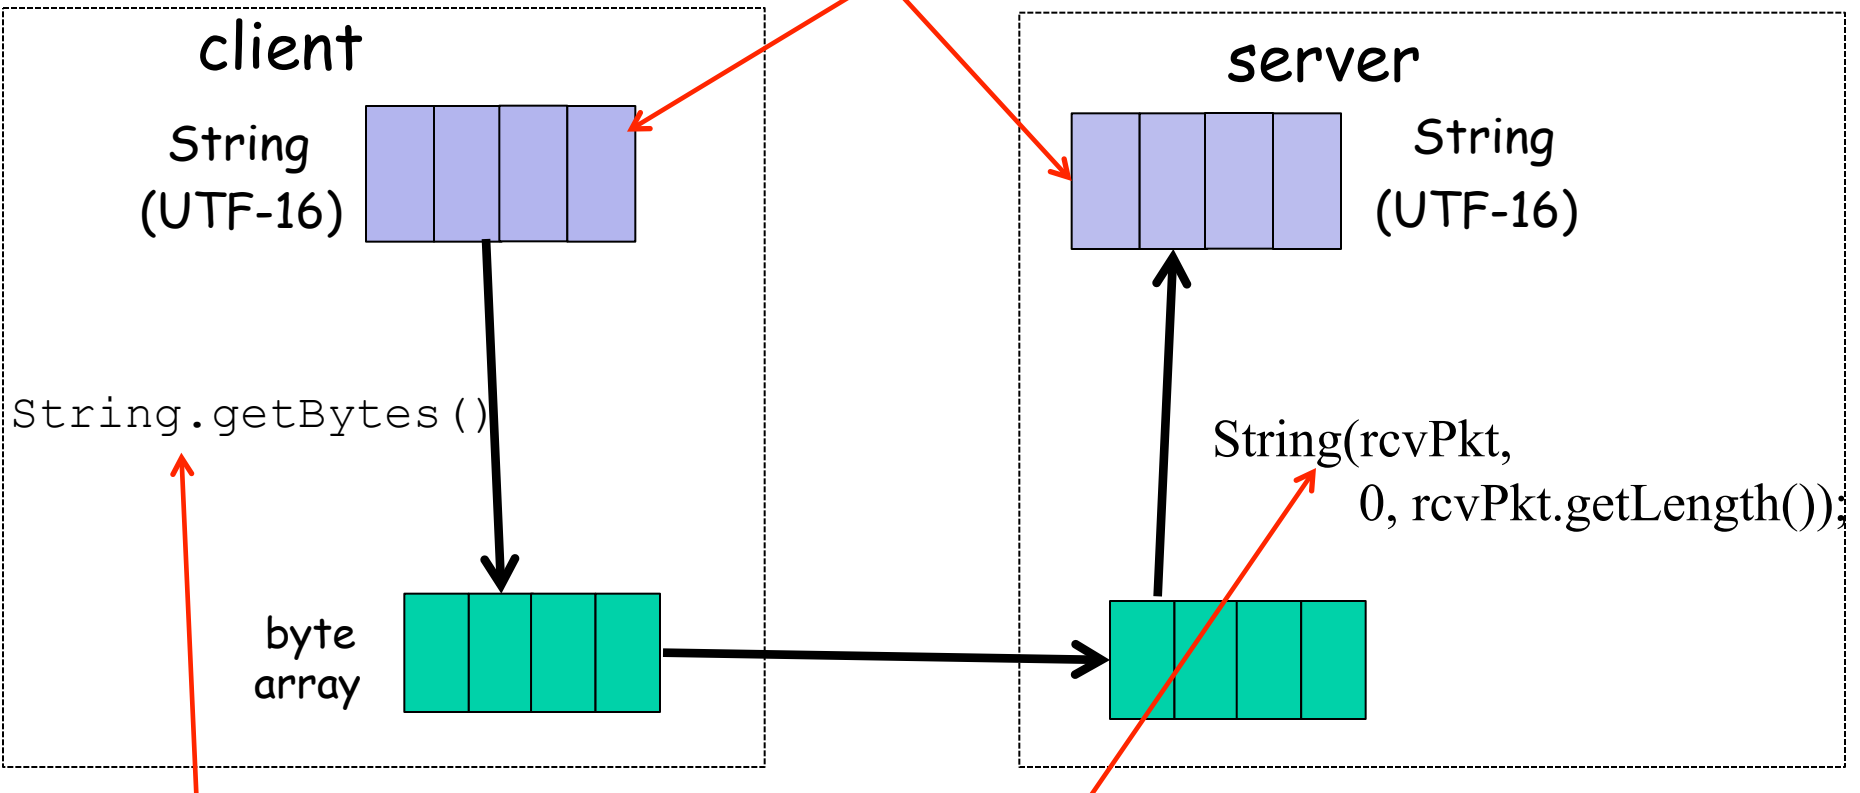
ARM, Power PC, Motorola 68k, IA-64



Intel x86

Example: String and Chars

Will we always get back the same string?



Depends on default local platform char set (why?) :
`java.nio.charset.Charset.defaultCharset()`

Example: Charset Troubles

- Try
 - java EncodingDecoding US-ASCII UTF-8

Encoding/Decoding as a Common Source of Errors

- ❑ Please read chapter 4 of Java Network Programming for more details
- ❑ Common mistake even in many (textbook) examples:
 - <http://www.java2s.com/Code/Java/Network-Protocol/UseDatagramSocketToSendOutAndReceiveDatagramPacket.htm>

DataStream

