

# Network Applications:

## Network App Programming: TCP

### FTP

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

2/8/2016

# Outline

---

- Admin and recap
- Network app programming
- FTP

## Recap: DNS Protocol, Messages

Many features: typically over **UDP** (can use **TCP**); **query** and **reply** messages with the **same message format**; **length/content encoding of names**; **simple compression**; **additional info as server push**

| Identification  | Flags                    | 12 bytes                                   |
|---|--------------------------|--|
| Number of questions   | Number of answer RRs     |  |
| Number of authority RRs   | Number of additional RRs |  |
| Questions<br>(variable number of questions)                     |                          | Name, type fields for a query              |
| Answers<br>(variable number of resource records)                |                          | RRs in response to query                   |
| Authority<br>(variable number of resource records)              |                          | Records for authoritative servers          |
| Additional information<br>(variable number of resource records) |                          | Additional “helpful” info that may be used |

# Recap: Connectionless UDP: Big Picture (Java version)

## Server (running on serv)

create socket,  
port=x, for  
incoming request:  
`serverSocket =  
DatagramSocket( x )`

```
read request from  
serverSocket  
generate reply, create  
datagram using client  
host address, port number
```

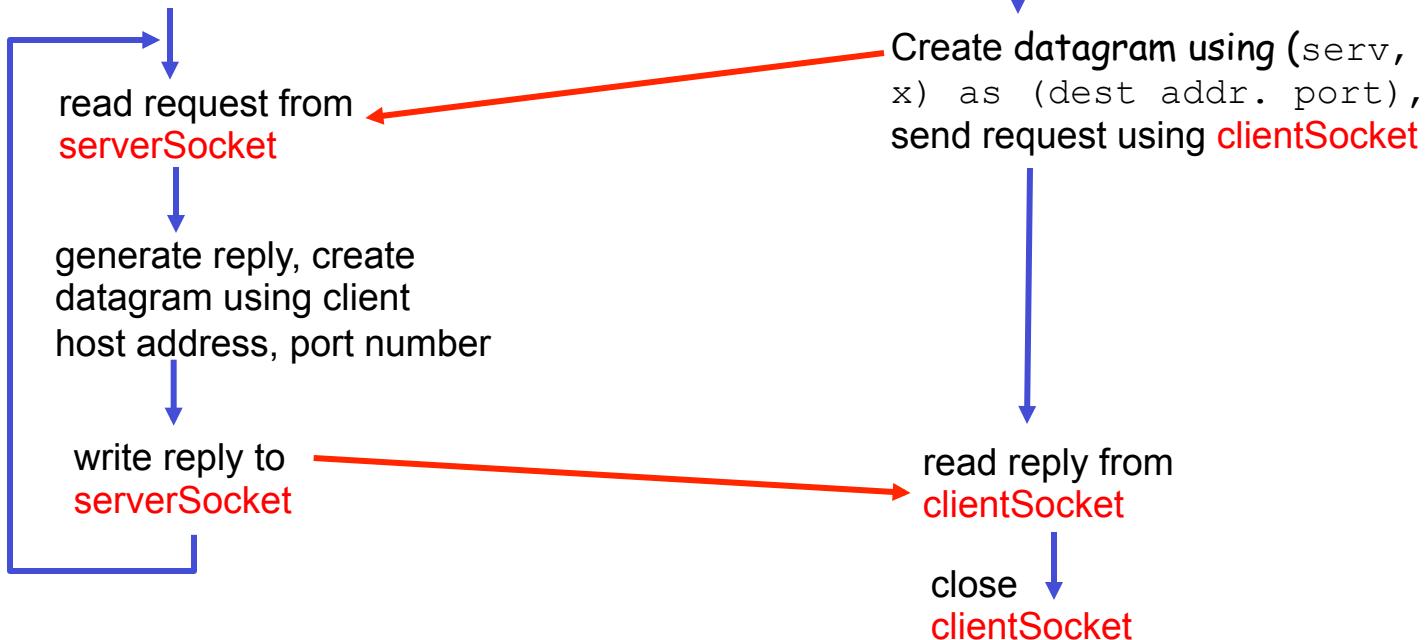
write reply to  
`serverSocket`

## Client

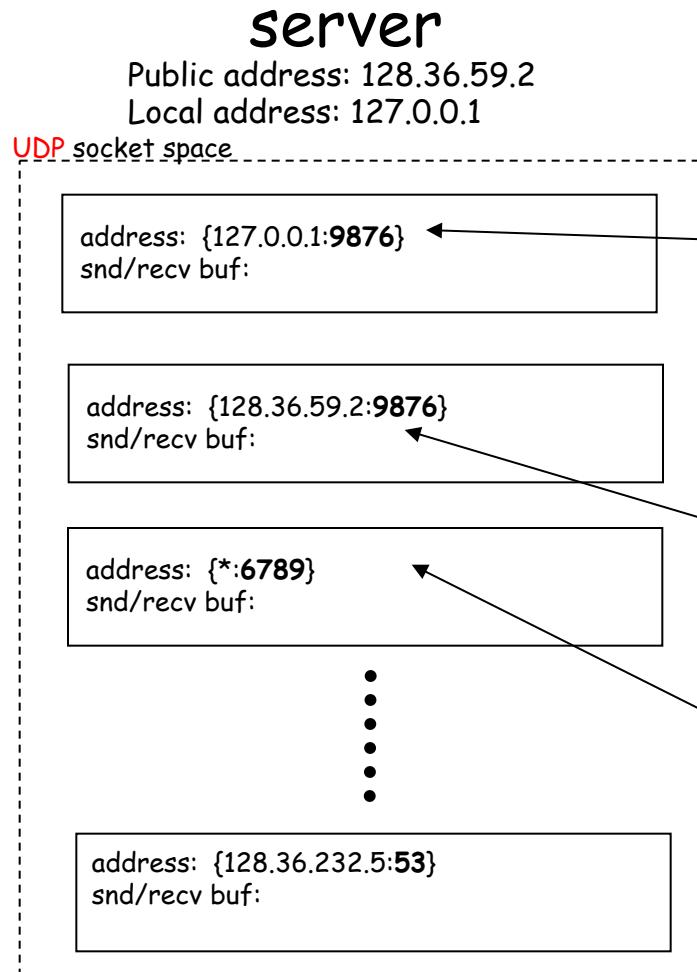
create socket,  
`clientSocket =  
DatagramSocket()`

Create datagram using (serv,  
x) as (dest addr. port),  
send request using `clientSocket`

```
read reply from  
clientSocket  
close  
clientSocket
```



# Recap: UDP Sockets



```
InetAddress sIP1 =  
InetAddress.getByName("localhost");  
DatagramSocket ssock1 =  
new DatagramSocket(9876, sIP1);
```

```
InetAddress sIP2 =  
InetAddress.getByName("128.36.59.2");  
DatagramSocket ssock2 =  
new DatagramSocket(9876, sIP2);
```

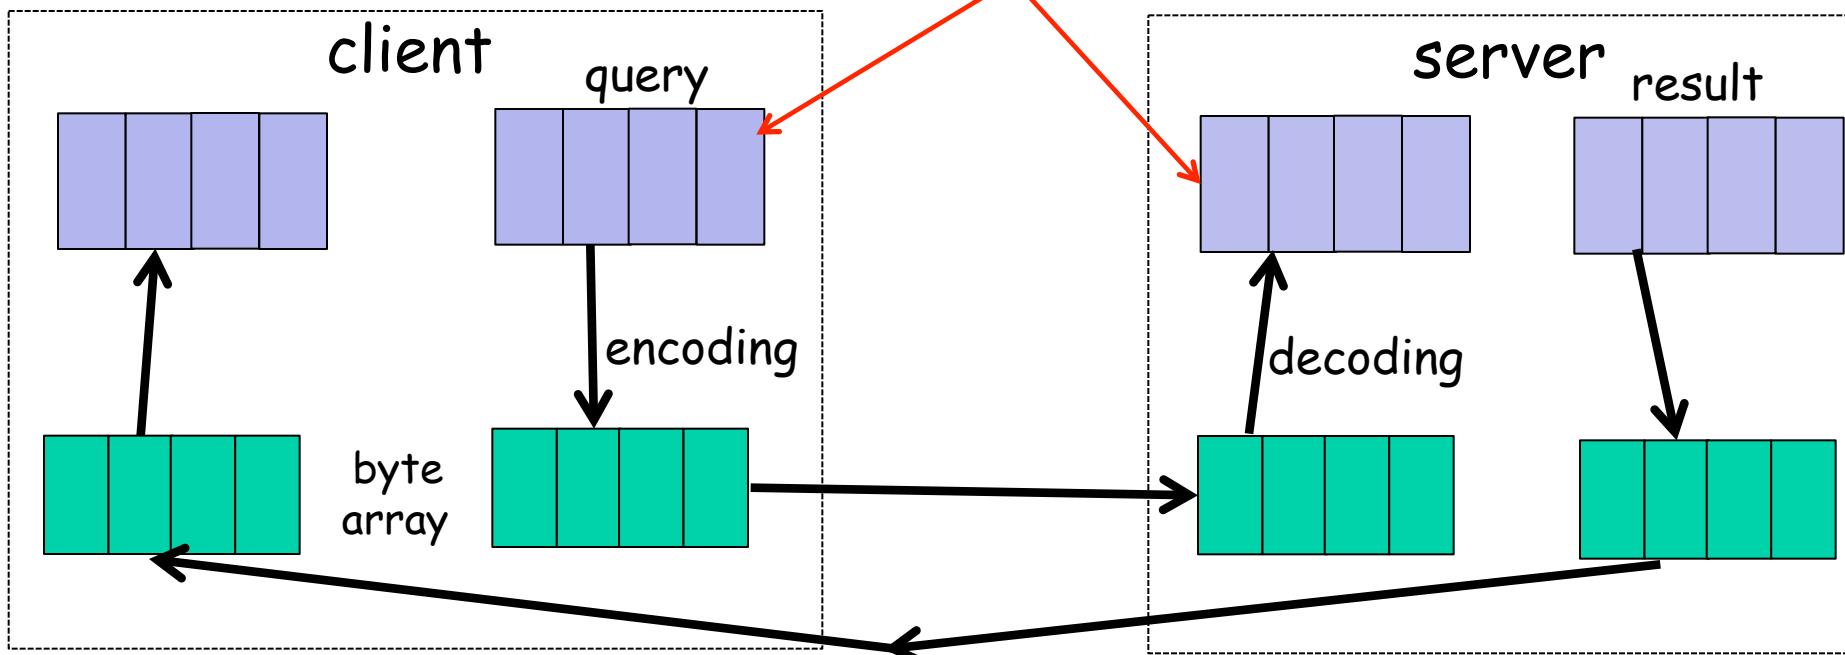
```
DatagramSocket serverSocket =  
new DatagramSocket(6789);
```

UDP demultiplexing is based on matching (dst address, dst port)

# Recap: Data Encoding/Decoding

- Pay attention to encoding/decoding of data:  
transport layer handles only a sequence of bytes

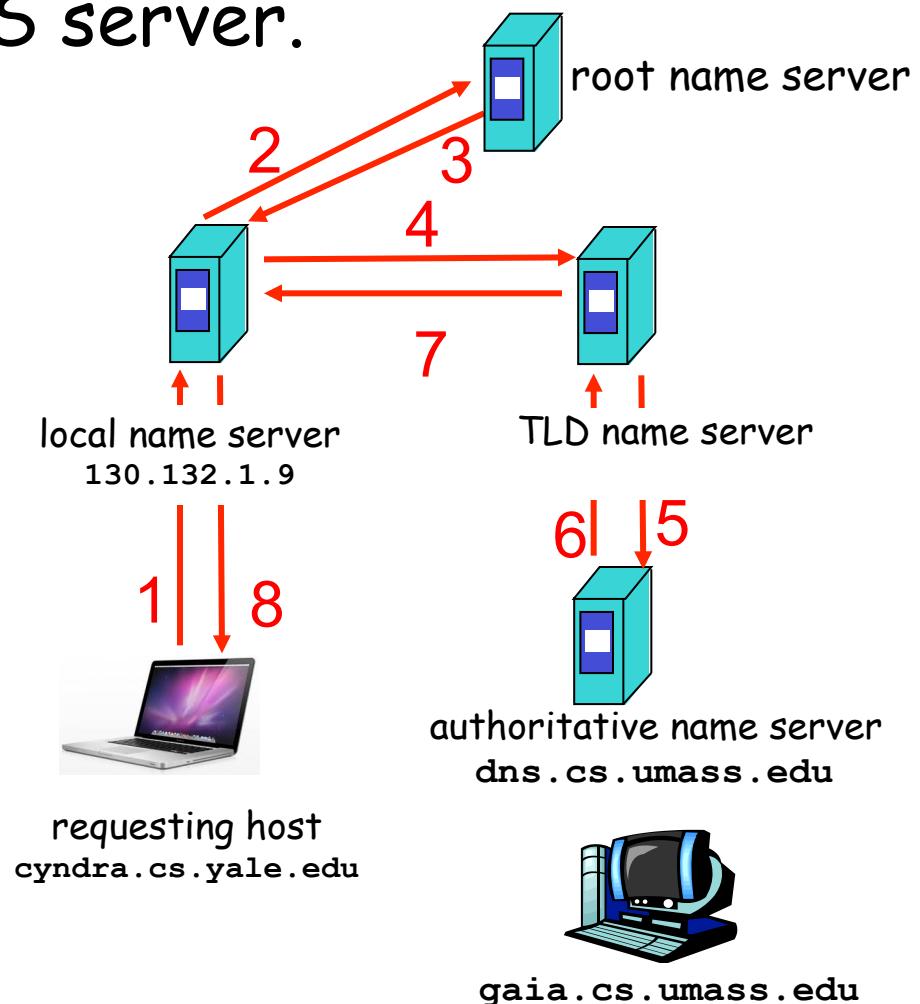
if not careful, query sent !=  
query received



# Discussion: UDP/DNS Server Pseudocode

- Modify the example UDP server code to implement a local DNS server.

|   |                          |  |
|---|--------------------------|--|
| Identification  | Flags                    | 12 bytes                                   |
| Number of questions   | Number of answer RRs     |  |
| Number of authority RRs   | Number of additional RRs |  |
| Questions<br>(variable number of questions)                     |                          | Name, type fields for a query              |
| Answers<br>(variable number of resource records)                |                          | RRs in response to query                   |
| Authority<br>(variable number of resource records)              |                          | Records for authoritative servers          |
| Additional information<br>(variable number of resource records) |                          | Additional "helpful" info that may be used |



# UDP/DNS Implementation

- Standard UDP demultiplexing (find out return address by src.addr/src.port of UDP packet) does not always work
- DNS solution:  
identification: remember the mapping

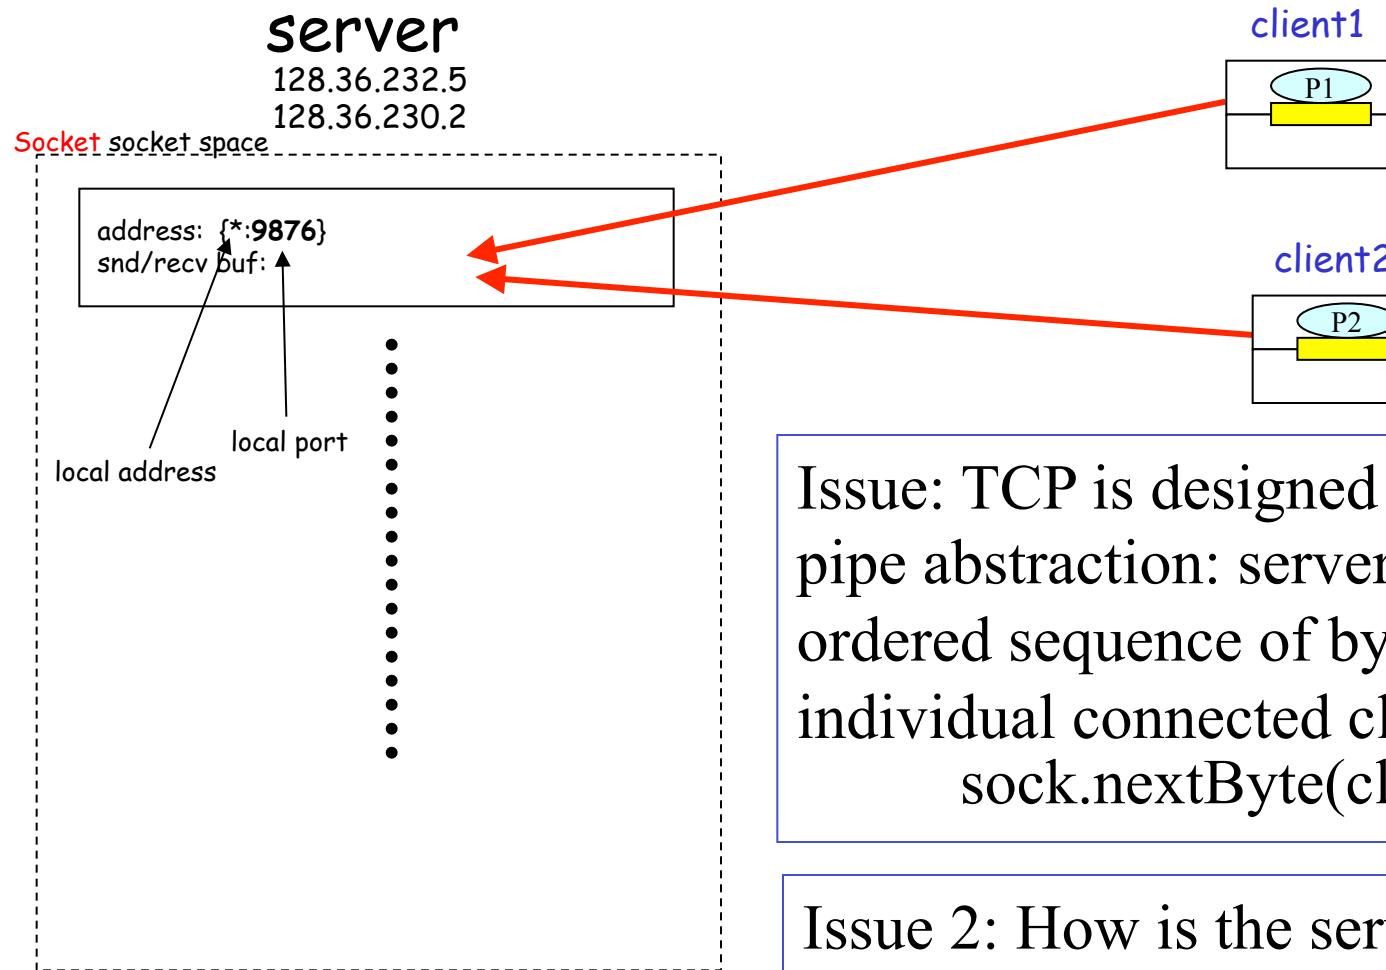
| Identification  | Flags                    | 12 bytes                                   |
|---|--------------------------|--|
| Number of questions   | Number of answer RRs     |  |
| Number of authority RRs   | Number of additional RRs |  |
| Questions<br>(variable number of questions)                     |                          | Name, type fields for a query              |
| Answers<br>(variable number of resource records)                |                          | RRs in response to query                   |
| Authority<br>(variable number of resource records)              |                          | Records for authoritative servers          |
| Additional information<br>(variable number of resource records) |                          | Additional "helpful" info that may be used |

# Outline

---

- Recap
- Network application programming
  - Overview
  - UDP
  - Basic TCP

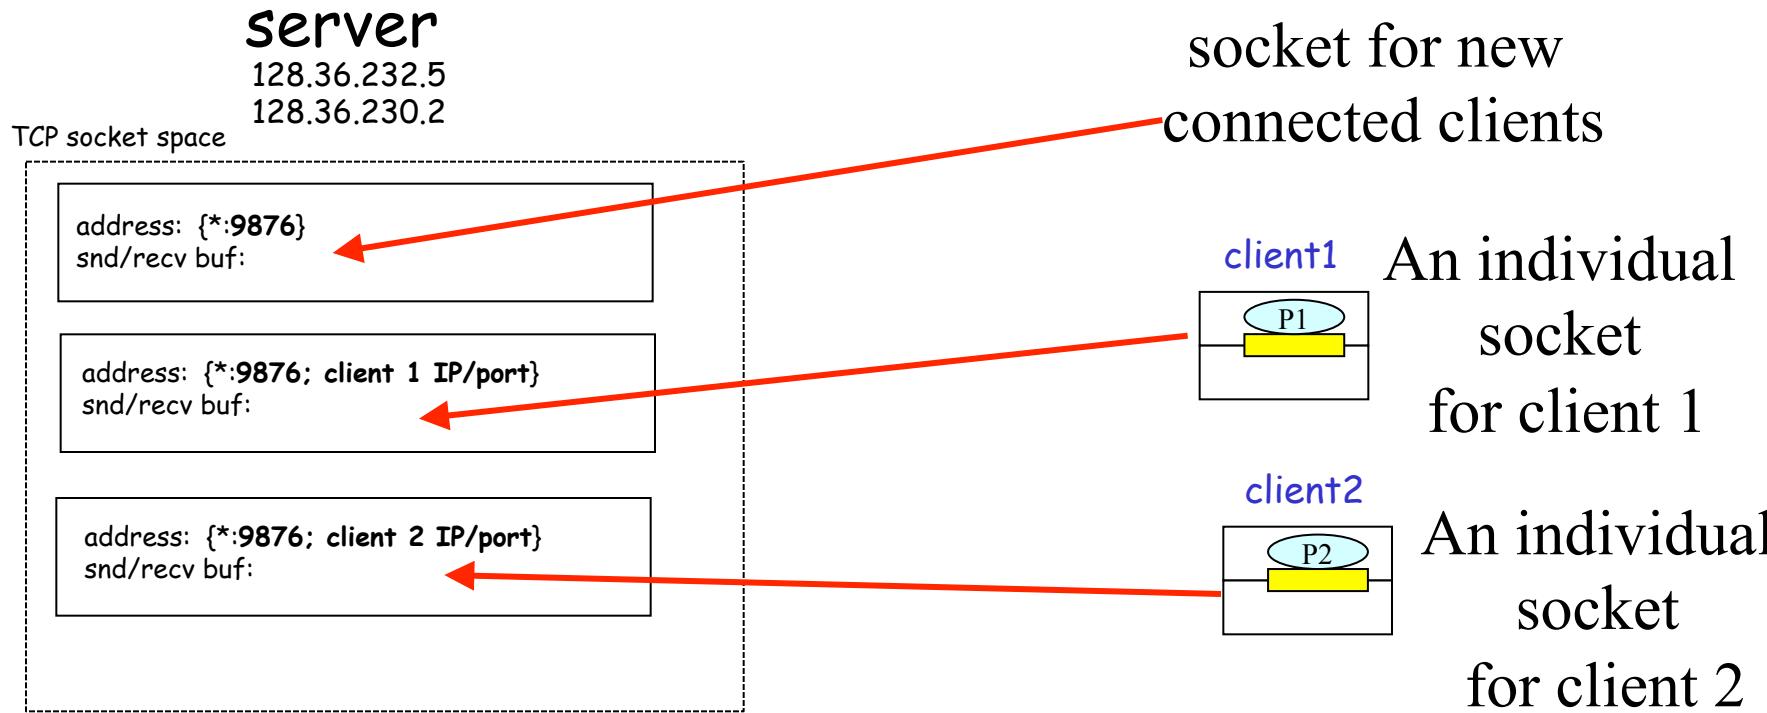
# TCP Socket Design: Starting w/ UDP



Issue: TCP is designed to provide a pipe abstraction: server reads an ordered sequence of bytes from each individual connected client  
`sock.nextByte(client1)?`

Issue 2: How is the server notified that a new client is connected?  
`newClient = sock.getNewClient()?`

# BSD TCP Socket API Design



Q: How to decide where to put a new packet?

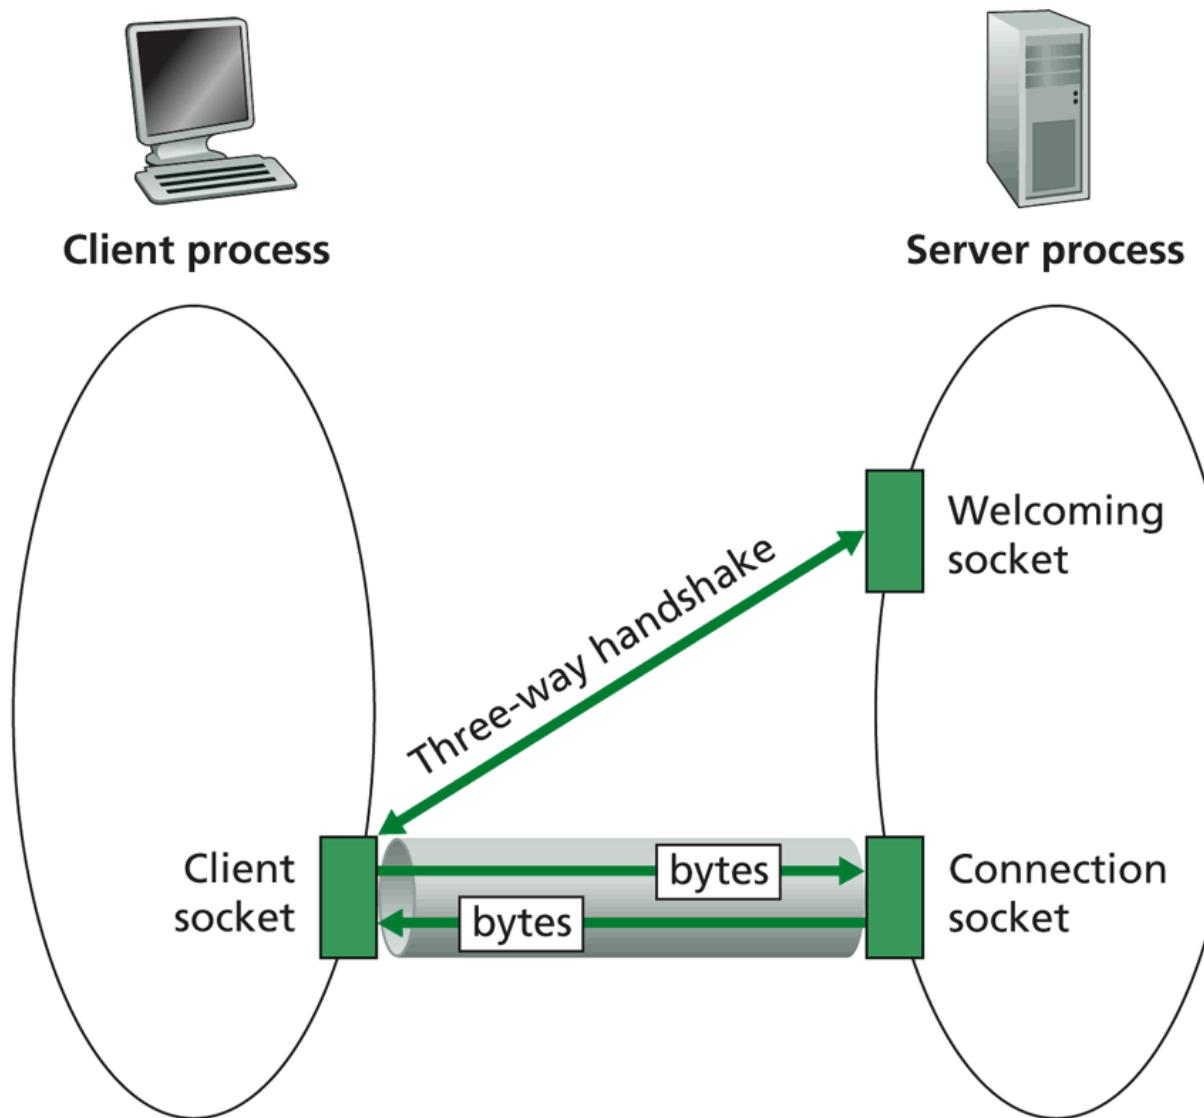
A: Packet demultiplexing is based on four tuples:  
(dst addr, dst port, src addr, src port)

# TCP Connection-Oriented Demux

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- recv host uses all four values to direct segment to appropriate socket
  - different connections/sessions are automatically separated into different sockets

-Welcome socket: the waiting room  
-connSocket: the operation room

# TCP Socket Big Picture



# Client/server Socket Workflow: TCP

## Server (running on hostid)

```
create socket,  
port=x, for  
incoming request:  
welcomeSocket =  
ServerSocket(x)
```

```
wait for incoming  
connection request  
connectionSocket =  
welcomeSocket.accept()
```

```
read request from  
connectionSocket
```

```
write reply to  
connectionSocket
```

```
close  
connectionSocket
```

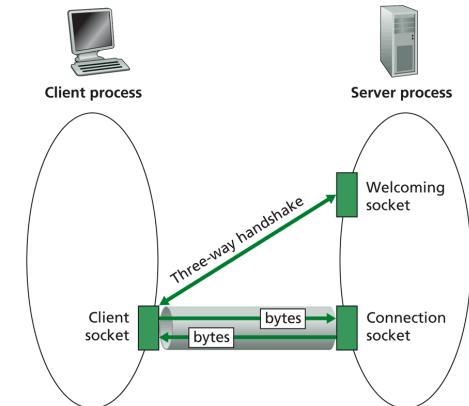
## Client

```
create socket,  
connect to hostid, port=x  
clientSocket =  
Socket()
```

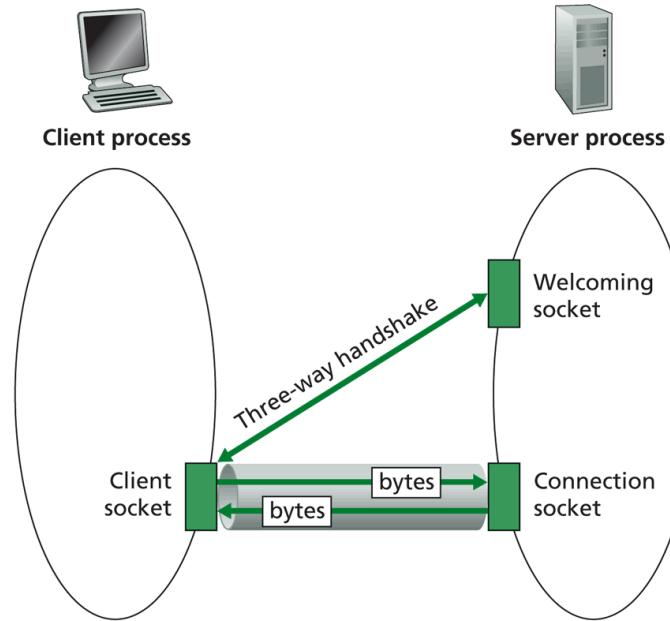
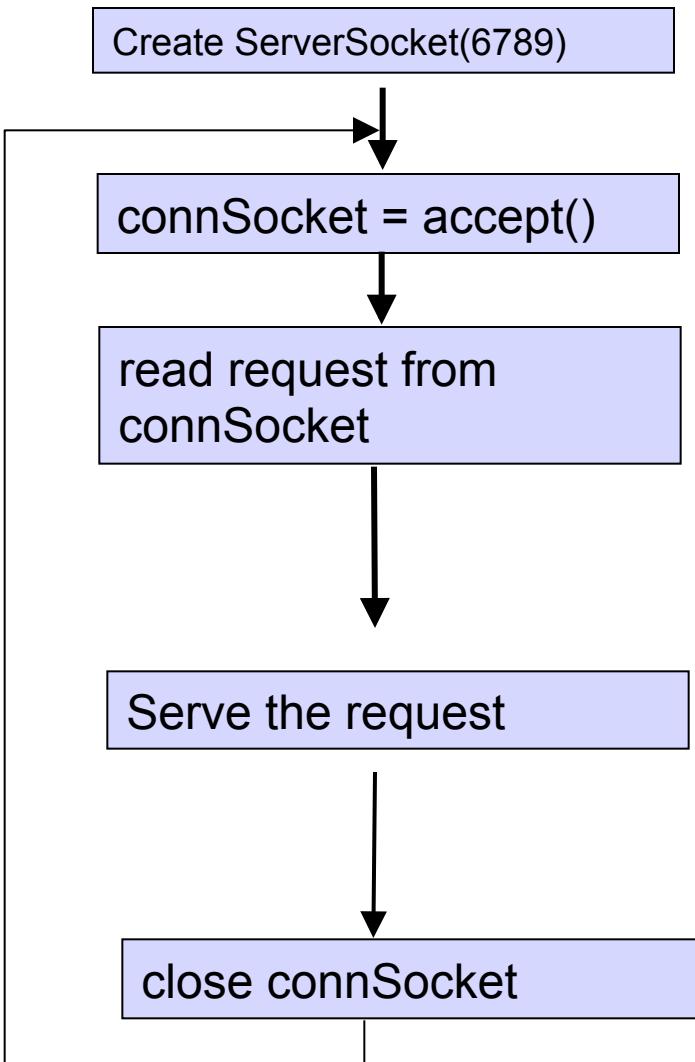
```
send request using  
clientSocket
```

```
read reply from  
clientSocket
```

```
close  
clientSocket
```



# Server Flow



-Welcome socket: the waiting room  
-connSocket: the operation room

# ServerSocket

- **ServerSocket()**
  - creates an unbound server socket.
- **ServerSocket(int port)**
  - creates a server socket, bound to the specified port.
- **ServerSocket(int port, int backlog)**
  - creates a server socket and binds it to the specified local port number, with the specified backlog.
- **ServerSocket(int port, int backlog, InetAddress bindAddr)**
  - creates a server with the specified port, listen backlog, and local IP address to bind to.
  
- **bind(SocketAddress endpoint)**
  - binds the ServerSocket to a specific address (IP address and port number).
- **bind(SocketAddress endpoint, int backlog)**
  - binds the ServerSocket to a specific address (IP address and port number).
  
- **Socket accept()**
  - listens for a connection to be made to this socket and accepts it.
  
- **close()**
  - closes this socket.

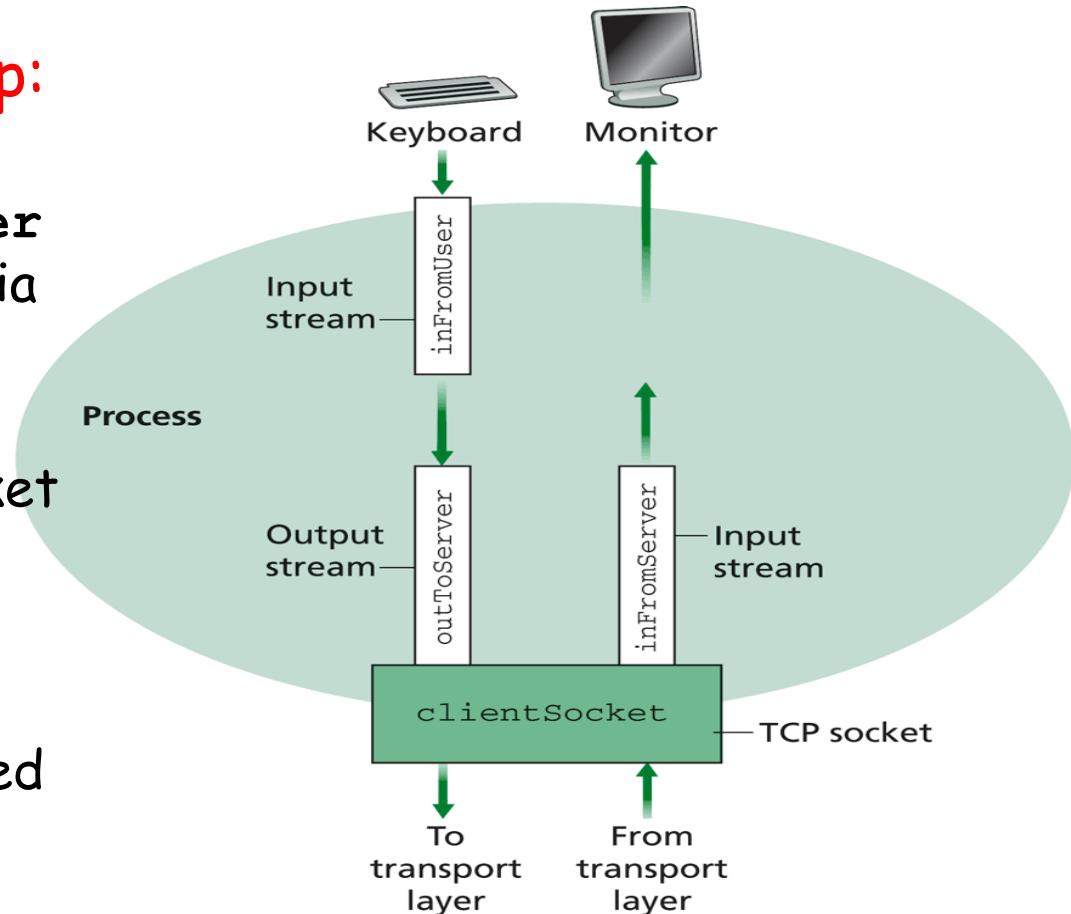
# (Client) Socket

- **Socket(InetAddress address, int port)**  
creates a stream socket and connects it to the specified port number at the specified IP address.
- **Socket(InetAddress address, int port, InetAddress localAddr, int localPort)**  
creates a socket and connects it to the specified remote address on the specified remote port.
- **Socket(String host, int port)**  
creates a stream socket and connects it to the specified port number on the named host.
- **bind(SocketAddress bindpoint)**  
binds the socket to a local address.
- **connect(SocketAddress endpoint)**  
connects this socket to the server.
- **connect(SocketAddress endpoint, int timeout)**  
connects this socket to the server with a specified timeout value.
- **InputStream get InputStream()**  
returns an input stream for this socket.
- **OutputStream get OutputStream()**  
returns an output stream for this socket.
- **close()**  
closes this socket.

# Simple TCP Example

## Example client-server app:

- 1) client reads line from standard input (`inFromUser` stream), sends to server via socket (`outToServer` stream)
- 2) server reads line from socket
- 3) server converts line to uppercase, sends back to client
- 4) client reads, prints modified line from socket (`inFromServer` stream)



# Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPClient {

    public static void main(String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;

        Create input stream → BufferedReader inFromUser =
            new BufferedReader(new InputStreamReader(System.in));
        sentence = inFromUser.readLine();

        Create client socket, connect to server → Socket clientSocket = new Socket("server.name", 6789);

        Create output stream attached to socket → DataOutputStream outToServer =
            new DataOutputStream(clientSocket.getOutputStream());
    }
}
```

# OutputStream

- **public abstract class OutputStream**
  - **public abstract void write(int b) throws IOException**
  - **public void write(byte[] data) throws IOException**
  - **public void write(byte[] data, int offset, int length) throws IOException**
  - **public void flush( ) throws IOException**
  - **public void close( ) throws IOException**

# InputStream

- public abstract class InputStream
  - public abstract int read( ) throws IOException
  - public int read(byte[] input) throws IOException
  - public int read(byte[] input, int offset, int length) throws IOException
  - public long skip(long n) throws IOException
  - public int available( ) throws IOException
  - public void close( ) throws IOException

## Example: Java client (TCP), cont.

```
Send line  
to server ]→ outToServer.writeBytes(sentence + '\n');

Create  
input stream  
attached to socket ]→ BufferedReader inFromServer =  
new BufferedReader(new  
InputStreamReader(clientSocket.getInputStream()));

modifiedSentence = inFromServer.readLine();

System.out.println("FROM SERVER: " + modifiedSentence);

clientSocket.close();

}

}
```

# Example: Java server (TCP)

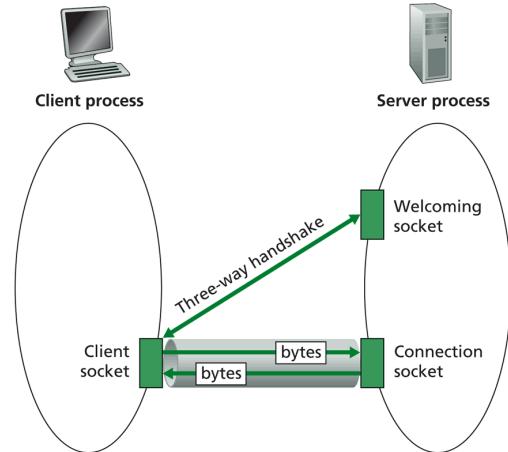
```
import java.io.*;
import java.net.*;

class TCPServer {

    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        ServerSocket welcomeSocket = new ServerSocket(6789);
```

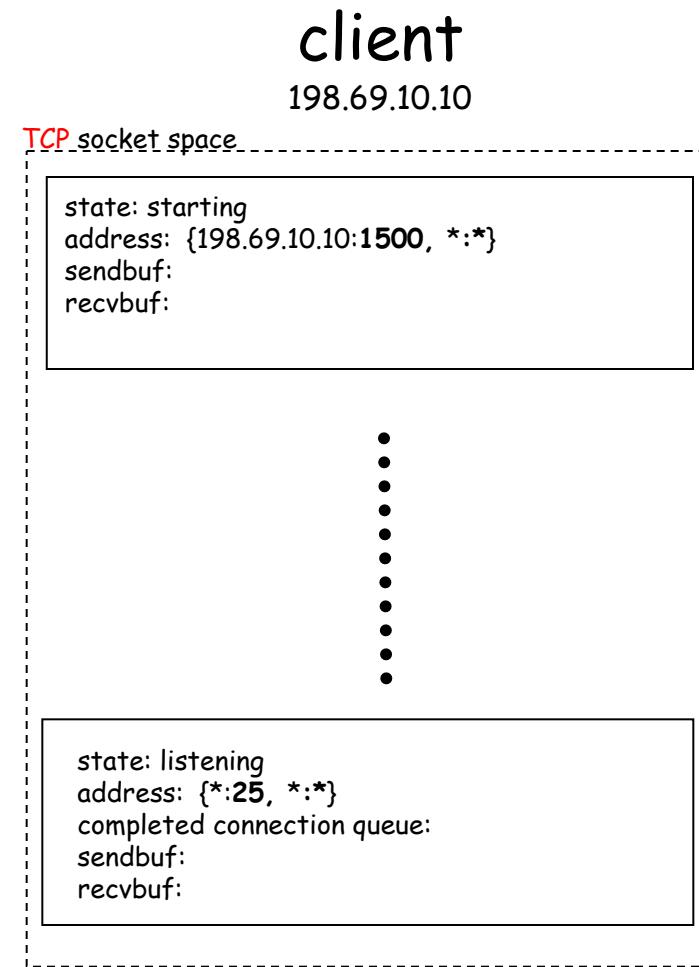
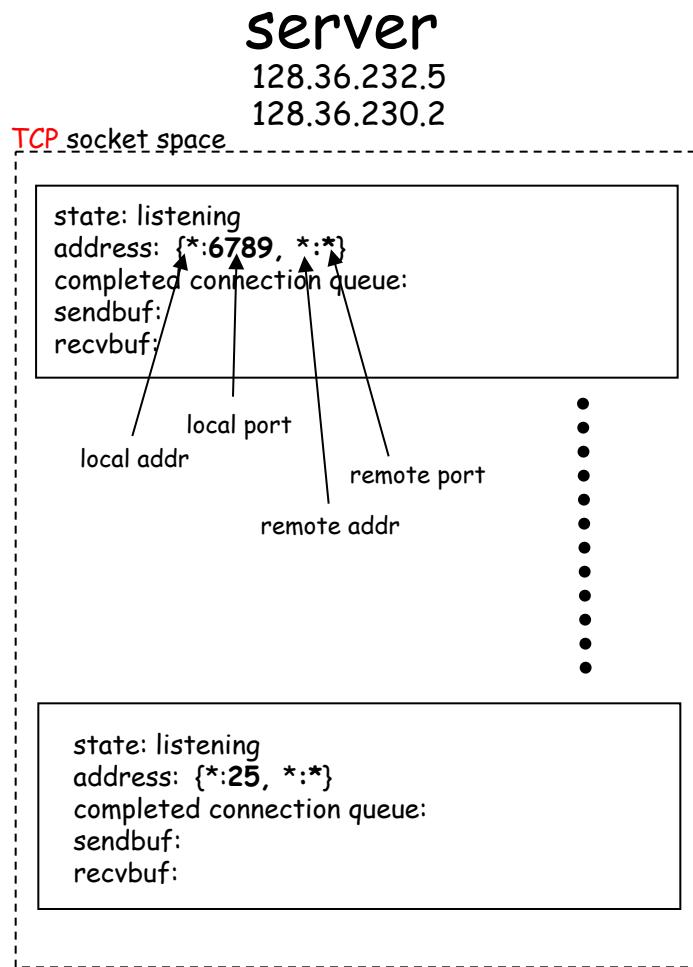
*Create welcoming socket at port 6789*



# Demo

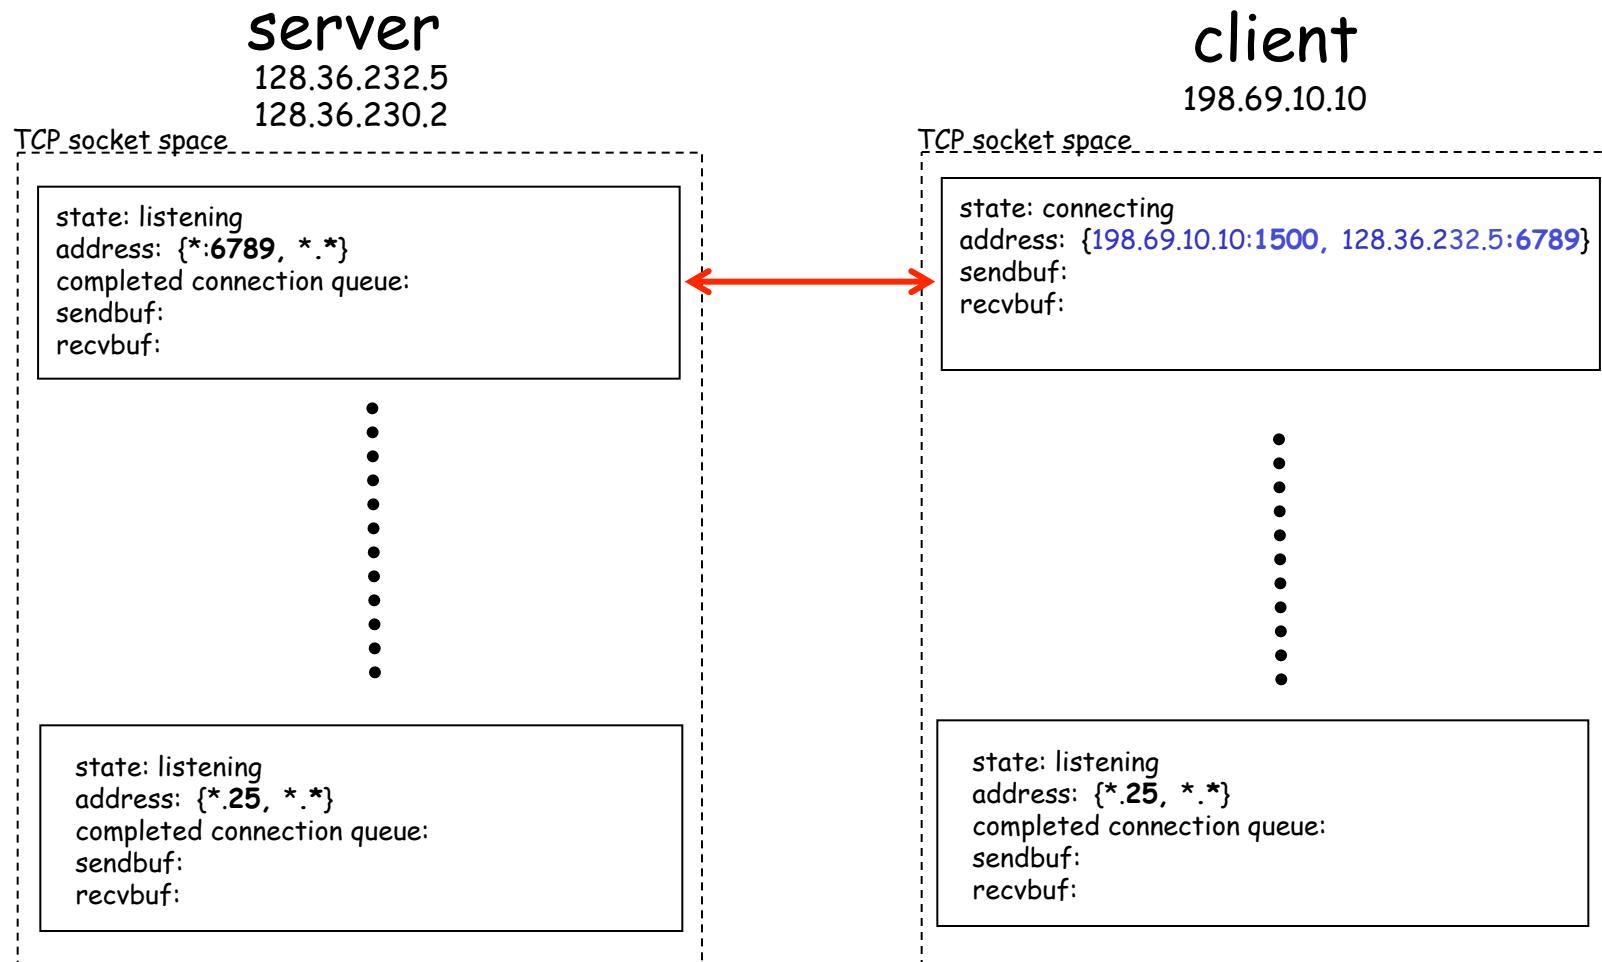
```
% wireshark to capture our TCP traffic  
tcp.srcport==6789 or tcp.dstport==6789
```

# Under the Hood: After Welcome (Server) Socket



%netstat -p tcp -n -a

# Client Initiates Connection



# Example: Client Connection

## Handshake Done

**server**

128.36.232.5  
128.36.230.2

TCP socket space

```
state: listening
address: {*:6789, *:*}
completed connection queue:
{128.36.232.5:6789, 198.69.10.10.1500}
sendbuf:
recvbuf:
```

•  
•  
•  
•  
•  
•  
•

```
state: listening
address: {*:25, *:*}
completed connection queue:
sendbuf:
recvbuf:
```

**client**

198.69.10.10

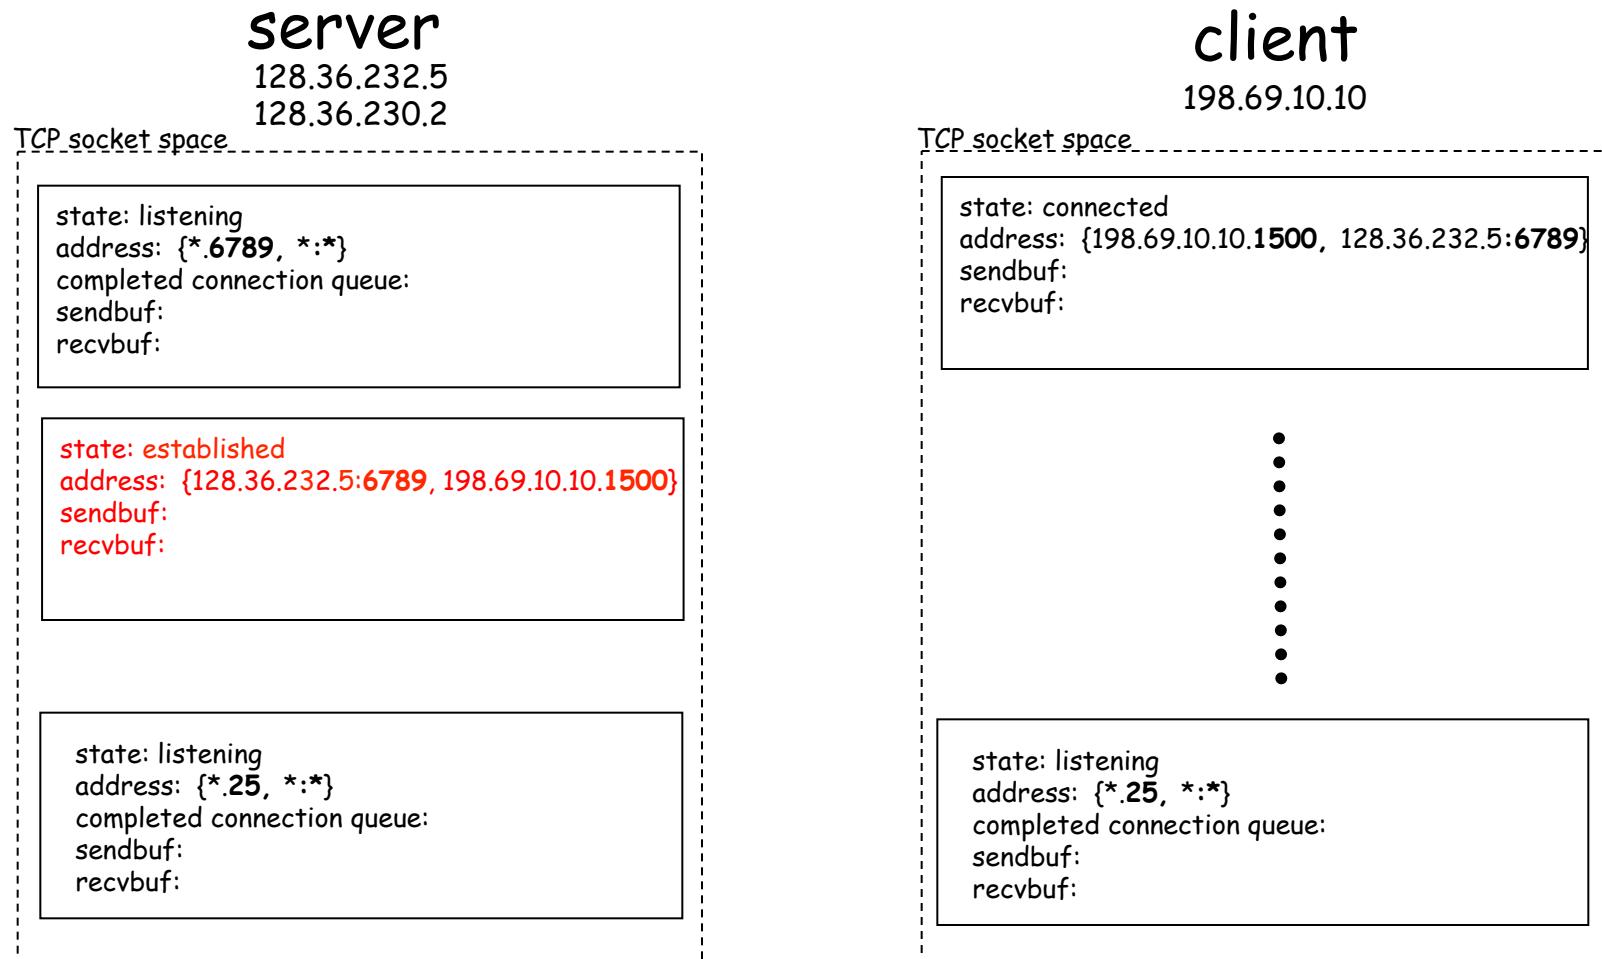
TCP socket space

```
state: connected
address: {198.69.10.10:1500, 128.36.232.5:6789}
sendbuf:
recvbuf:
```

•  
•  
•  
•  
•  
•

```
state: listening
address: {*:25, *:*}
completed connection queue:
sendbuf:
recvbuf:
```

# Example: Client Connection Handshake Done



Packet demultiplexing is based on (dst addr, dst port, src addr, src port)  
Packet sent to the socket with **the best match!**

# Demo

---

- ❑ What if more client connections than backlog allowed?

# Example: Java server (TCP)

```
import java.io.*;
import java.net.*;

class TCPServer {

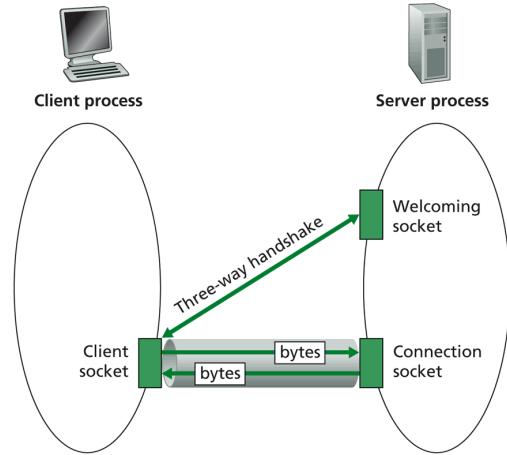
    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        ServerSocket welcomeSocket = new ServerSocket(6789);

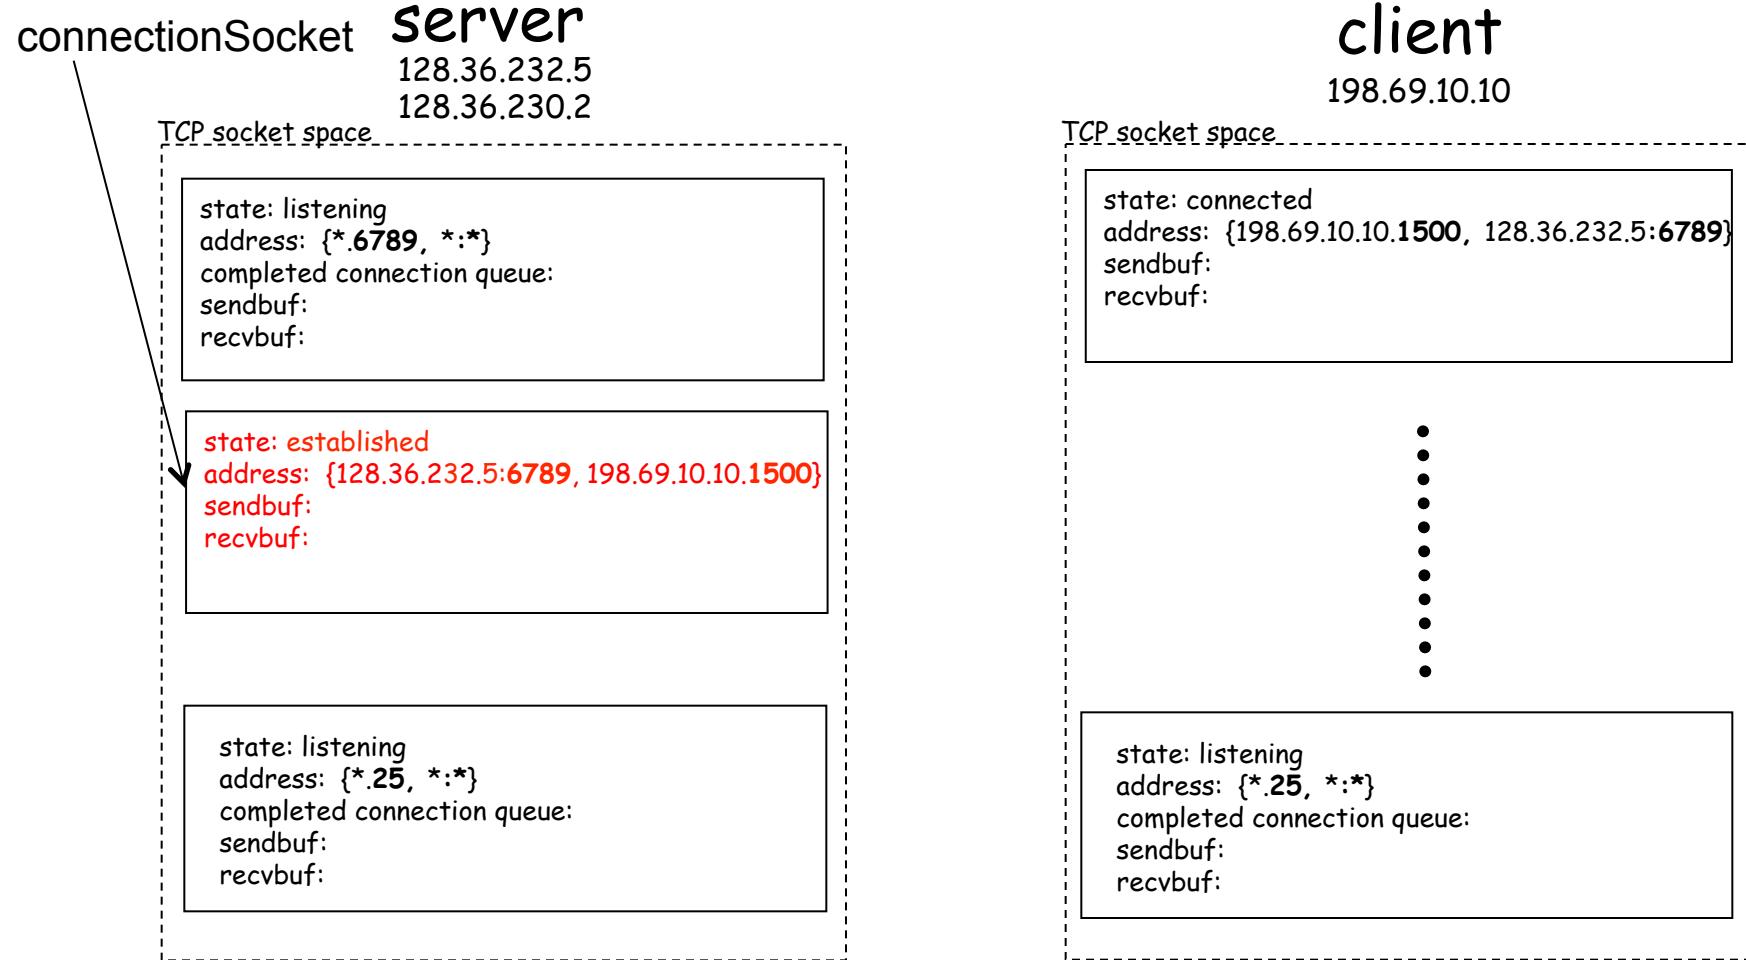
        while(true) {
            Socket connectionSocket = welcomeSocket.accept();

```

Wait, on welcoming socket for contact by client



# Example: Server accept()



# Example: Java server (TCP):

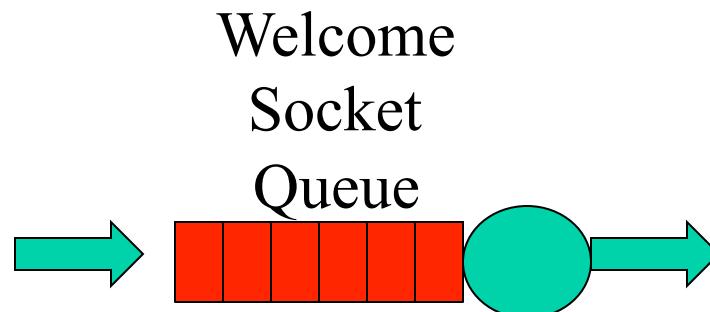
## Processing

```
Create input stream, attached to socket → BufferedReader inFromClient =  
new BufferedReader(new InputStreamReader(connectionSocket.getInputStream()));  
  
Read in line from socket → clientSentence = inFromClient.readLine();  
capitalizedSentence = clientSentence.toUpperCase() + '\n';  
  
}  
}  
}
```

## Example: Java server (TCP): Output

# Analysis

- Assume that client requests arrive at a rate of  $\lambda$ /second
- Assume that each request takes  $1/\mu$  seconds
- Some basic questions
  - How big is the backlog (welcome queue)



# Analysis

---

- ❑ Is there any interop issue in the sample program?

# Analysis

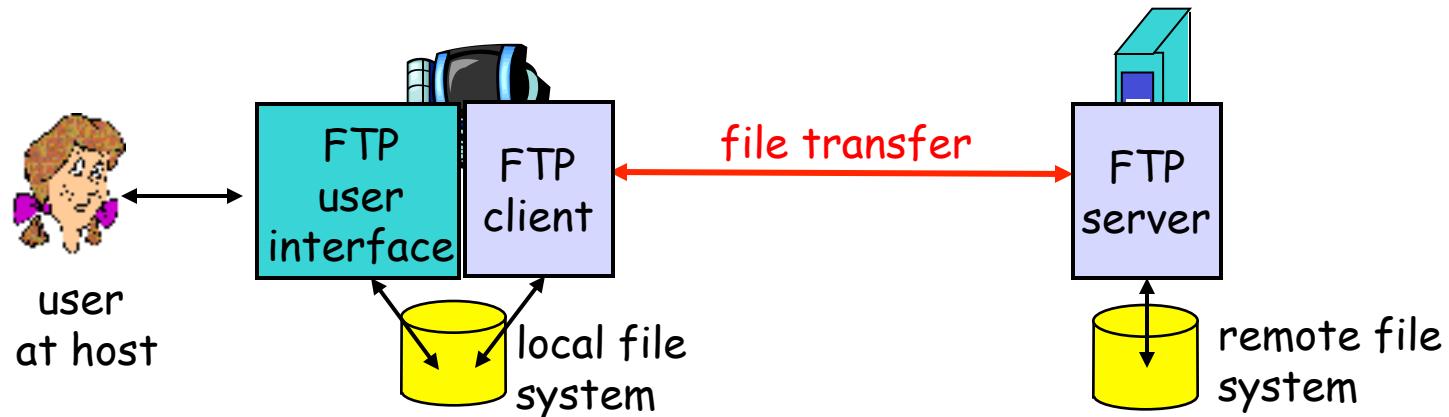
- Is there any interop issue in the sample program?
  - DataOutputStream writeBytes(String) truncates
    - [http://docs.oracle.com/javase/1.4.2/docs/api/java/io/DataOutputStream.html#writeBytes\(java.lang.String\)](http://docs.oracle.com/javase/1.4.2/docs/api/java/io/DataOutputStream.html#writeBytes(java.lang.String))

# Outline

---

- Recap
- Network application programming
- FTP

# FTP: the File Transfer Protocol



- Transfer files to/from remote host
- Client/server model
  - client*: side that initiates transfer (either to/from remote)
  - server*: remote host
- ftp: RFC 959
- ftp server: port 21/20 (smtp 25, http 80)

# FTP Commands, Responses

## Sample commands:

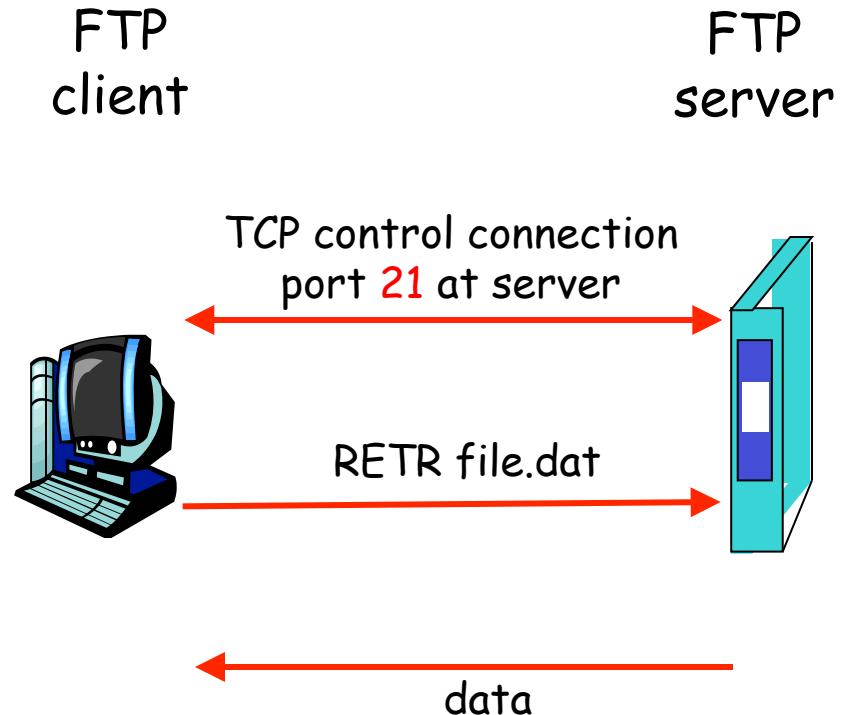
- sent as ASCII text over control channel
- USER *username***
- PASS *password***
- PWD** returns current dir
- STAT** shows server status
- LIST** returns list of file in current directory
- RETR *filename*** retrieves (gets) file
- STOR *filename*** stores file

## Sample return codes

- status code and phrase
- 331 Username OK, password required**
- 125 data connection already open; transfer starting**
- 425 Can't open data connection**
- 452 Error writing file**

# FTP Protocol Design

- What is the simplest design of data transfer?



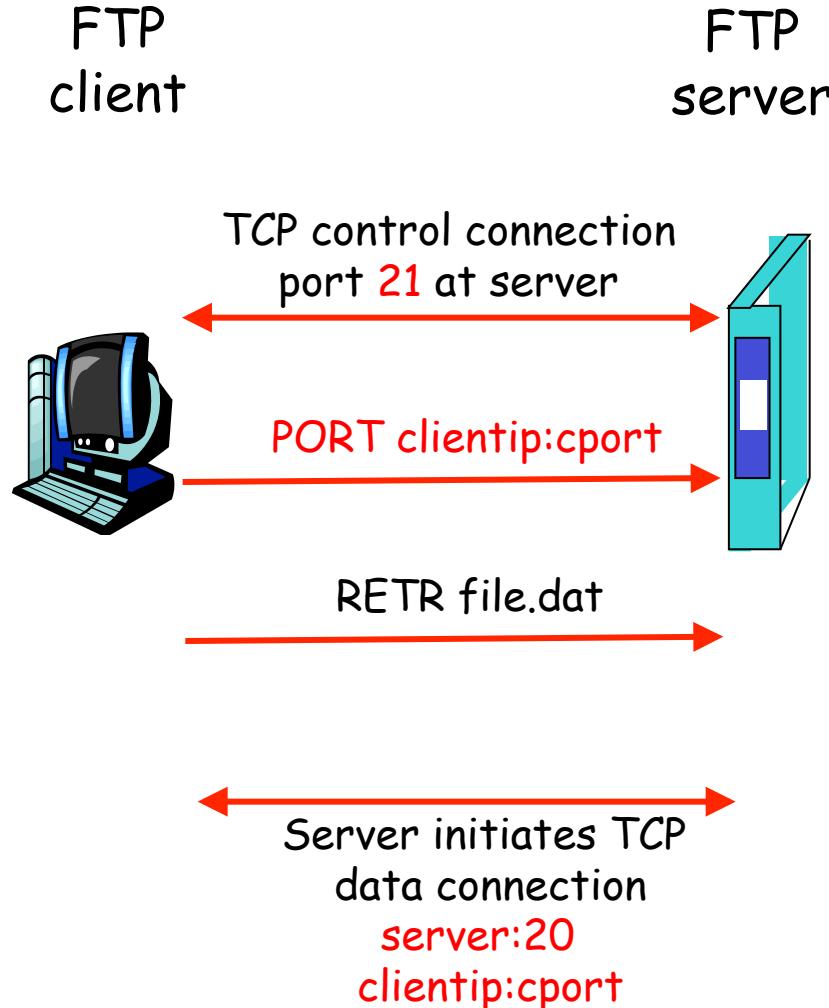
# FTP: A Client-Server Application with Separate Control, Data Connections

- Two types of TCP connections opened:
  - A control connection: exchange commands, responses between client, server.  
“out of band control”
  - Data connections: each for file data to/  
from server

Discussion: why does FTP separate control/data connections?

Q: How to create a new data connection?

# Traditional FTP: Client Specifies Port for Data Connection

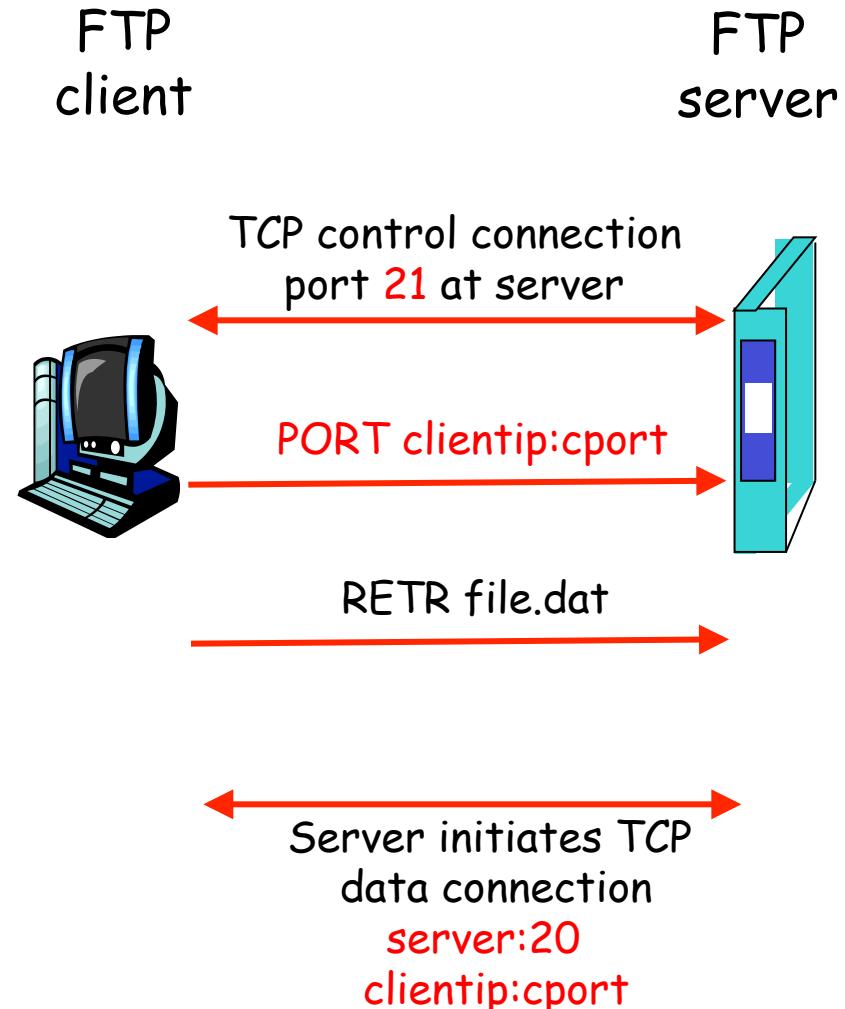


## Example using telnet/nc

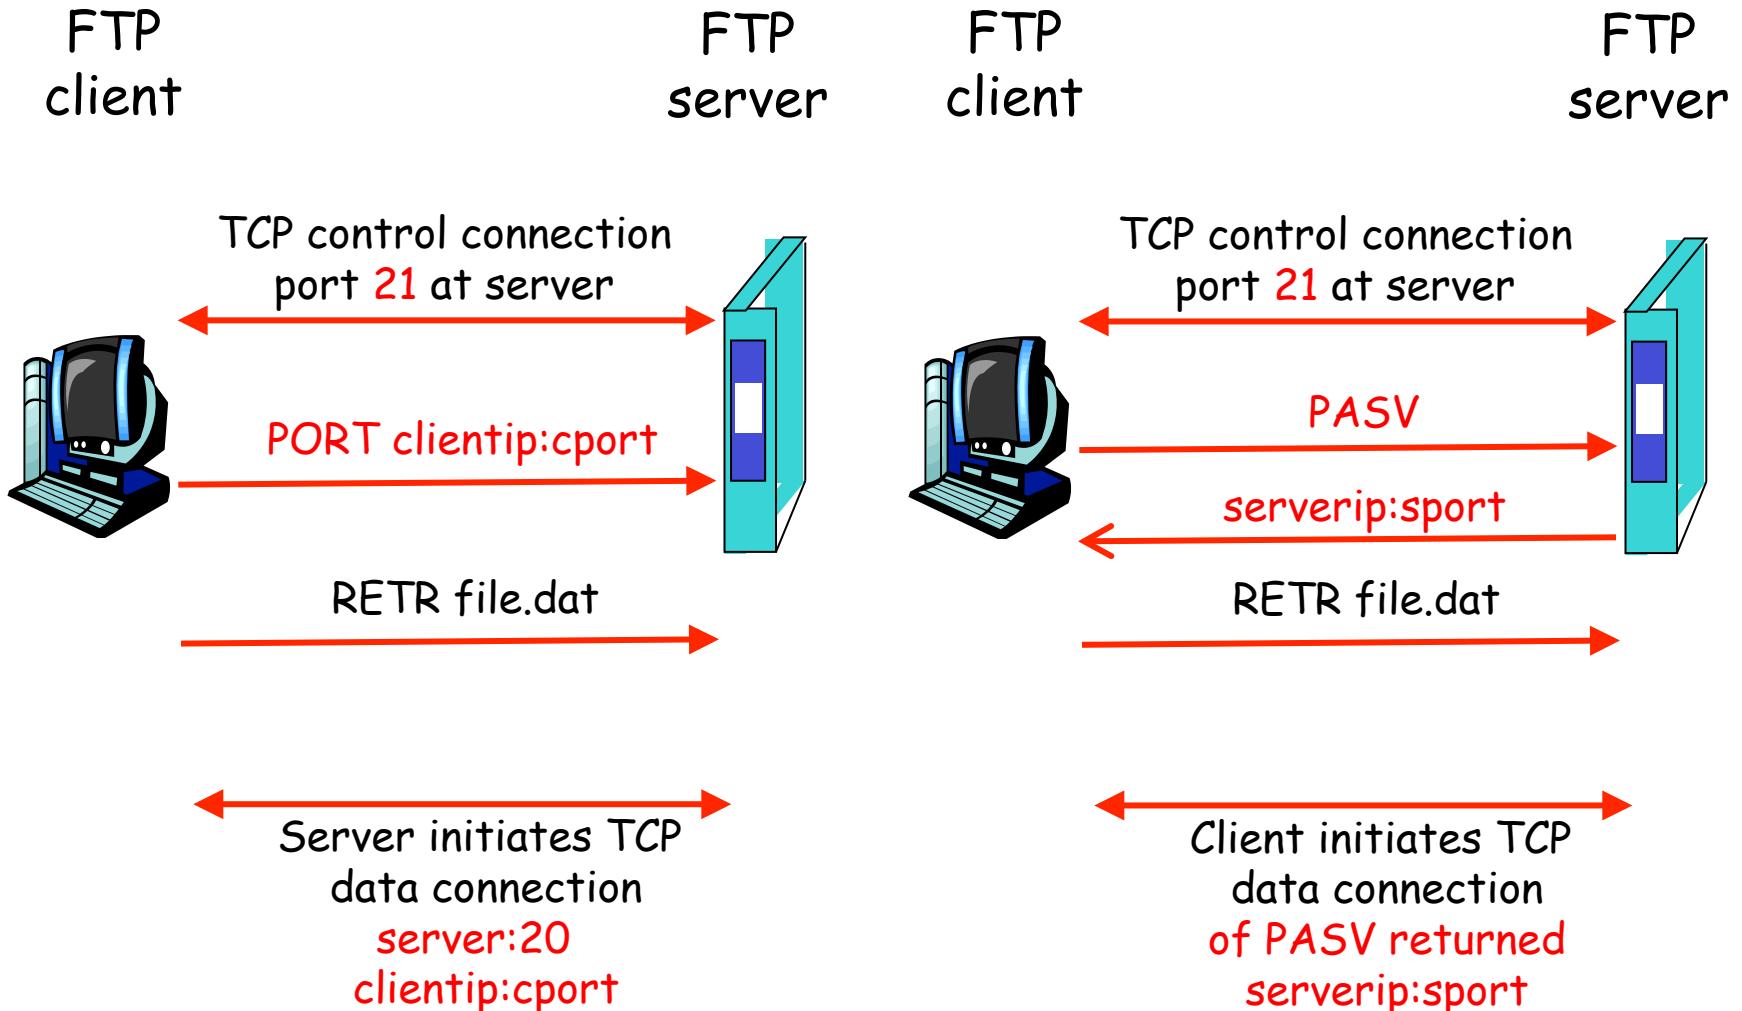
- Use telnet for the control channel
  - telnet ftp.freebsd.org 21
  - user, pass
  - port 172,27,5,145,4,1
  - list
  
- use nc (NetCat) to receive/send data with server
  - nc -v -l 1025

# Problem of the Client PORT Approach

- Many Internet hosts are behind **NAT/firewalls** that block connections initiated from outside



# FTP PASV: Server Specifies Data Port, Client Initiates Connection



# Example

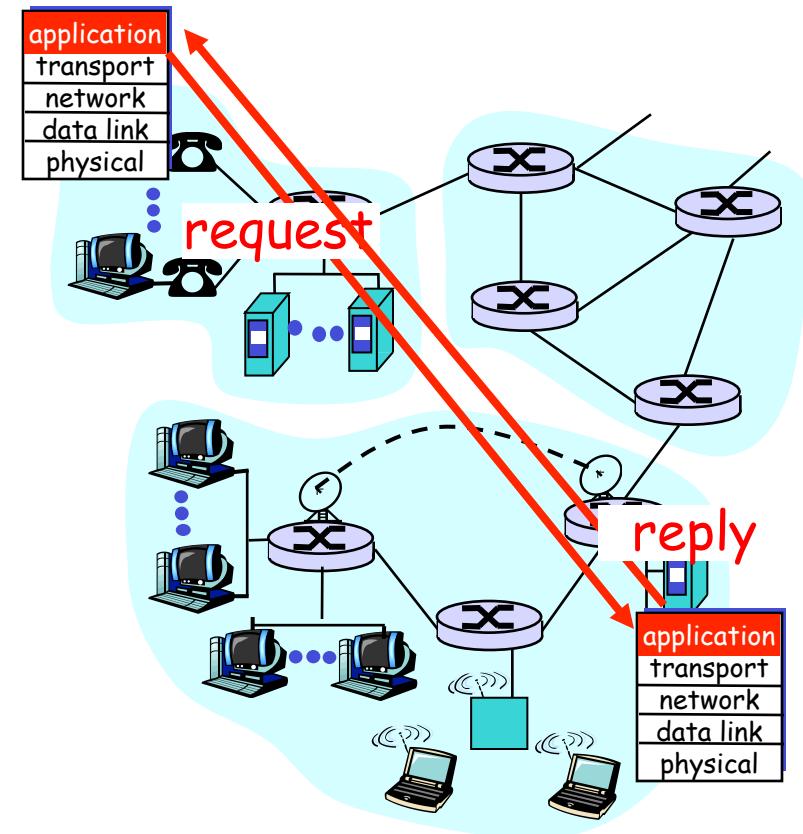
---

- Use Wireshark to capture traffic
  - Using chrome to visit  
<ftp://ftp.freebsd.org/pub/FreeBSD/README.TXT>

# FTP Evaluation

Key questions to ask about a C-S application

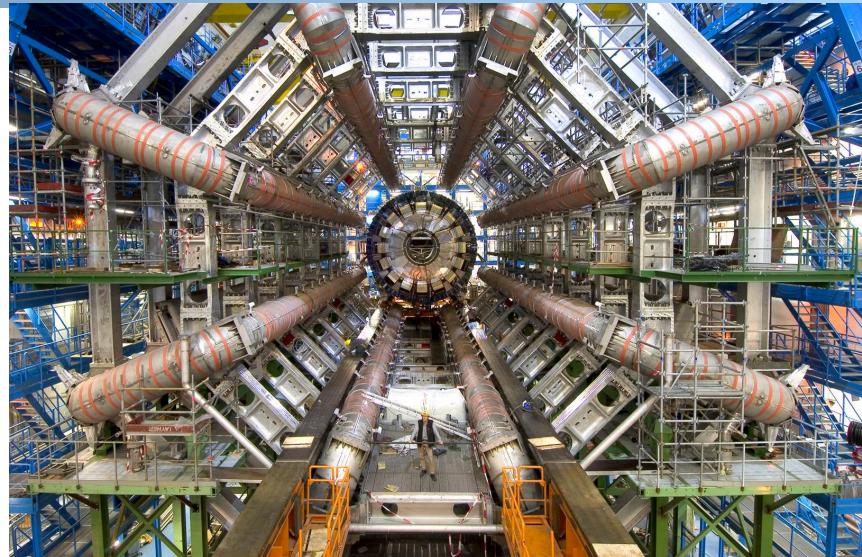
- Is the application **extensible**?
- Is the application **scalable**?
- How does the application handle server failures (being **robust**)?
- How does the application provide **security**?



What are some interesting design features of the FTP protocol?

# FTP Extensions

- ❑ FTP extensions are still being used extensively in large data set transfers (e.g., LHC)
  
- ❑ See GridFTP to FTP extensions
  - <https://www.ogf.org/documents/GFD.20.pdf>



# Outline

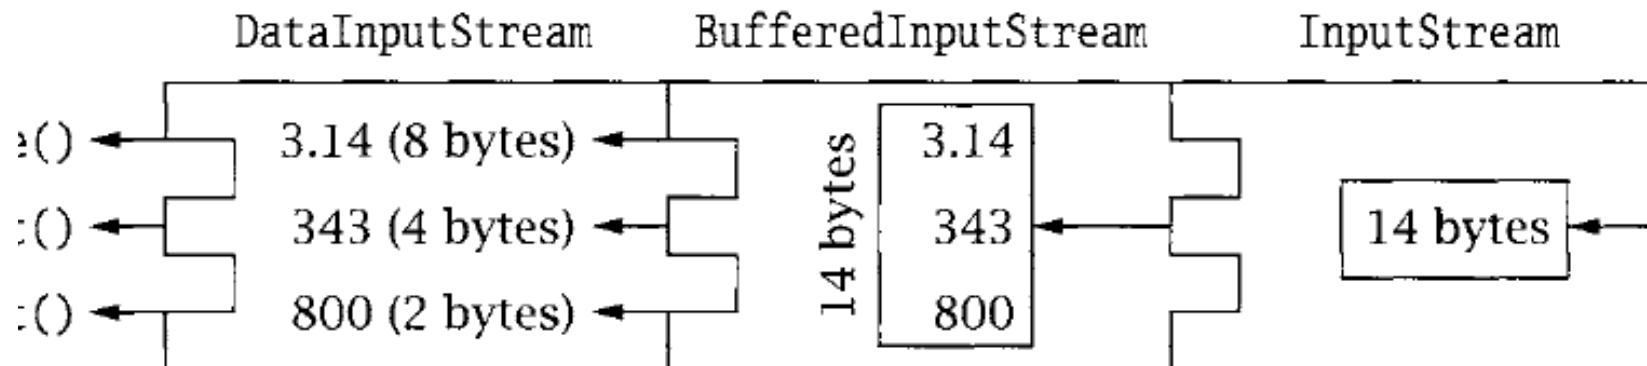
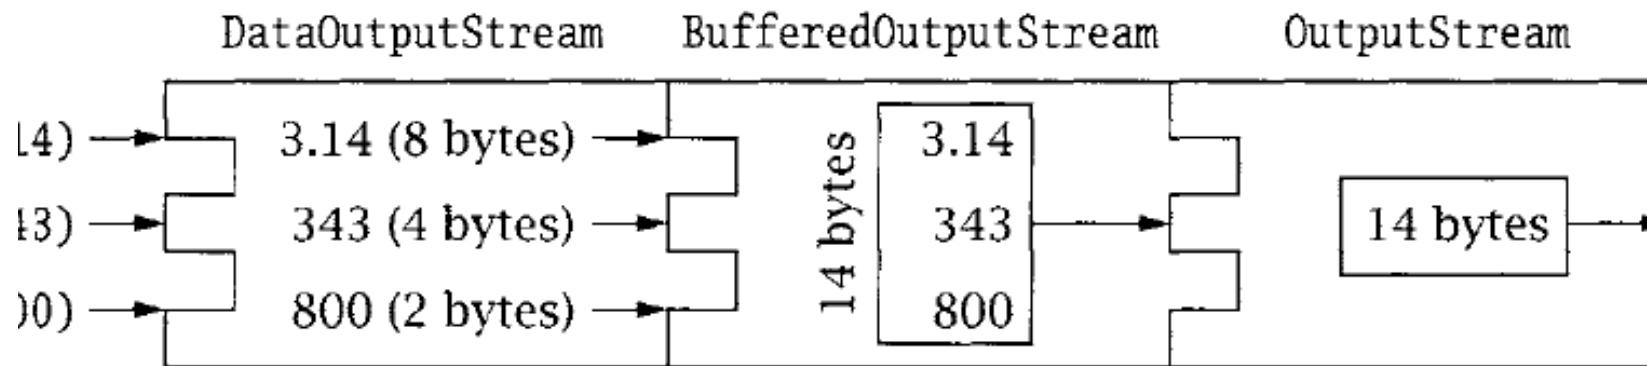
---

- Recap
- Network application programming
- FTP
- HTTP

---

# Backup Slides

# DataStream



# Demo

% wireshark to capture traffic  
tcp.srcport==6789 or tcp.dstport==6789