
Network Applications: High-performance Server Design

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

2/22/2016

Admin

- ❑ Assignment three will be posted tomorrow.
- ❑ Please start to think about projects
- ❑ Dates for exam 1?

Recap: Thread-Based Network Servers

- ❑ Why: threads (execution sequences) so that only one thread is blocked
- ❑ How:
 - Per-request thread
 - problem: large # of threads and their creations/deletions may let overhead grow out of control
 - Thread pool
 - Design 1: Service threads compete on the welcome socket
 - Design 2: Service threads and the main thread coordinate on the shared queue
 - polling (busy wait)
 - suspension: wait/notify

Recap: Program Correctness Analysis

❑ Safety

- No read/write; write/write conflicts
 - holding lock `Q` before reading or modifying shared data `Q` and `Q.wait_list`
- `Q.remove()` is not on an empty queue

❑ Liveness (progress)

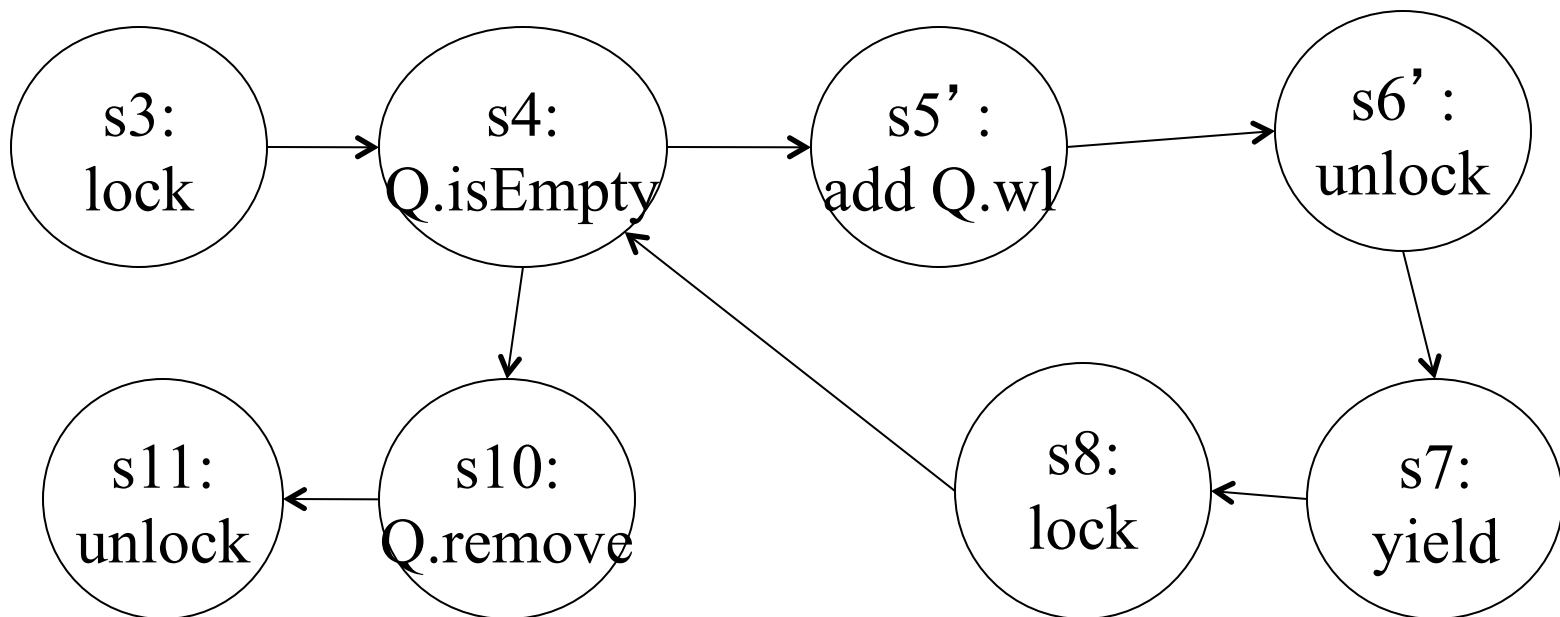
- main thread can always add to `Q`
- every connection in `Q` will be processed

❑ Fairness

- For example, in some settings, a designer may want the threads to share load equally

Main thread can always add to Q

- ❑ Assume main is blocked
- ❑ Suppose Q is not empty, then each iteration removes one element from Q
- ❑ In finite number of iterations, all elements in Q are removed and all service threads unlock and block



Each connection in Q is processed


- ❑ Cannot be guaranteed unless
 - there is fairness in the thread scheduler, or
 - put a limit on Q size to block the main thread

Blocking Queues in Java

- ❑ Design Pattern for producer/consumer pattern with blocking
- ❑ Two handy implementations
 - `LinkedBlockingQueue` (FIFO, may be bounded)
 - `ArrayBlockingQueue` (FIFO, bounded)
 - (plus a couple more)

<https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/BlockingQueue.html>

Summary: Using Threads



Advantages

- Intuitive (sequential) programming model
- Shared address space simplifies optimizations


Disadvantages

- Overhead: thread stacks, synchronization
- Thread pool parameter (how many threads) difficult to tune

Should You Use Threads?


- ❑ Typically avoid threads for io
 - Use event-driven, not threads, for GUIs, distributed systems, low-end servers.

- ❑ Use threads where true CPU concurrency is needed.
 - Where threads needed, isolate usage in threaded application kernel: keep most of code single-threaded.





(Extreme) Event-Driven Programming

- ❑ One execution stream: no CPU concurrency
- ❑ A single-thread event loop issues commands, waits for events, invokes handlers (callbacks)
 - Handlers issue asynchronous (non-blocking) I/O
 - No preemption of event handlers (handlers generally short-lived)



Event-Driven Programming

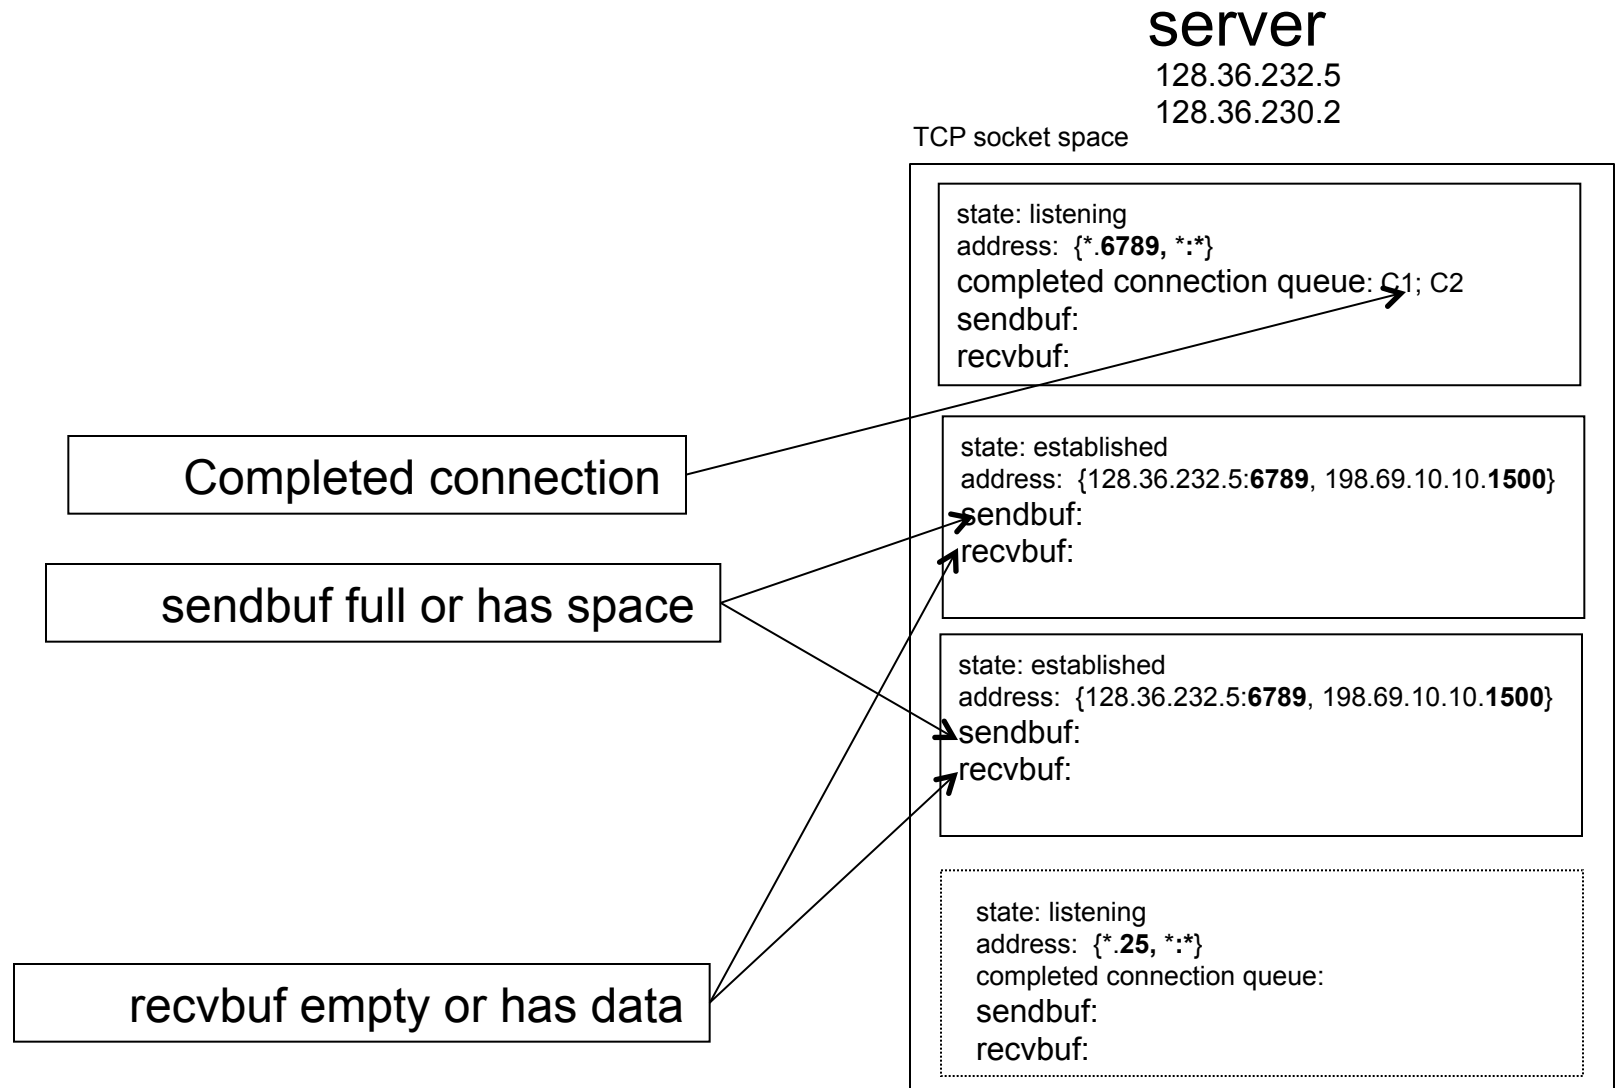


- ❑ Advantages
 - Single address space
 - No synchronization
- ❑ Disadvantages
 - Program complexity
 - In practice, disk reads/page fault still block
- ❑ Many examples: Click router, Flash web server, TP Monitors, NOX controller, Google Chrome (libevent), Dropbox (libevent),


Async Server I/O Basis

- ❑ Modern operating systems, such as Windows, Mac and Linux, provide facilities for fast, scalable IO based on the use of **asynchronous initiation** (e.g., aio_read) and **notifications of ready IO operations** taking place in the operating system layers.
 - Windows: IO Completion Ports
 - Linux: select, epoll (2.6)
 - Mac/FreeBSD: kqueue
- ❑ An Async IO package (e.g., Java Nio, Boost ASOI, Netty) aims to make (**some of**) these facilities available to applications

Async I/O Example: Ready Operations



Async IO with OS Notification



- ❑ Software framework on top of OS notification
 - Register handlers with dispatcher on sources (e.g., which sockets) and events (e.g., acceptable, readable, writable) to monitor
 - Dispatcher asks OS to check if any ready source/event
 - Dispatcher calls the registered handler of each ready event/source

Dispatcher Structure

```
//clients register interests/handlers on events/sources
while (true) {
    - ready events = select() /* or selectNow(), or
                               select(int timeout)
                               to check the
                               ready events from the
                               registered interest
                               events of sources */

    - foreach ready event {
        switch event type:
            accept: call accept handler
            readable: call read handler
            writable: call write handler
    }
}
```

Outline

- ❑ Admin and recap
- ❑ High performance servers
 - Thread
 - Per-request thread
 - Thread pool
 - Busy wait
 - Wait/notify
 - Asynchronous servers
 - Overview
 - Java async io

Async I/O in Java


- ❑ Java AIO provides some platform-independent abstractions on top of OS notification mechanisms (e.g., select/epoll)
- ❑ A typical network server (or package) builds on top of AIO abstractions
 - Sources (channel)
 - Selector
 - Buffers

Async I/O in Java: Sources

- ❑ A source that can generate events in Java is called a `SelectableChannel` object:
 - Example `SelectableChannels`:
`DatagramChannel`, `ServerSocketChannel`,
`SocketChannel`, `Pipe.SinkChannel`,
`Pipe.SourceChannel`
 - use `configureBlocking(false)` to make a channel non-blocking
- ❑ Note: Java `SelectableChannel` does not include file I/O

Async I/O in Java: Selector

- ❑ An important class is the class `Selector`, which is a base of the multiplexer/dispatcher
- ❑ Constructor of `Selector` is protected; create by invoking the `open` method to get a selector (why?)




Selector and Registration

- ❑ A selectable channel registers events to be monitored with a `selector` with the `register` method
- ❑ The registration returns an object called a `SelectionKey`:

```
SelectionKey key =  
    channel.register(selector, ops);
```

Java Async I/O Structure

- A `SelectionKey` object stores:
 - **interest set**: events to check:
`key.interestOps (ops)`
 - **ready set**: after calling `select`, it contains the events that are ready, e.g.
`key.isReadable ()`
 - **an attachment** that you can store anything you want
`key.attach (myObj)`



Checking Events

- ❑ A program calls `select` (or `selectNow()`, or `select(int timeout)`) to check for ready events from the registered `SelectableChannels`
 - Ready events are called the selected key set

```
selector.select();  
Set readyKeys = selector.selectedKeys();
```
- ❑ The program iterates over the selected key set to process all ready events


Dispatcher Structure

```
while (true) {  
    - selector.select()  
    - Set readyKeys = selector.selectedKeys();  
  
    - foreach key in readyKeys {  
        switch event type of key:  
            accept: call accept handler  
            readable: call read handler  
            writable: call write handler  
    }  
}
```

Async I/O in Java: ByteBuffer


- ❑ Java SelectableChannels typically use ByteBuffer for read and write
 - `channel.read(byteBuffer);`
 - `channel.write(byteBuffer);`
- ❑ ByteBuffer is a powerful class that can be used for both read and write
- ❑ It is derived from the class Buffer
- ❑ You can use it either as an absolute indexed or relatively indexed buffer

Buffer (relative index)




- ❑ Each Buffer has three numbers: position, limit, and capacity
 - Invariant: $0 \leq \text{position} \leq \text{limit} \leq \text{capacity}$
- ❑ `Buffer.clear()`: `position = 0; limit=capacity`

channel.read(Buffer)




- ❑ Put data into Buffer, starting at position, not to reach limit

channel.write(Buffer)




- ❑ Move data from Buffer to channel, starting at `position`, not to reach `limit`

Buffer.flip()



- ❑ `Buffer.flip()`: `limit=position; position=0`
- ❑ Why flip: used to switch from preparing data to output, e.g.,
 - `buf.put(header); // add header data to buf`
 - `in.read(buf); // read in data and add to buf`
 - `buf.flip(); // prepare for write`
 - `out.write(buf);`

Buffer.compact()



- ❑ Move [position , limit) to 0
- ❑ Set position to limit-position, limit to capacity

```
buf.clear(); // Prepare buffer for use
for (;;) {
    if (in.read(buf) < 0 && !buf.hasRemaining())
        break; // No more bytes to transfer
    buf.flip();
    out.write(buf);
    buf.compact(); // In case of partial write
}
```

Example

- ❑ See AsyncEchoServer/v1-2/
EchoServer.java

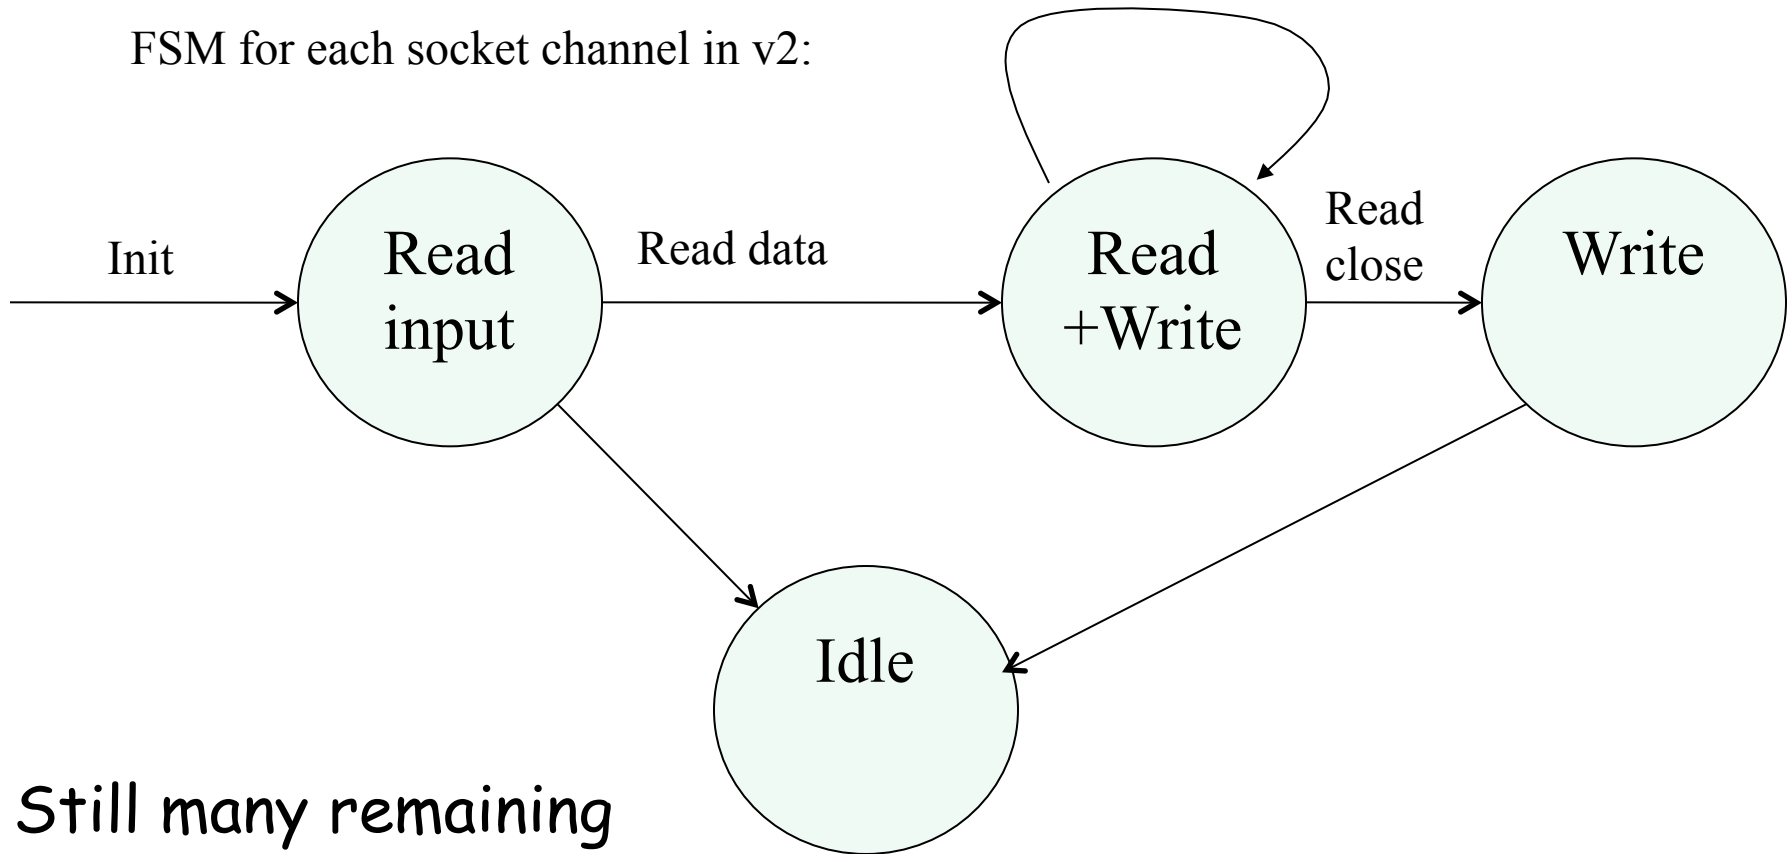
Problems of Async Echo Server v1

- ❑ Empty write: Callback to `handleWrite()` is unnecessary when nothing to write
 - Imagine empty write with 10,000 sockets
 - Solution: initially read only, later allow write

- ❑ `handleRead()` still reads after the client closes
 - ❑ Solution: after reading end of stream (read returns -1), deregister read interest for the channel

(Partial) Finite State Machine (FSM)

FSM for each socket channel in v2:



Still many remaining
issues such as idle
instead of close

Finite-State Machine and Thread

- Why no need to introduce FSM for a thread version?

