

---

Network Applications:  
High-performance Server Design:  
Async Servers/Operational Analysis

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

2/24/2016

# Admin

---

- ❑ Assignment three posted.
- ❑ Decisions
  - Projects or exam 2
  - Date for exam 1

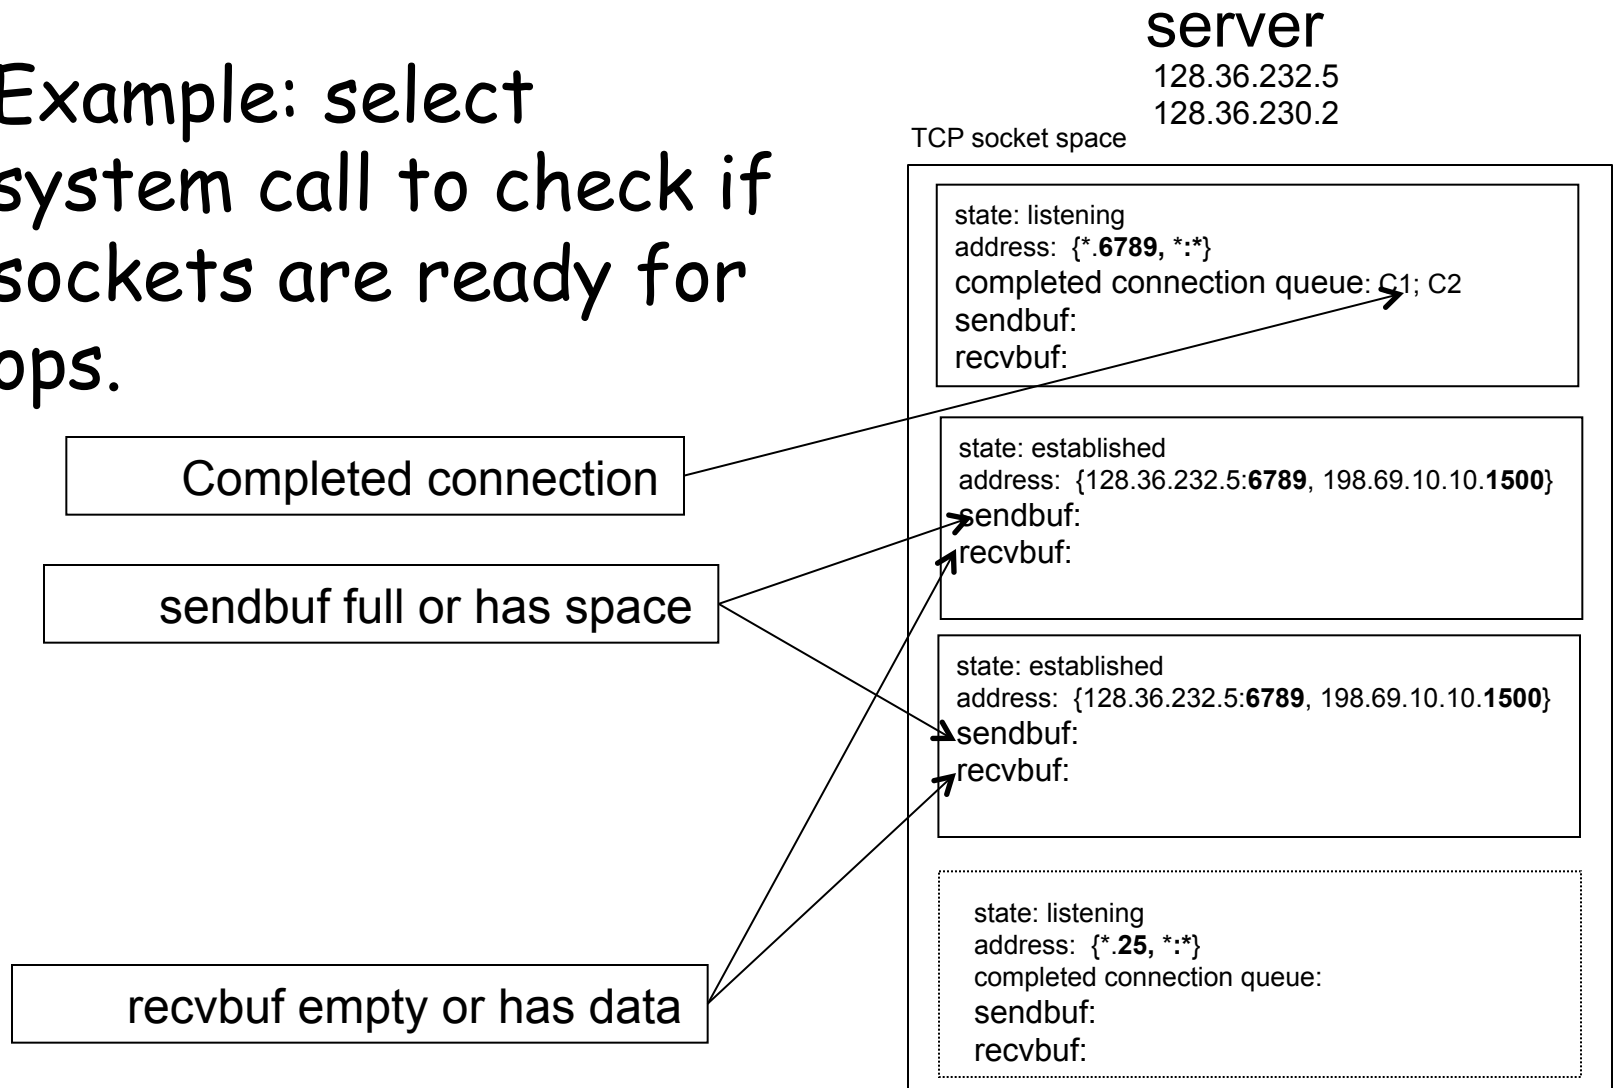
# Recap: Async Network Server

---

- ❑ Basic ideas: non-blocking operations
  1. **peek** system state (using a **select** mechanism) to issue only **ready** operations
  2. **asynchronous initiation** (e.g., `aio_read`) and completion **notification (callback)**


# Recap: Async Network Server using Select

- Example: select system call to check if sockets are ready for ops.



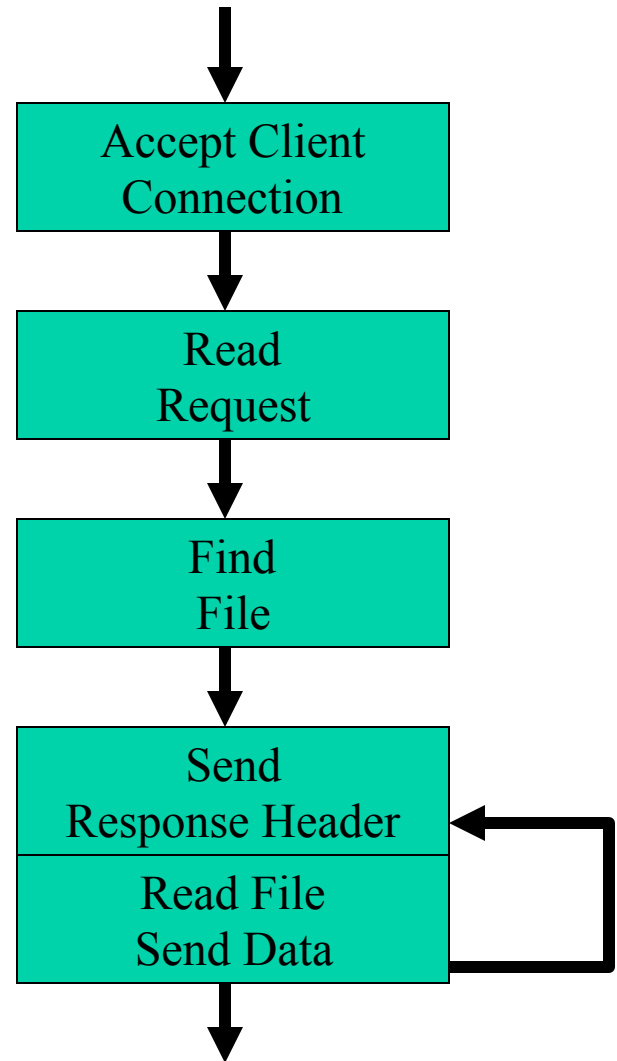

# Recap: Async Network Server using Select

- ❑ A event loop issues commands, waits for events, invokes handlers (callbacks)

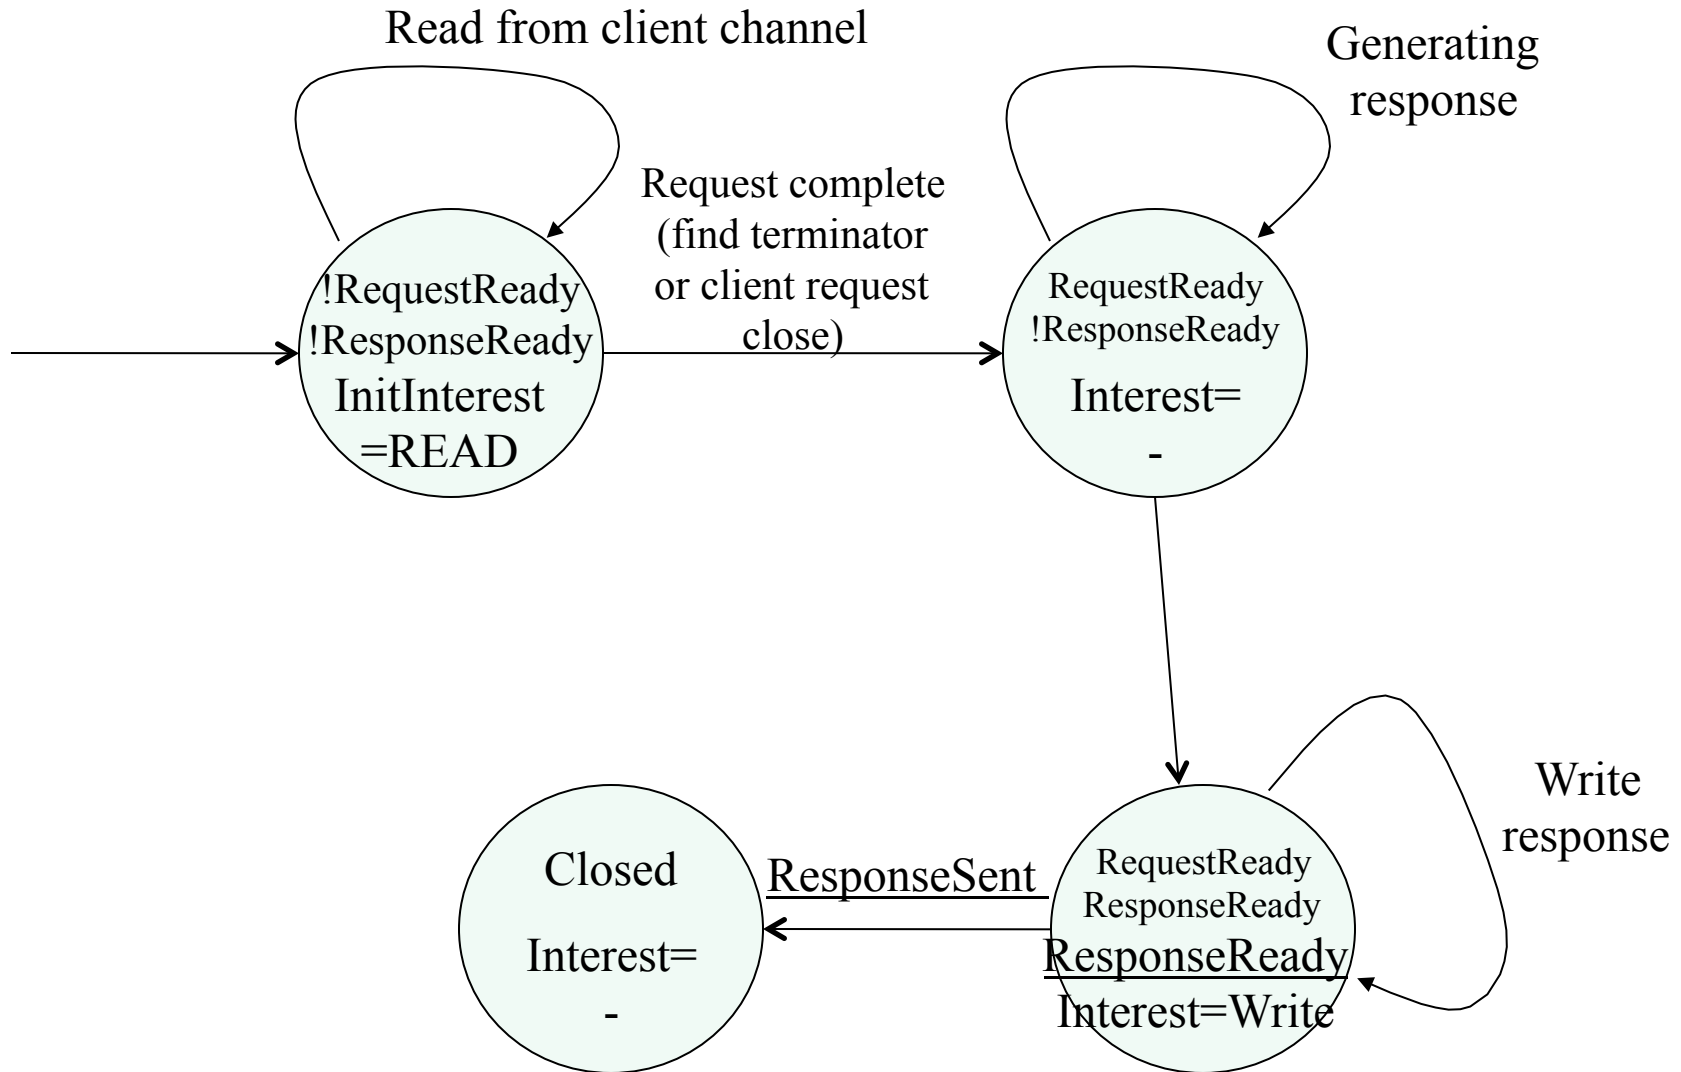


```
// clients register interests/  
handlers on events/sources  
while (true) {  
    - ready events = select()  
      /* or selectNow(),  
        or select(int timeout) */  
  
    - foreach ready event {  
        switch event type:  
        accept: call accept handler  
        readable: call read handler  
        writable: call write handler  
    }  
}
```

# Recap: Need to Manage Finite State Machine



# Another Finite State Machine



# Finite State Machine Design

- ❑ EchoServerV2:
  - Mixed read and write
- ❑ Example last slide: staged
  - First read request and then write response
- ❑ Choice depends on protocol and tolerance of complexity, e.g.,
  - HTTP/1.0 channel may use staged
  - HTTP/1.1/2/Chat channel may use mixed



# Toward More General Server Framework

- ❑ Our example EchoServer is for a specific protocol
- ❑ A general async/io programming framework tries to introduce structure to allow substantial reuse
  - Async io programming framework is among the more complex software systems
  - We will see one simple example, using EchoServer as a basis

# A More Extensible Dispatcher Design

- ❑ Fixed accept/read/write functions are not general design
- ❑ Requirement: map from key to handler
- ❑ A solution: Using attachment of each channel
  - Attaching a `ByteBuffer` to each channel is a narrow design for simple echo servers
  - A more general design can use the attachment to store a callback that indicates not only data (state) but also the handler (function)

# A More Extensible Dispatcher Design

## □ Attachment stores generic event handler

### ○ Define interfaces

- IAcceptHandler and
- IReadWriteHandler

### ○ Retrieve handlers at run time

```
if (key.isAcceptable()) { // a new connection is ready
    IAcceptHandler aH = (IAcceptHandler) key.attachment();
    aH.handleAccept(key);
}

if (key.isReadable() || key.isWritable()) {
    IReadWriteHandler rwH = (IReadWriteHandler) key.attachment();
    if (key.isReadable()) rwH.handleRead(key);
    if (key.isWritable()) rwH.handleWrite(key);
}
```

# Handler Design: Acceptor

---

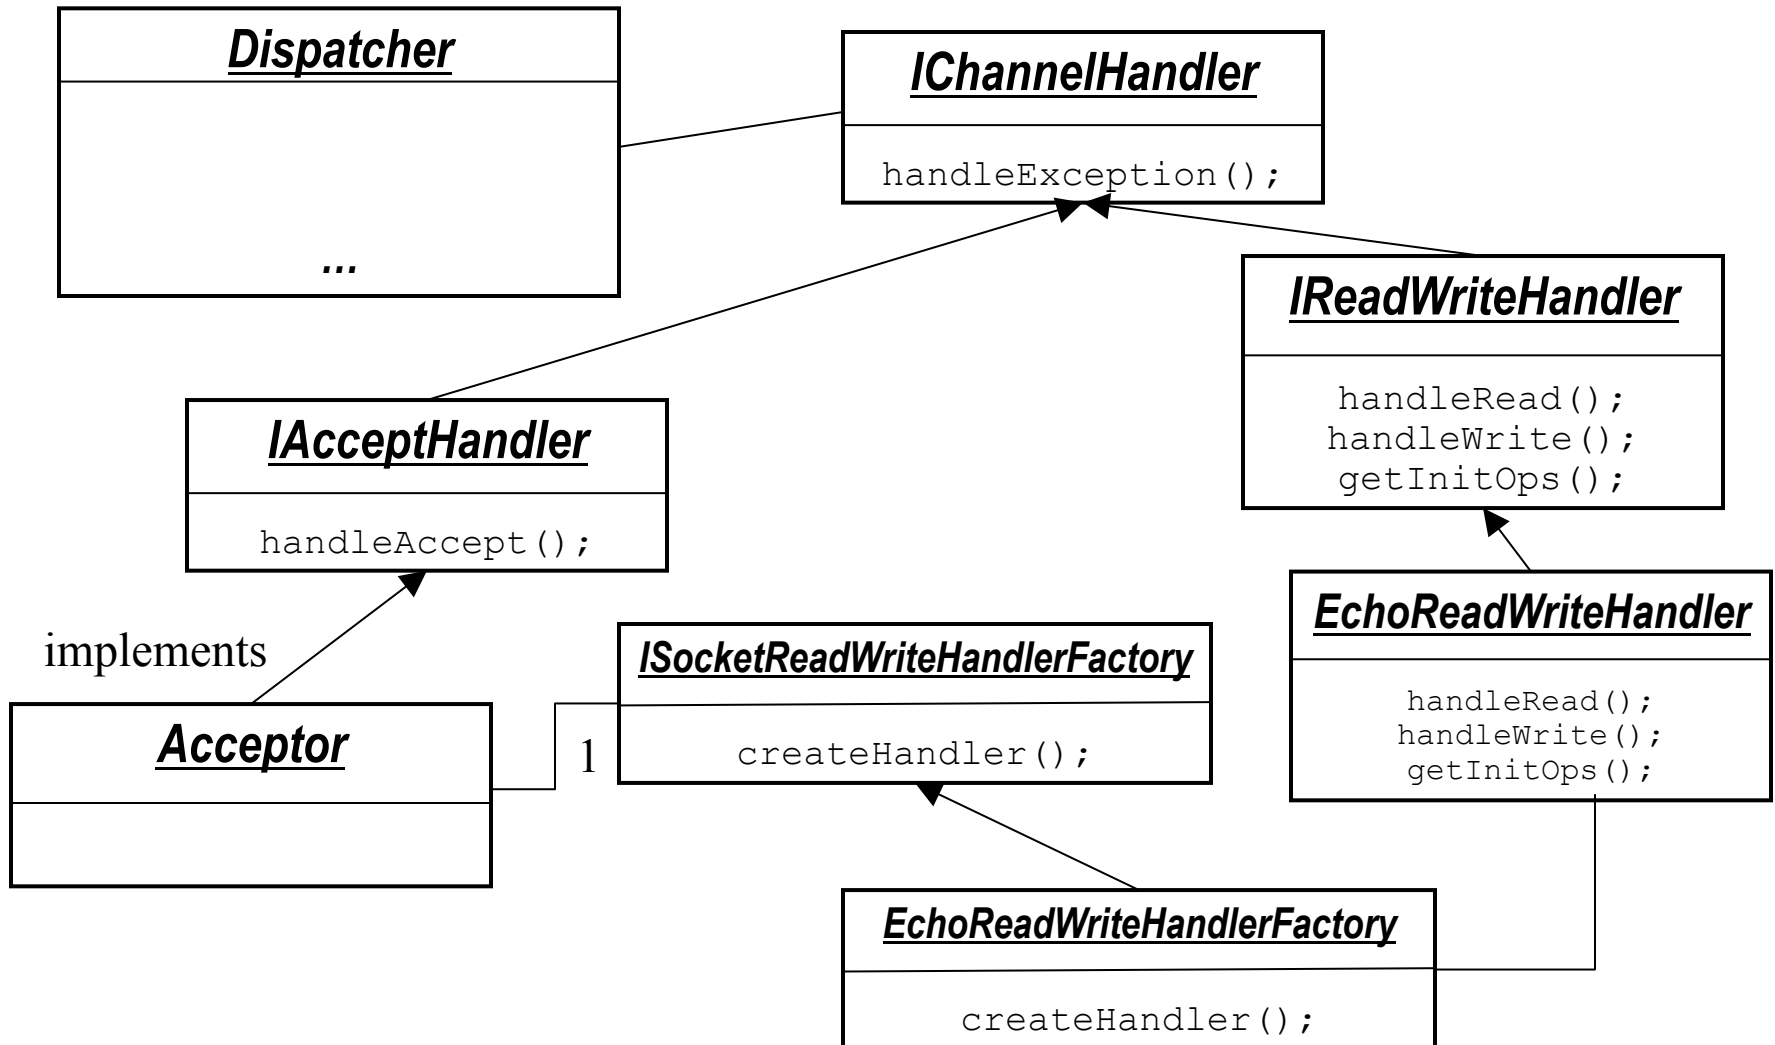
- ❑ What should an accept handler object know?
  - ServerSocketChannel (so that it can call accept)
    - Can be derived from SelectionKey in the call back
  - Selector (so that it can register new connections)
    - Can be derived from SelectionKey in the call back
  - What ReadWrite object to create (different protocols may use different ones)?
    - Pass a Factory object: SocketReadWriteHandlerFactory

# Handler Design: ReadWriteHandler

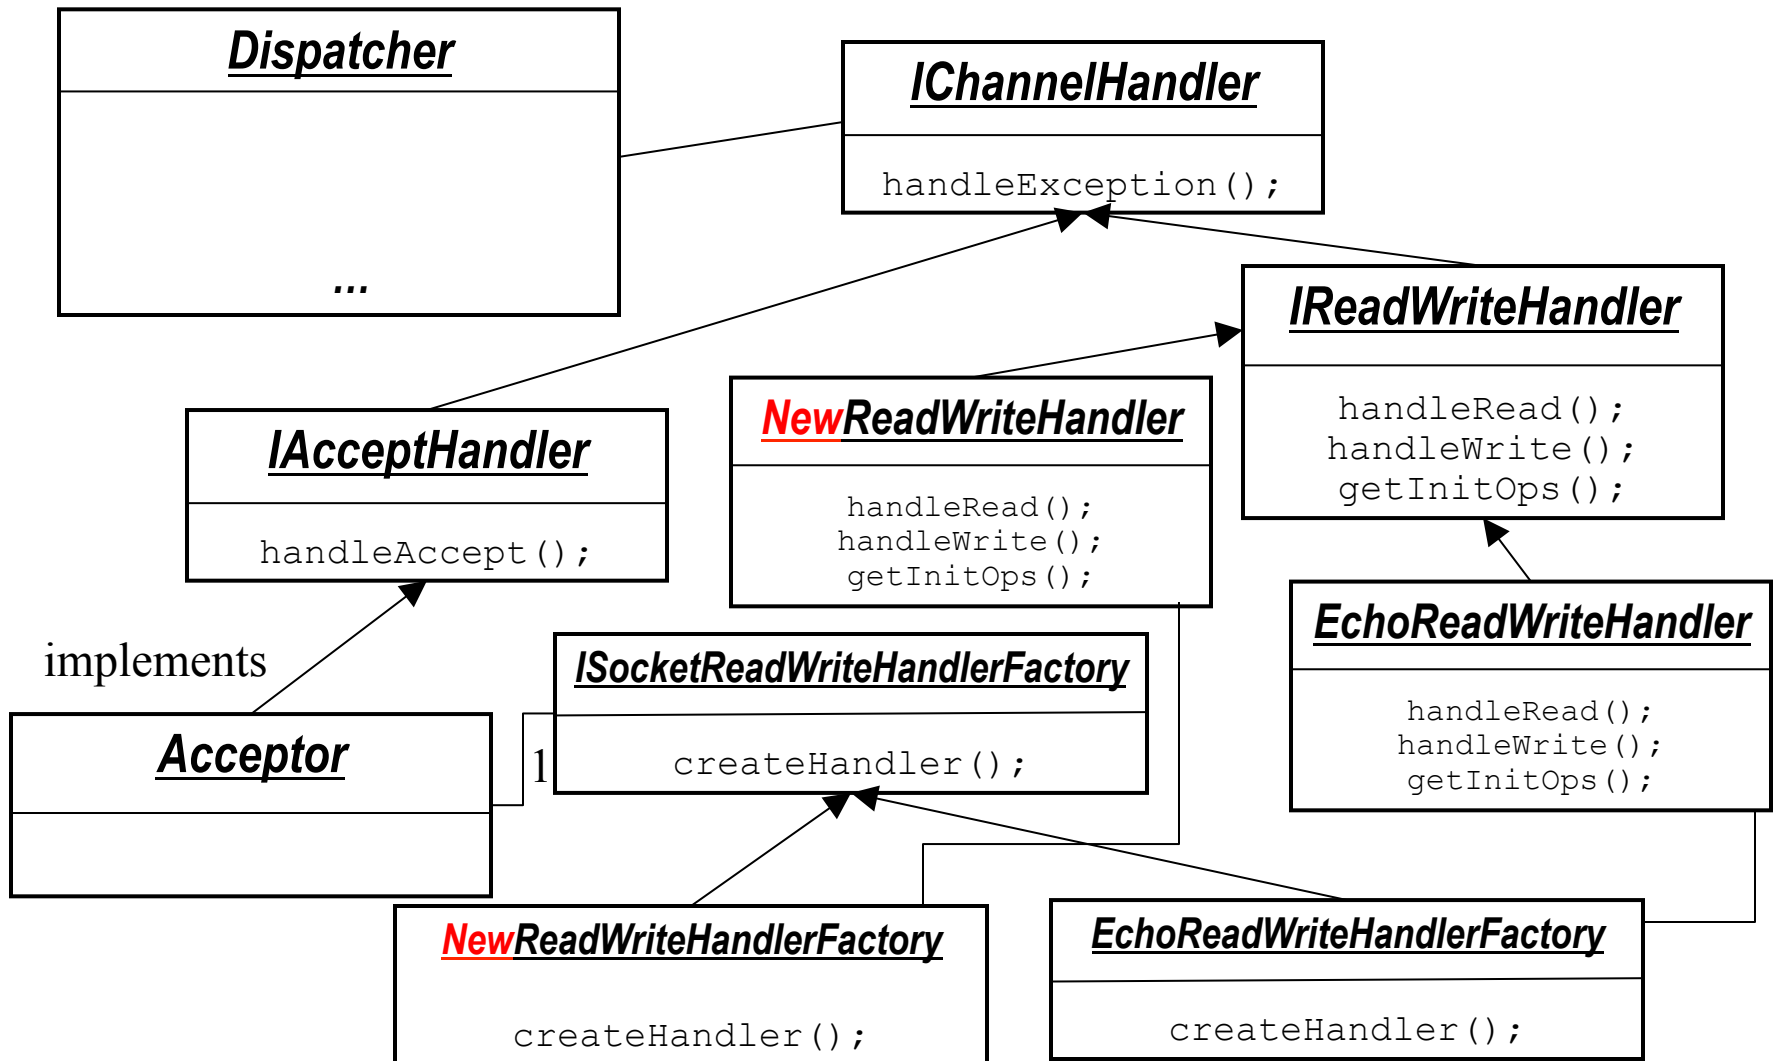
---

- ❑ What should a ReadWrite handler object know?
  - SocketChannel (so that it can read/write data)
    - Can be derived from SelectionKey in the call back
  - Selector (so that it can change state)
    - Can be derived from SelectionKey in the call back

# Class Diagram of SimpleNAIO



# Class Diagram of SimpleNAIO



# SimpleNAIO

---

□ See `AsyncEchoServer/v3/*.java`



# Discussion on SimpleNAIO

## ❑ In our current implementation (Server.java)

1. Create dispatcher
2. Create server socket channel and listener
3. Register server socket channel to dispatcher
4. Start dispatcher thread

Can we switch 3 and 4?

# Extending SimpleNAIO

- ❑ A production network server often closes a connection if it does not receive a complete request in TIMEOUT
  - ❑ One way to implement time out is that
    - the read handler registers a timeout event with a timeout watcher thread with a call back
    - the watcher thread invokes the call back upon TIMEOUT
    - the callback closes the connection
- Any problem?

# Extending Dispatcher Interface

- ❑ Interacting from another thread to the dispatcher thread can be tricky
- ❑ Typical solution: async command queue

```
while (true) {  
    - process async. command queue  
    - ready events = select (or selectNow(), or  
      select(int timeout)) to check for ready events  
      from the registered interest events of  
      SelectableChannels  
  
    - foreach ready event  
      call handler  
}
```

# Question

- ❑ How may you implement the async command queue to the selector thread?

```
public void invokeLater(Runnable run) {  
    synchronized (pendingInvocations) {  
        pendingInvocations.add(run);  
    }  
    selector.wakeup();  
}
```

# Question

---

- ❑ What if another thread wants to wait until a command is finished by the dispatcher thread?

```

public void invokeAndWait(final Runnable task)
    throws InterruptedException
{
    if (Thread.currentThread() == selectorThread) {
        // We are in the selector's thread. No need to schedule
        // execution
        task.run();
    } else {
        // Used to deliver the notification that the task is executed
        final Object latch = new Object();
        synchronized (latch) {
            // Uses the invokeLater method with a newly created task
            this.invokeLater(new Runnable() {
                public void run() {
                    task.run();
                }
                // Notifies
                synchronized(latch) { latch.notify(); }
            });
            // Wait for the task to complete.
            latch.wait();
        }
        // Ok, we are done, the task was executed. Proceed.
    }
}

```

# Recap: Async Network Server

---

- ❑ Basic idea: non-blocking operations
  1. **peek** system state (select) to issue only **ready** operations
  2. **asynchronous initiation** (e.g., aio\_read) and completion **notification (callback)**

# Alternative Design: Asynchronous Channel using Future/Listener

- ❑ Java 7 introduces  
    `AsynchronousServerSocketChannel` and  
    `AsynchronousSocketChannel` beyond  
    `ServerSocketChannel` and `SocketChannel`
  - `accept`, `connect`, `read`, `write` return `Futures` or have a callback. Selectors are not used

<https://docs.oracle.com/javase/7/docs/api/java/nio/channels/AsynchronousServerSocketChannel.html>

<https://docs.oracle.com/javase/7/docs/api/java/nio/channels/AsynchronousSocketChannel.html>



```

SocketAddress address
    = new InetSocketAddress(args[0], port);
AsynchronousSocketChannel client
    = AsynchronousSocketChannel.open();
Future<Void> connected
    = client.connect(address);

ByteBuffer buffer = ByteBuffer.allocate(100);

// wait for the connection to finish
connected.get();

// read from the connection
Future<Integer> future = client.read(buffer);

// do other things...

// wait for the read to finish...
future.get();

// flip and drain the buffer
buffer.flip();
WritableByteChannel out
    = Channels.newChannel(System.out);
out.write(buffer);

```

```

class LineHandler implements
CompletionHandler<Integer, ByteBuffer> {

    @Override
    public void completed(Integer result, ByteBuffer buffer)
    {
        buffer.flip();
        WritableByteChannel out
            = Channels.newChannel(System.out);
        try {
            out.write(buffer);
        } catch (IOException ex) {
            System.err.println(ex);
        }
    }

    @Override
    public void failed(Throwable ex,
                        ByteBuffer attachment) {
        System.err.println(ex.getMessage());
    }
}

```

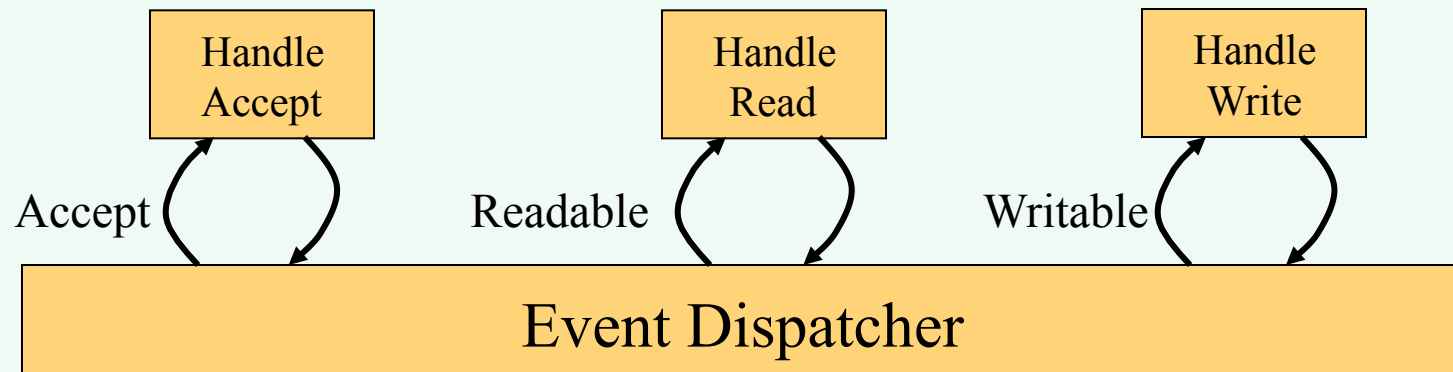
```

ByteBuffer buffer = ByteBuffer.allocate(100);
CompletionHandler<Integer, ByteBuffer>
    handler = new LineHandler();
channel.read(buffer, buffer, handler);

```

# Extending FSM

- ❑ In addition to management threads, a system may still need multiple threads for performance (why?)
  - FSM code can never block, but page faults, file io, garbage collection may still force blocking
  - CPU may become the bottleneck and there maybe multiple cores supporting multiple threads (typically 2 n threads)



# Summary: Architecture

---

- ❑ Architectures
  - Multi threads
  - Asynchronous
  - Hybrid
  
- ❑ Assigned reading: SEDA

# Problems of Event-Driven Server

---


- ❑ Obscure control flow for programmers and tools
- ❑ Difficult to engineer, modularize, and tune
- ❑ Difficult for performance/failure isolation between FSMs

# Another view

- Events obscure control flow
  - For programmers *and* tools

| <i>Threads</i>  | <i>Events</i>   |
|---|---|
| <pre>thread_main(int sock) {     struct session s;     accept_conn(sock, &amp;s);     read_request(&amp;s);     pin_cache(&amp;s);     write_response(&amp;s);     unpin(&amp;s); }  pin_cache(struct session *s) {     pin(&amp;s);     if( !in_cache(&amp;s) )         read_file(&amp;s); }</pre> | <pre>AcceptHandler(event e) {     struct session *s = new_session(e);     RequestHandler.enqueue(s); }  RequestHandler(struct session *s) {     ...; CacheHandler.enqueue(s); }  CacheHandler(struct session *s) {     pin(s);     if( !in_cache(s) ) ReadFileHandler.enqueue(s);     else                ResponseHandler.enqueue(s); }  ...  ExitHandler(struct session *s) {     ...; unpin(&amp;s); free_session(s); }</pre> |

## *Web Server*



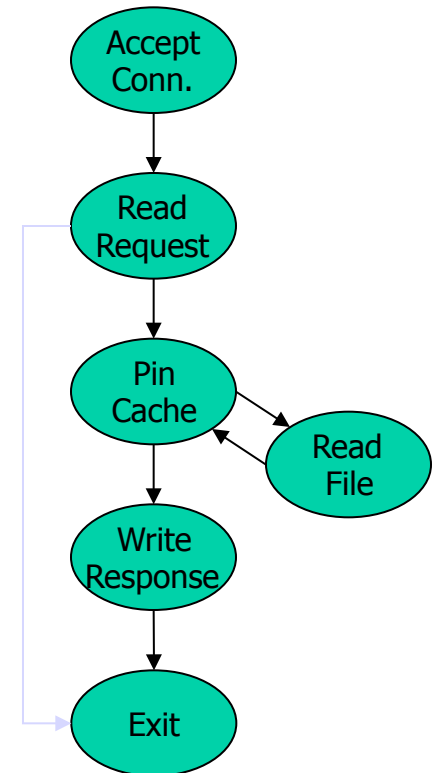
[von Behren]

# State Management

- ❑ Events require manual state management
- ❑ Hard to know when to free
  - Use GC or risk bugs

| <i>Threads</i>   | <i>Events</i>  |
|--|--|
| <pre>thread_main(int sock) {<br/>    struct session s;<br/>    accept_conn(sock, &amp;s);<br/>    if( !read_request(&amp;s) )<br/>        return;<br/>    pin_cache(&amp;s);<br/>    write_response(&amp;s);<br/>    unpin(&amp;s);<br/>}</pre><br><pre>pin_cache(struct session *s) {<br/>    pin(&amp;s);<br/>    if( !in_cache(&amp;s) )<br/>        read_file(&amp;s);<br/>}</pre> | <pre>CacheHandler(struct session *s) {<br/>    pin(s);<br/>    if( !in_cache(s) ) ReadFileHandler.enqueue(s);<br/>    else<br/>        ResponseHandler.enqueue(s);<br/>}<br/><br/>RequestHandler(struct session *s) {<br/>    ...; if( error ) return; CacheHandler.enqueue(s);<br/>}<br/>...<br/>ExitHandler(struct session *s) {<br/>    ...; unpin(&amp;s); free_session(s);<br/>}<br/><br/>AcceptHandler(event e) {<br/>    struct session *s = new_session(e);<br/>    RequestHandler.enqueue(s); }</pre> |

## *Web Server*



[von Behren]

# Summary: The High-Performance Network Servers Journey

- ❑ Avoid blocking (so that we can reach bottleneck throughput)
  - Introduce threads
- ❑ Limit unlimited thread overhead
  - Thread pool, async io
- ❑ Coordinating data access
  - synchronization (lock, synchronized)
- ❑ Coordinating behavior: avoid busy-wait
  - Wait/notify; select FSM, Future/Listener
- ❑ Extensibility/robustness
  - Language support/Design for interfaces



# Beyond Class: Design Patterns

---

- ❑ We have seen Java as an example
- ❑ C++ and C# can be quite similar. For C++ and general design patterns:
  - <http://www.cs.wustl.edu/~schmidt/PDF/OOCP-tutorial4.pdf>
  - <http://www.stal.de/Downloads/ADC2004/pr03.pdf>



# Some Questions

---

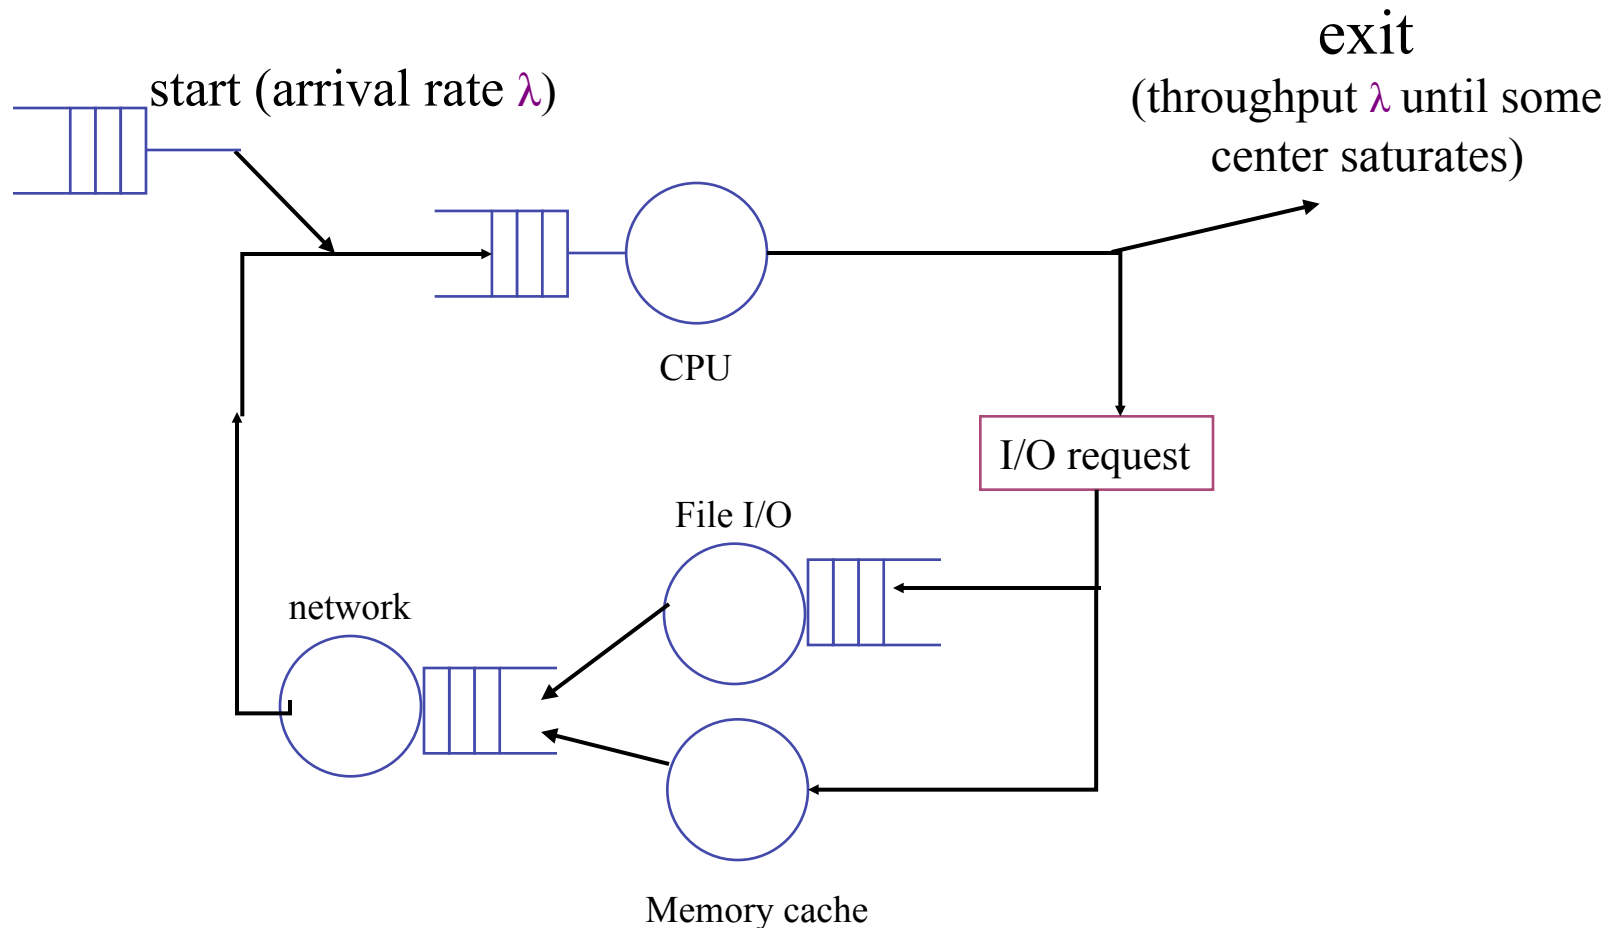
- ❑ When is CPU the bottleneck for scalability?
  - So that we need to add helpers
- ❑ How do we know that we are reaching the limit of scalability of a single machine?
- ❑ These questions drive network server architecture design

# Operational Analysis

---


- ❑ Relationships that do not require any assumptions about the distribution of service times or inter-arrival times.
- ❑ Identified originally by Buzen (1976) and later extended by Denning and Buzen (1978).
- ❑ We touch only some techniques/results
  - In particular, bottleneck analysis
- ❑ More details see linked reading

# Under the Hood (An example FSM)




# Operational Analysis: Resource Demand of a Request

CPU




$V_{\text{CPU}}$  visits for  $S_{\text{CPU}}$  units of resource time per visit

Network




$V_{\text{Net}}$  visits for  $S_{\text{Net}}$  units of resource time per visit

Disk



$V_{\text{Disk}}$  visits for  $S_{\text{Disk}}$  units of resource time per visit

Memory



$V_{\text{Mem}}$  visits for  $S_{\text{Mem}}$  units of resource time per visit

# Operational Quantities

- T: observation interval
- $B_i$ : busy time of device  $i$
- $i = 0$  denotes system

$A_i$ : # arrivals to device  $i$

$C_i$ : # completions at device  $i$

$$\text{arrival rate } \lambda_i = \frac{A_i}{T}$$

$$\text{Throughput } X_i = \frac{C_i}{T}$$

$$\text{Utilization } U_i = \frac{B_i}{T}$$

$$\text{Mean service time } S_i = \frac{B_i}{C_i}$$

# Utilization Law

$$\begin{aligned}\text{Utilization } U_i &= \frac{B_i}{T} \\ &= \frac{C_i}{T} \frac{B_i}{C_i} \\ &= X_i S_i\end{aligned}$$

- The law is independent of any assumption on arrival/service process
- Example: Suppose NIC processes 125 pkts/sec, and each pkt takes 2 ms. What is utilization of the network NIC?

# Deriving Relationship Between R, U, and S for one Device

- Assume flow balanced (arrival=throughput), Little's Law:

$$Q = \lambda R = X R$$

- Assume PASTA (Poisson arrival--memory-less arrival--sees time average), a new request sees Q ahead of it, and FIFO

$$R = S + Q S = S + X R S$$

- According to utilization law,  $U = X S$

$$R = S + U R \longrightarrow R = \frac{S}{1-U}$$