

---

Network Applications:  
Operational Analysis;  
Load Balancing among Multiple Servers

Y. Richard Yang

<http://zoo.cs.yale.edu/classes/cs433/>

2/29/2016

# Admin

---

- ❑ Assignment three status and questions.

# Recap: Operational Analysis

□ Objective: derive relationships among (measured) operational performance metrics

- T: observation interval
- $A_i$ : # arrivals to device  $i$
- $B_i$ : busy time of device  $i$
- $C_i$ : # completions at device  $i$
- $i = 0$  denotes system

$$\text{Arrival rate } \lambda_i = \frac{A_i}{T}$$

$$\text{Throughput } X_i = \frac{C_i}{T}$$

$$\text{Utilization } U_i = \frac{B_i}{T}$$

$$\text{Mean service time } S_i = \frac{B_i}{C_i}$$

$$\text{Utilization } U_i = X_i S_i$$

# Forced Flow Law

- Assume each request visits device  $i$   $V_i$  times

$$\begin{aligned}\text{Throughput } X_i &= \frac{C_i}{T} \\ &= \frac{C_i}{C_0} \frac{C_0}{T} \\ &= V_i X\end{aligned}$$

# Bottleneck Device

$$\begin{aligned}\text{Utilization } U_i &= X_i S_i \\ &= V_i X S_i \\ &= X V_i S_i\end{aligned}$$

- Define  $D_i = V_i S_i$  as the total demand of a request on device  $i$
- The device with the highest  $D_i$  has the highest utilization, and thus is called the **bottleneck**

# Bottleneck vs System Throughput

---

$$\text{Utilization } U_i = XV_iS_i \leq 1$$

$$\rightarrow X \leq \frac{1}{D_{\max}}$$

# Example 1

---

- ❑ A request may need
  - 10 ms CPU execution time
  - 1 Mbytes network bw
  - 1 Mbytes file access where
    - 50% hit in memory cache
- ❑ Suppose network bw is 100 Mbps, disk I/O rate is 1 ms per 8 Kbytes (assuming the program reads 8 KB each time)
- ❑ Where is the bottleneck?

# Example 1 (cont.)

## □ CPU:

- $D_{\text{CPU}} = 10 \text{ ms}$  (e.q. 100 requests/s)

## □ Network:

- $D_{\text{Net}} = 1 \text{ Mbytes} / 100 \text{ Mbps} = 80 \text{ ms}$  (e.q., 12.5 requests/s)

## □ Disk I/O:

- $D_{\text{disk}} = 0.5 * 1 \text{ ms} * 1\text{M}/8\text{K} = 62.5 \text{ ms}$   
(e.q. = 16 requests/s)



# Example 2

- ❑ A request may need
  - 150 ms CPU execution time (e.g., **dynamic content**)
  - 1 Mbytes network bw
  - 1 Mbytes file access where
    - 50% hit in memory cache
- ❑ Suppose network bw is 100 Mbps, disk I/O rate is 1 ms per 8 Kbytes (assuming the program reads 8 KB each time)
- ❑ Bottleneck: CPU -> use multiple threads to use more CPUs, if available, to avoid CPU as bottleneck

# Interactive Response Time Law

## □ System setup

- Closed system with  $N$  users
- Each user sends in a request, after response, think time, and then sends next request

## ○ Notation

- $Z$  = user think-time,  $R$  = Response time

- The total cycle time of a user request is  $R+Z$

In duration  $T$ , #requests generated by each user:  $T/(R+Z)$  requests

# Interactive Response Time Law

□ *If N users and flow balanced:*

System Throughput  $X = \text{Total \# req.} / T$

$$= \frac{N \frac{T}{R+Z}}{T}$$

$$= \frac{N}{R+Z}$$

$$R = \frac{N}{X} - Z$$

# Bottleneck Analysis

---

$$X(N) \leq \min \left\{ \frac{1}{D_{\max}}, \frac{N}{D+Z} \right\}$$

$$R(N) \geq \max \{ D, ND_{\max} - Z \}$$

□ Here  $D$  is the sum of  $D_i$

## Proof

$$X(N) \leq \min \left\{ \frac{1}{D_{\max}}, \frac{N}{D+Z} \right\}$$

$$R(N) \geq \max \{ D, ND_{\max} - Z \}$$

□ We know

$$X \leq \frac{1}{D_{\max}} \quad R(N) \geq D$$

Using interactive response time law:

$$R = \frac{N}{X} - Z \quad \longrightarrow \quad R \geq ND_{\max} - Z$$

$$X = \frac{N}{R+Z} \quad \longrightarrow \quad X \leq \frac{N}{D+Z}$$

# Summary: Operational Laws


- ❑ Utilization law:  $U = XS$
- ❑ Forced flow law:  $X_i = V_i X$
- ❑ Bottleneck device: largest  $D_i = V_i S_i$
- ❑ Little's Law:  $Q_i = X_i R_i$
- ❑ Bottleneck bound of interactive response (for the given closed model):

$$X(N) \leq \min \left\{ \frac{1}{D_{\max}}, \frac{N}{D+Z} \right\}$$

$$R(N) \geq \max \{ D, ND_{\max} - Z \}$$

# In Practice: Common Bottlenecks

- ❑ No more file descriptors
- ❑ Sockets stuck in `TIME_WAIT`
- ❑ High memory use (swapping)
- ❑ CPU overload
- ❑ Interrupt (IRQ) overload



[Aaron Bannert]

# YouTube Design Alg.

---

```
while (true)
{
    identify_and_fix_bottlenecks();
    drink();
    sleep();
    notice_new_bottleneck();
}
```



# Summary: High-Perf. Network Server

- ❑ Avoid blocking (so that we can reach bottleneck throughput)
  - Introduce threads
- ❑ Limit unlimited thread overhead
  - Thread pool, async io
- ❑ Shared variables
  - Synchronization (lock, synchronized)
- ❑ Avoid busy-wait
  - Wait/notify; FSM
- ❑ Extensibility/robustness
  - Language support/Design for interfaces
- ❑ System modeling and measurements
  - Queueing analysis, operational analysis

# Outline

---

- ❑ Recap
- ❑ Single network server
- ❑ Multiple network servers
  - Why multiple network servers

# Why Multiple Servers?

---

## □ Scalability

### ○ Scaling beyond single server capability

- There is a fundamental limit on what a single server can
  - process (CPU/bw/disk throughput)
  - store (disk/memory)

### ○ Scaling beyond geographical location capability

- There is a limit on the speed of light
- Network detour and delay further increase the delay

# Why Multiple Servers?

---

- ❑ Redundancy and fault tolerance
  - Administration/Maintenance (e.g., incremental upgrade)
  - Redundancy (e.g., to handle failures)

# Why Multiple Servers?

---

- ❑ System/software architecture
  - Resources may be naturally distributed at different machines (e.g., run a single copy of a database server due to single license; access to resource from third party)
  - Security (e.g., front end, business logic, and database)

# Discussion: Key Technical Challenges in Using Multiple Servers

---


# Outline

---

- ❑ Recap
- ❑ Single network server
- ❑ Multiple network servers
  - Why multiple servers
  - Request routing mechanisms

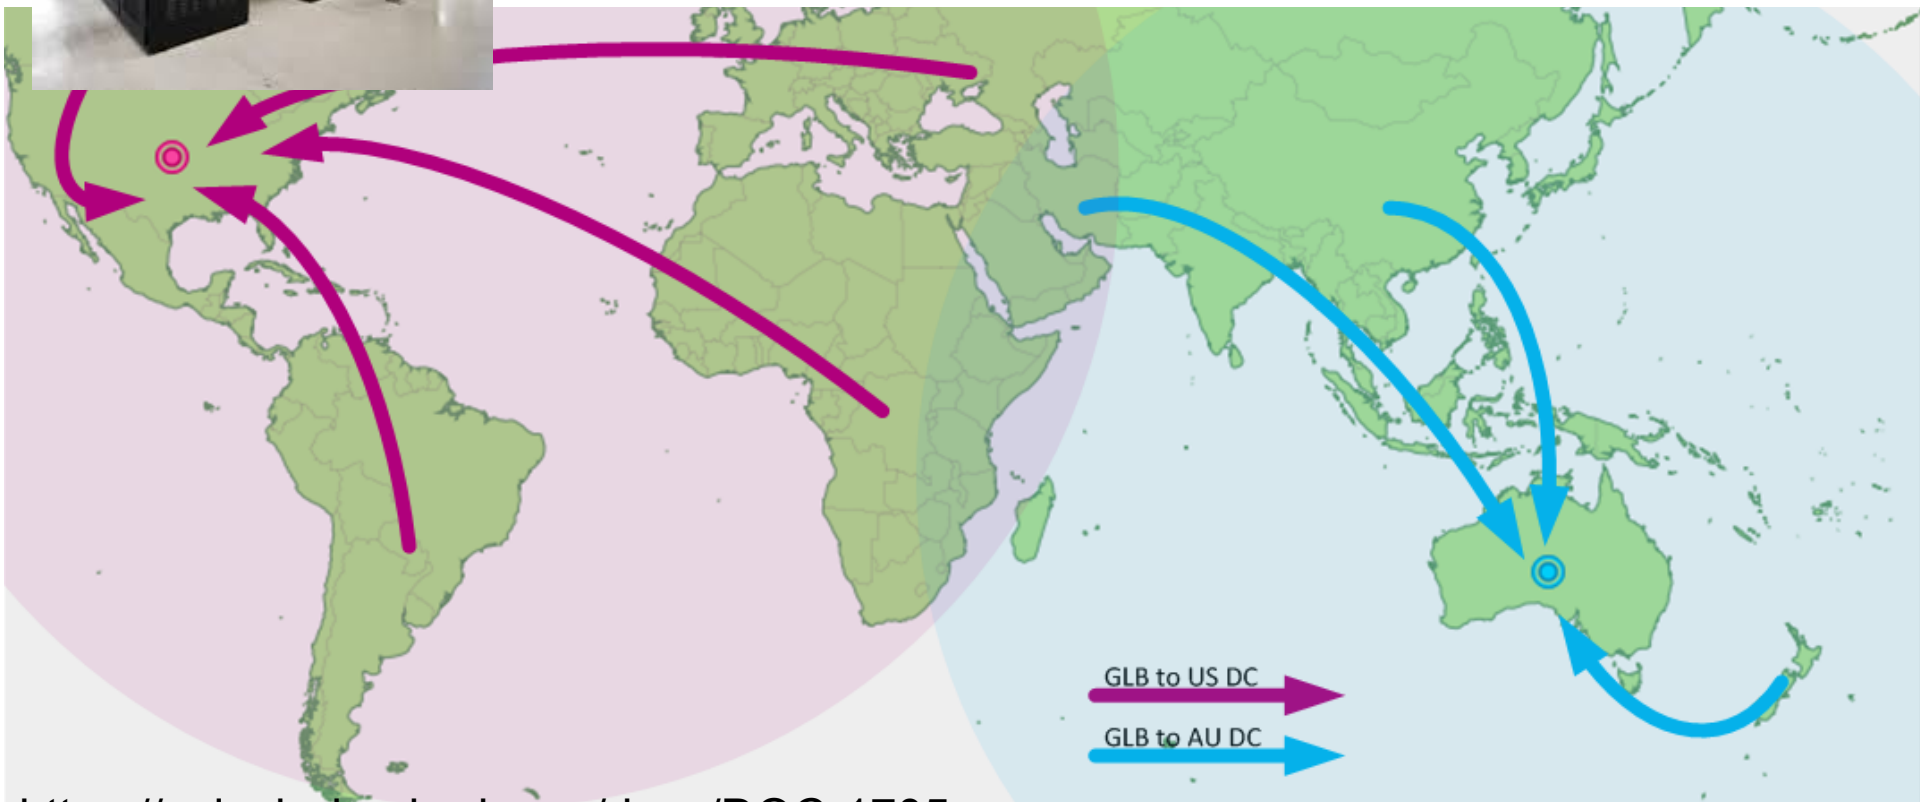
# Request Routing: Overview

- Global request routing: select a server site for each request
- Local request routing: select a specific server at the chosen site






# Request Routing: Overview



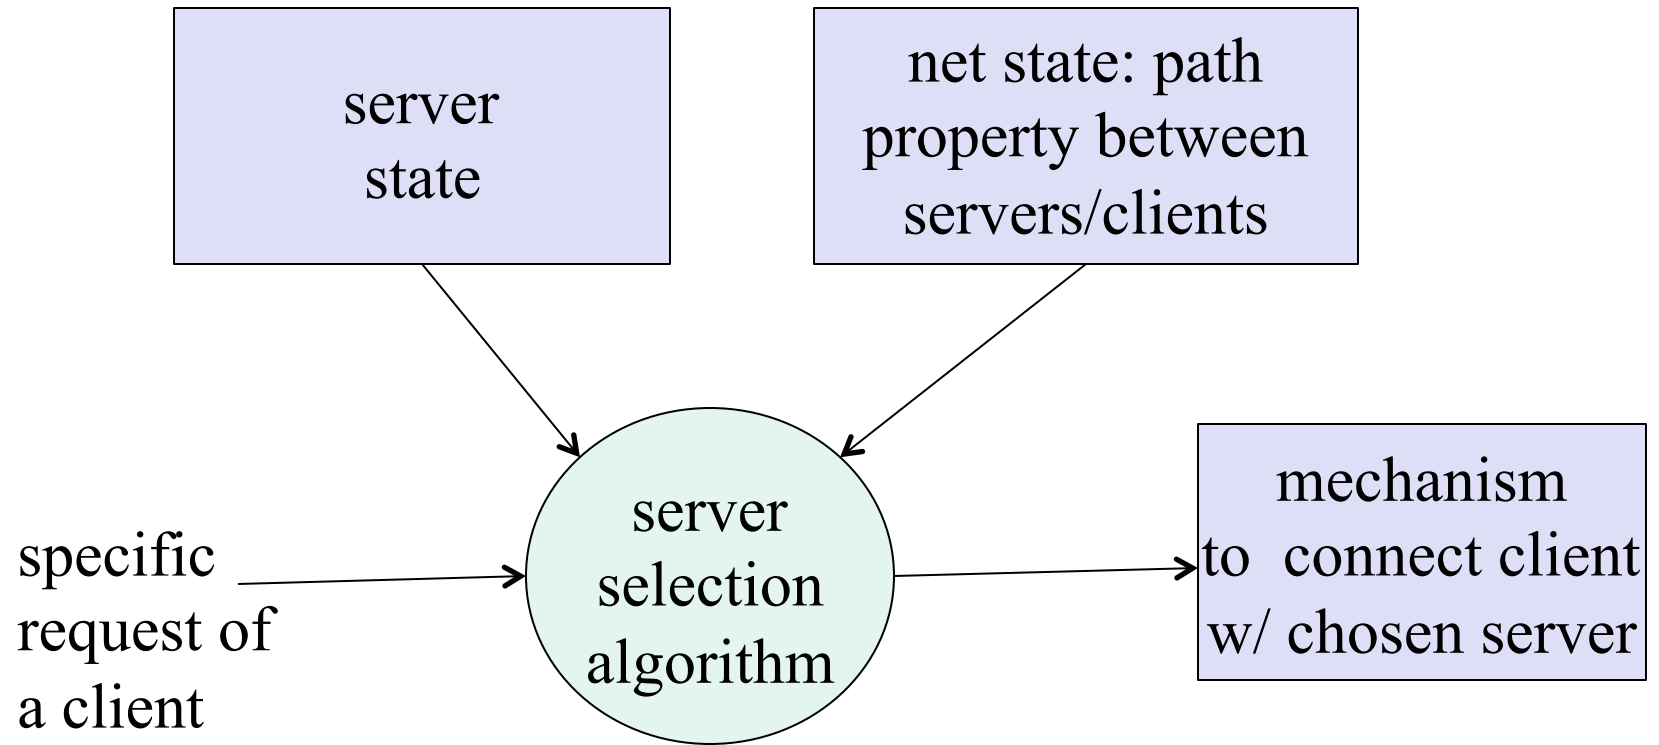
# Request Routing: Basic Architecture

## □ Major components


- Server state monitoring
  - Load (incl. failed or not); what requests it can serve
- Network path properties between clients and servers
  - E.g., bw, delay, loss, network cost
- Server selection alg.
- Request direction mechanism



# Request Routing: Basic Arch




# Network Path Properties



## □ Why is the problem difficult?

- Scalability: if do measurements, complete measurements grow with  $N * M$ , where
  - $N$  is # of clients
  - $M$  is # of servers




- Complexity/feasibility in computing path metrics

# Network Path Properties: Improve Scalability

## □ Aggregation:

- merge a set of IP addresses (reduce  $N$  and  $M$ )
  - E.g., when computing path properties, Akamai aggregates all clients sharing the same local DNS server




## □ Sampling and prediction


- Instead of measuring  $N \times M$  entries, we measure a subset and **predict** the **unmeasured** paths
- We will cover it later in the course

# Server Selection

- Why is the problem difficult?
  - What are potential problems of just sending each new client to the lightest load server?

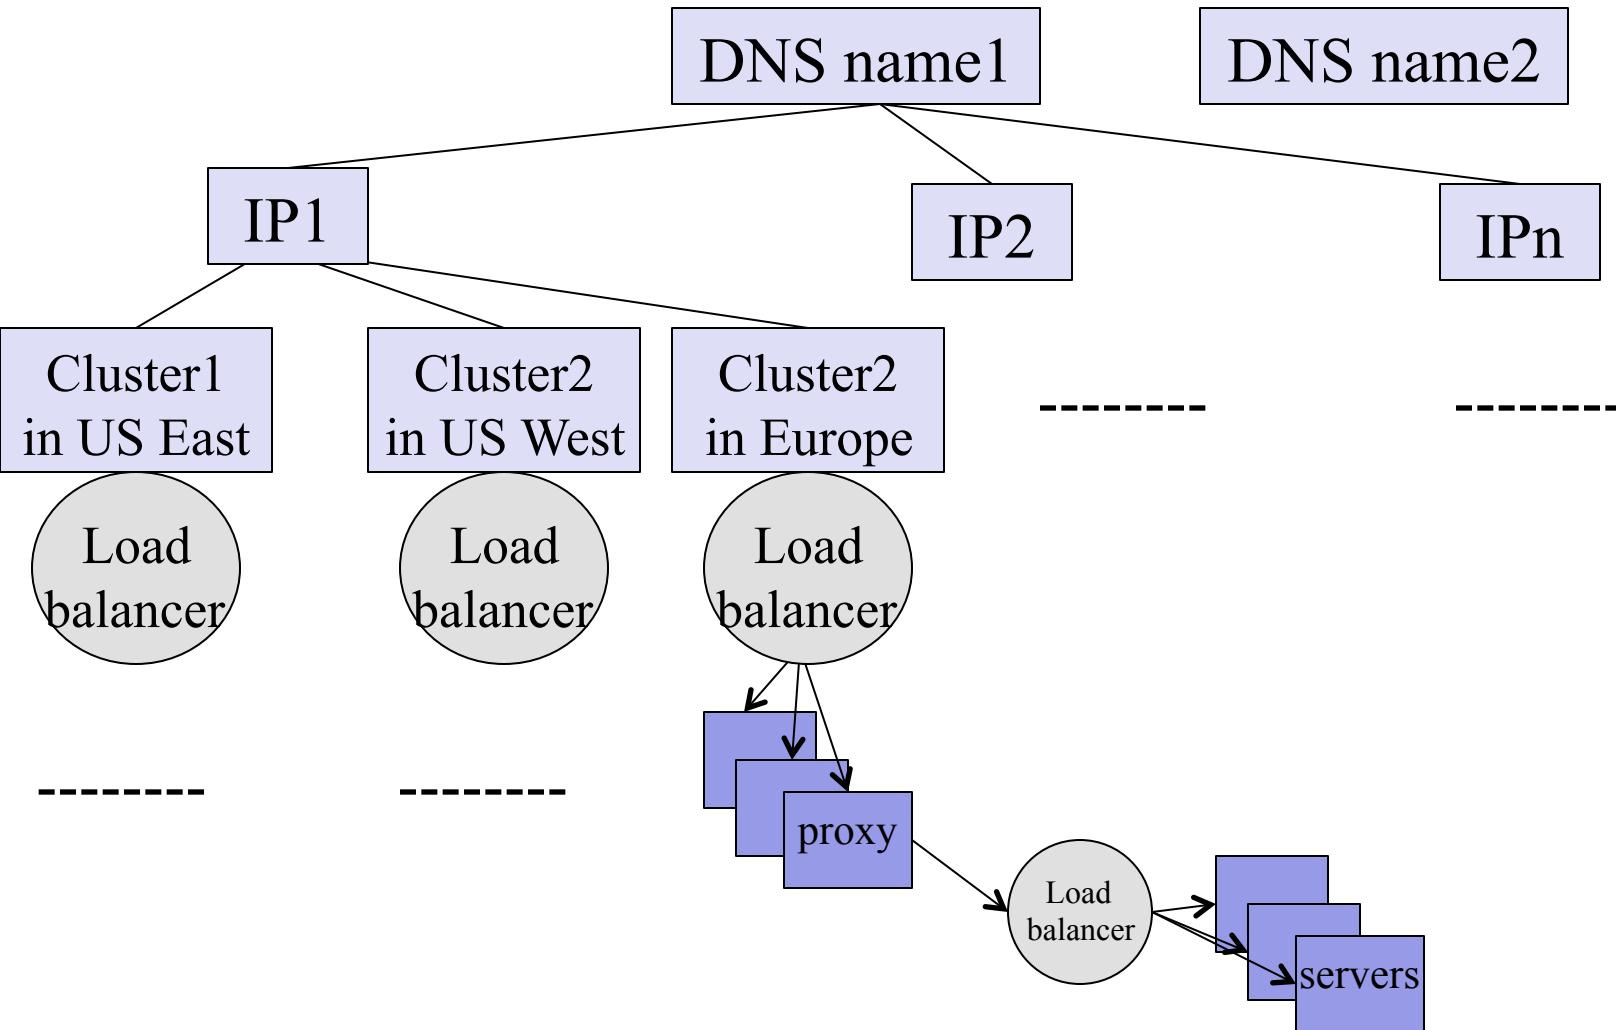


# Client Direction Mechanisms



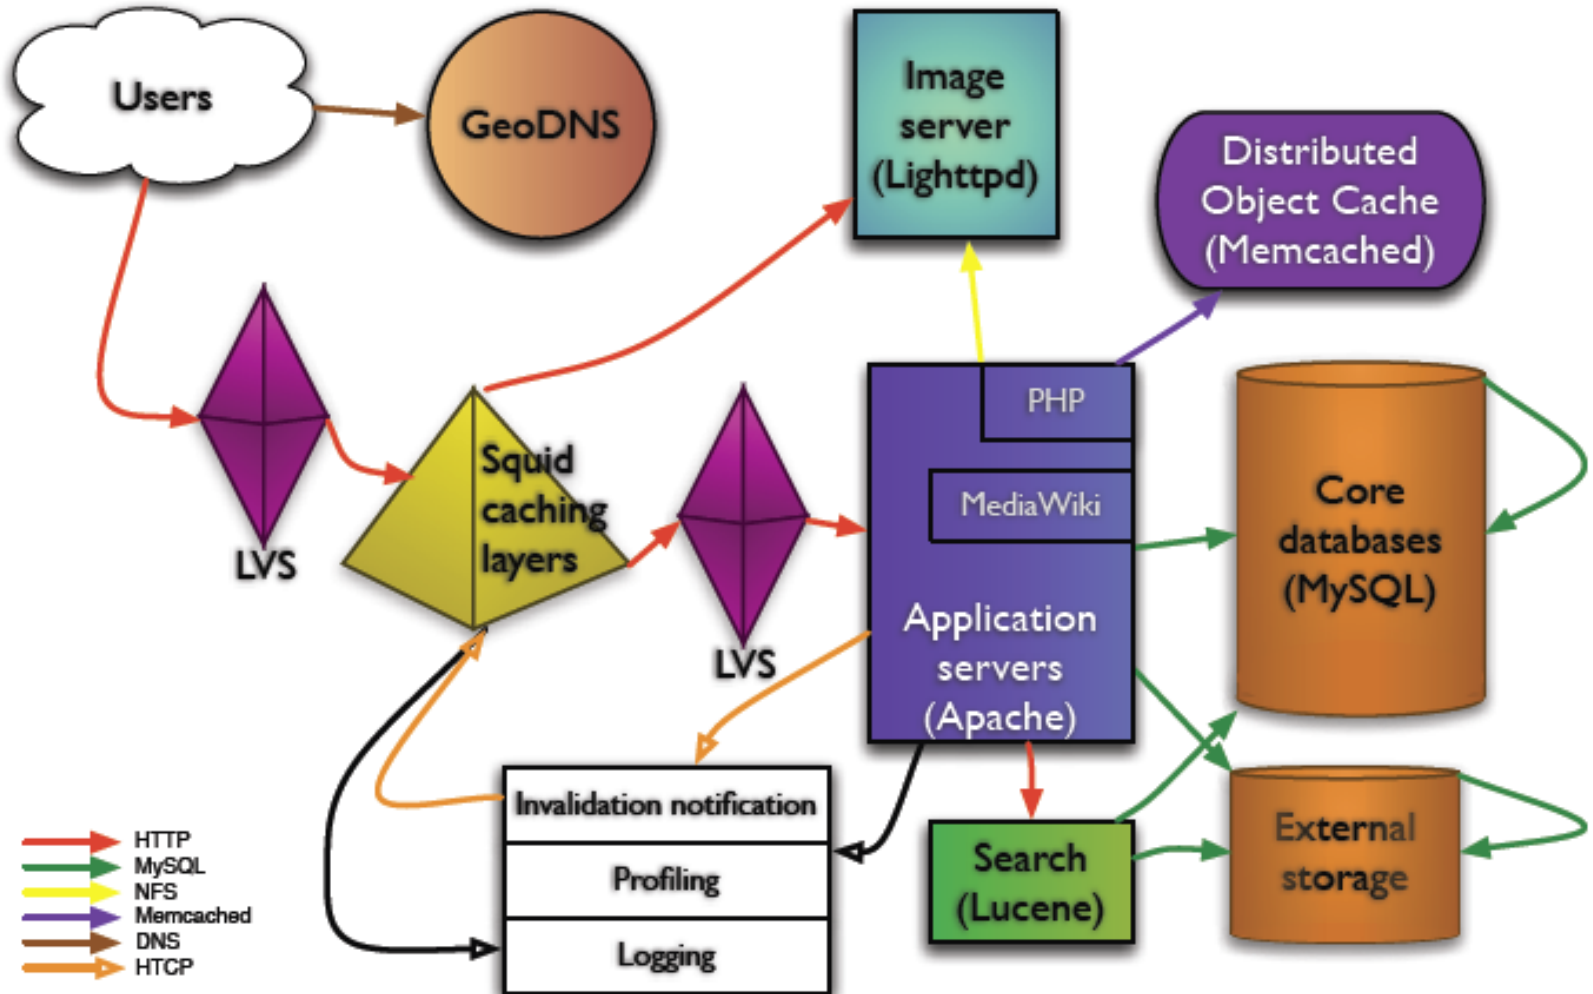
- ❑ Application layer
  - App/user is given a list of candidate server names
  - HTTP redirector
- ❑ DNS: name resolution gives a list of server addresses
- ❑ IP layer: Same IP address represents multiple physical servers
  - IP **anycast**: Same IP address shared by multiple servers and announced at different parts of the Internet. Network directs different clients to different servers
  - Smart-switch indirection: a server IP address may be a **virtual IP** address for a cluster of physical servers

# Direction Mechanisms are Often Combined





# Example: Wikipedia Architecture




# Outline

---

- ❑ Recap
- ❑ Single network server
- ❑ Multiple network servers
  - Why multiple servers
  - Request routing mechanisms
    - Overview
    - Application-layer

## Example: Netflix

Hostname	Organization
www.netflix.com	Netflix
signup.netflix.com	Amazon
movies.netflix.com	Amazon
agmoviecontrol.netflix.com	Amazon
nflx.i.87f50a04.x.lcdn.nflximg.com	Level 3
netflix-753.vo.llnwd.net	Limelight
netflix753.as.nflximg.com.edgesuite.net	Akamai



# Example: Netflix Manifest File

- Client player authenticates and then downloads manifest file from servers at Amazon Cloud

```
<nccp:cdns>
  <nccp:cdn>
    <nccp:name>level3</nccp:name>
    <nccp:cdnid>6</nccp:cdnid>
    <nccp:rank>1</nccp:rank>
    <nccp:weight>140</nccp:weight>
  </nccp:cdn>
  <nccp:cdn>
    <nccp:name>limelight</nccp:name>
    <nccp:cdnid>4</nccp:cdnid>
    <nccp:rank>2</nccp:rank>
    <nccp:weight>120</nccp:weight>
  </nccp:cdn>
  <nccp:cdn>
    <nccp:name>akamai</nccp:name>
    <nccp:cdnid>9</nccp:cdnid>
    <nccp:rank>3</nccp:rank>
    <nccp:weight>100</nccp:weight>
  </nccp:cdn>
</nccp:cdns>
```

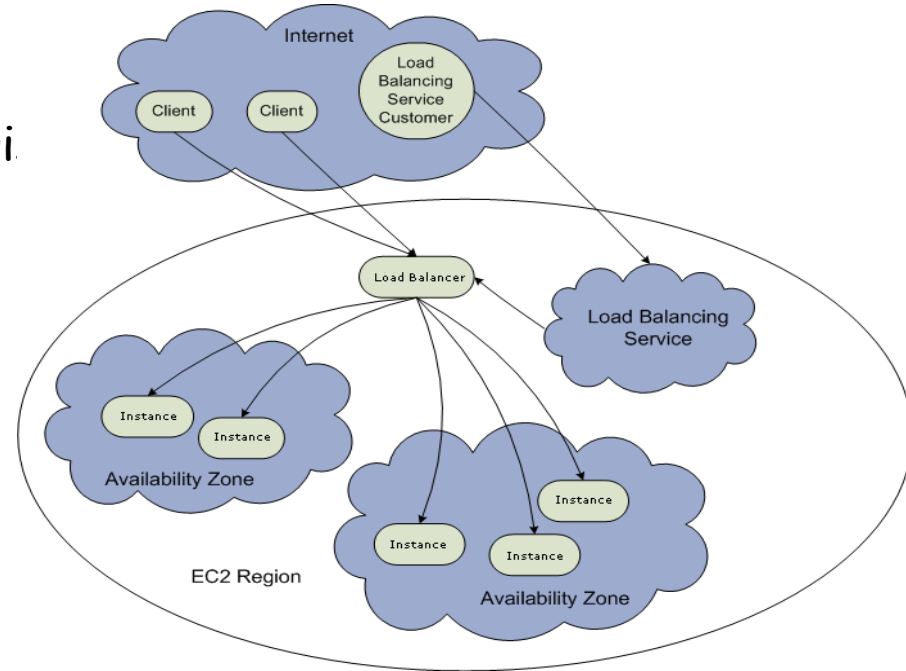
# Example: Netflix Manifest File

```
<nccp:bitrate>560</nccp:bitrate>
<nccp:videoprofile>
  playready-h264mpl30-dash
</nccp:videoprofile>
<nccp:resolution>
  <nccp:width>512</nccp:width>
  <nccp:height>384</nccp:height>
</nccp:resolution>
<nccp:pixelaspect>
  <nccp:width>4</nccp:width>
  <nccp:height>3</nccp:height>
</nccp:pixelaspect>v
<nccp:downloadurls>
  <nccp:downloadurl>
    <nccp:expiration>131xxx</nccp:expiration>
    <nccp:cdnid>6</nccp:cdnid>
    <nccp:url>http://nflx.i.../...</nccp:url>
  </nccp:downloadurl>
  <nccp:downloadurl>
    <nccp:expiration>131xxx</nccp:expiration>
    <nccp:cdnid>4</nccp:cdnid>
    <nccp:url>http://netflix.../...</nccp:url>
  </nccp:downloadurl>
  <nccp:downloadurl>
    <nccp:expiration>131xxx</nccp:expiration>
    <nccp:cdnid>9</nccp:cdnid>
    <nccp:url>http://netflix.../...</nccp:url>
  </nccp:downloadurl>
</nccp:downloadurls>
```

# Example: Amazon Elastic Cloud 2 (EC2)

## Elastic Load Balancing

- ❑ Use the *create-load-balancer* command to create an Elastic Load Balancer.
- ❑ Use the *register-instances-with-load-balancer* command to register Amazon EC2 instances that you want to load balance with the Elastic Load Balancer.
- ❑ Elastic Load Balancing automatically checks the health of your load balancing Amazon EC2 instances. You can optionally customize the health checks by using the *configure-healthcheck* command.
- ❑ Traffic to the DNS name provided by the Elastic Load Balancer is automatically distributed across healthy Amazon EC2 instances.



# Details: Create Load Balancer

The operation returns the DNS name of your LoadBalancer.  
You can then map that to any other domain name (such as  
`www.mywebsite.com`)  
(how?)

```
%aws elb create-load-balancer --load-  
balancer-name my-load-balancer --listeners  
"Protocol=HTTP,LoadBalancerPort=80,InstanceP  
rotocol=HTTP,InstancePort=80" --  
availability-zones us-west-2a us-west-2b
```

Result:

```
{ "DNSName": "my-load-balancer-123456789.us-west-2.elb.amazonaws.com" }
```

# Details: Configure Health Check

The operation configures how instances are monitored, e.g.,

```
%aws elb configure-health-check --load-balancer-name my-load-balancer --health-check Target=HTTP:80/png,Interval=30,UnhealthyThreshold=2,HealthyThreshold=2,Timeout=3
```

Result:

```
{
  "HealthCheck": {
    "HealthyThreshold": 2,
    "Interval": 30,
    "Target": "HTTP:80/png",
    "Timeout": 3,
    "UnhealthyThreshold": 2
  }
}
```



# Details: Register Instances

The operation registers instances that can receive traffic,

```
%aws elb register-instances-with-load-balancer --load-balancer-name my-load-balancer --instances i-d6f6fae3
```

Result:

```
{  "Instances": [  
    {"InstanceId": "i-d6f6fae3"},  
    {"InstanceId": "i-207d9717"},  
    {"InstanceId": "i-afefb49b"}  
  ]  
}
```

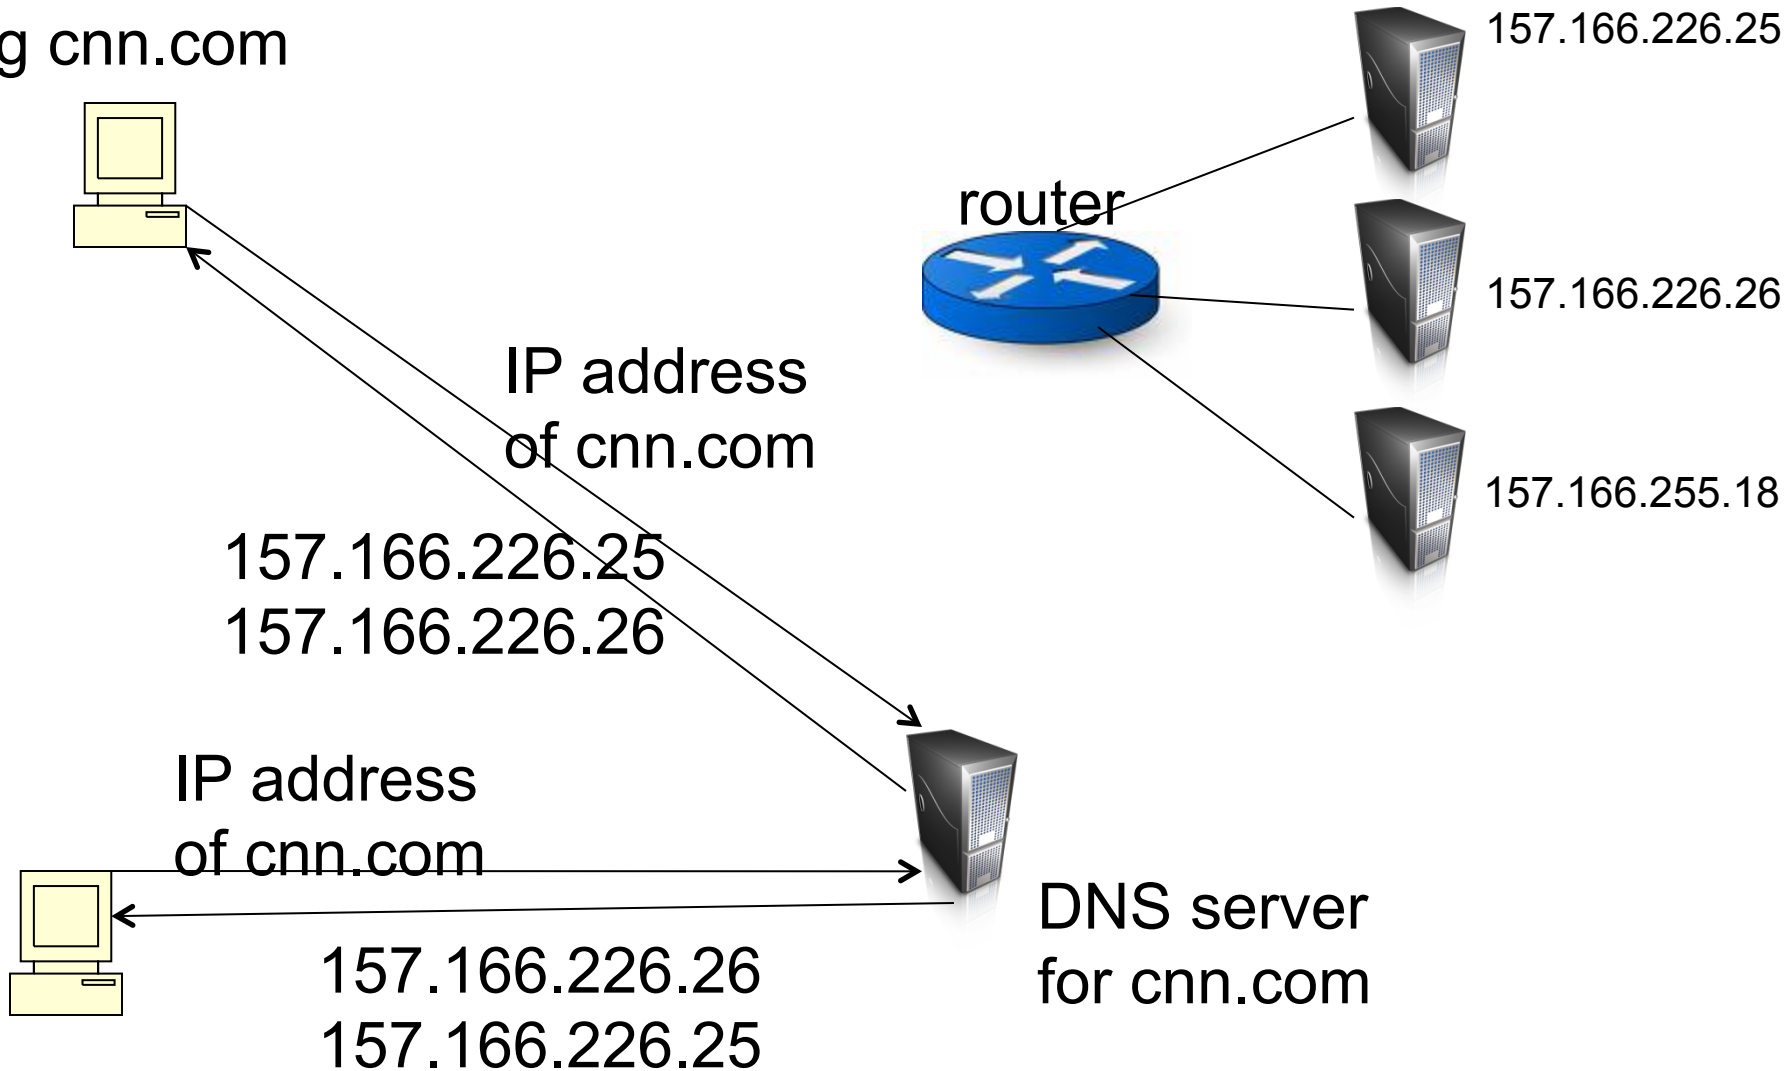
# Outline

---

- ❑ Recap
- ❑ Single network server
- ❑ Multiple network servers
  - Why multiple servers
  - Request routing mechanisms
    - Overview
    - Application-layer
    - DNS

# Basic DNS Indirection and Rotation

%dig cnn.com



# CDN Using DNS (Akamai Architecture as an Example)

- ❑ Content publisher (e.g., NYTimes)
  - provides base HTML documents
  - runs **origin** server(s)
  
- ❑ Akamai runs
  - (~200,000) **edge** servers for hosting content
    - Deployment into 110 countries and 1400 networks
  - customized **DNS redirection servers** to select edge servers based on
    - closeness to client browser
    - server load

# Linking to Akamai

- ❑ Originally, URL Akamaization of embedded content: e.g.,

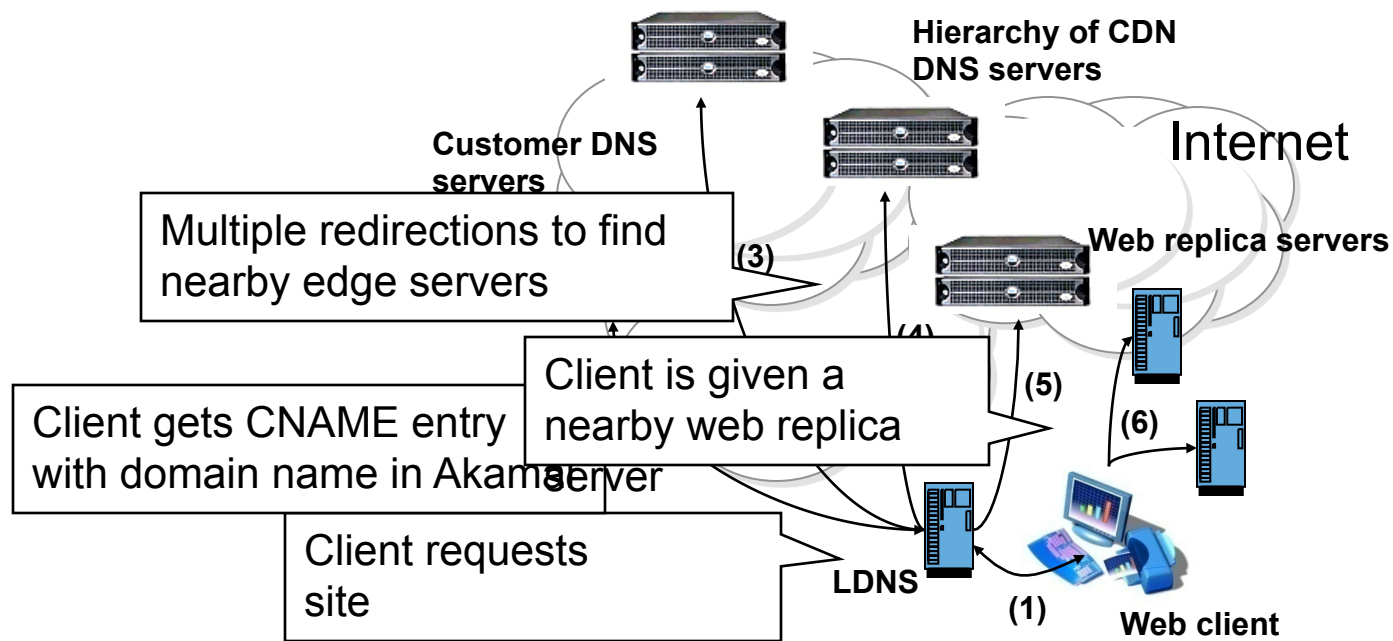
<IMG SRC= http://www.provider.com/image.gif >  
changed to

<IMG SRC = http://a661. g.akamai.net/hash/image.gif>

Note that this DNS redirection unit is per customer, not individual files.

- ❑ URL Akamaization is becoming obsolete and supported mostly for legacy reasons
  - Currently most content publishers prefer to use DNS CNAME to link to Akamai servers

# Akamai Load Direction Flow



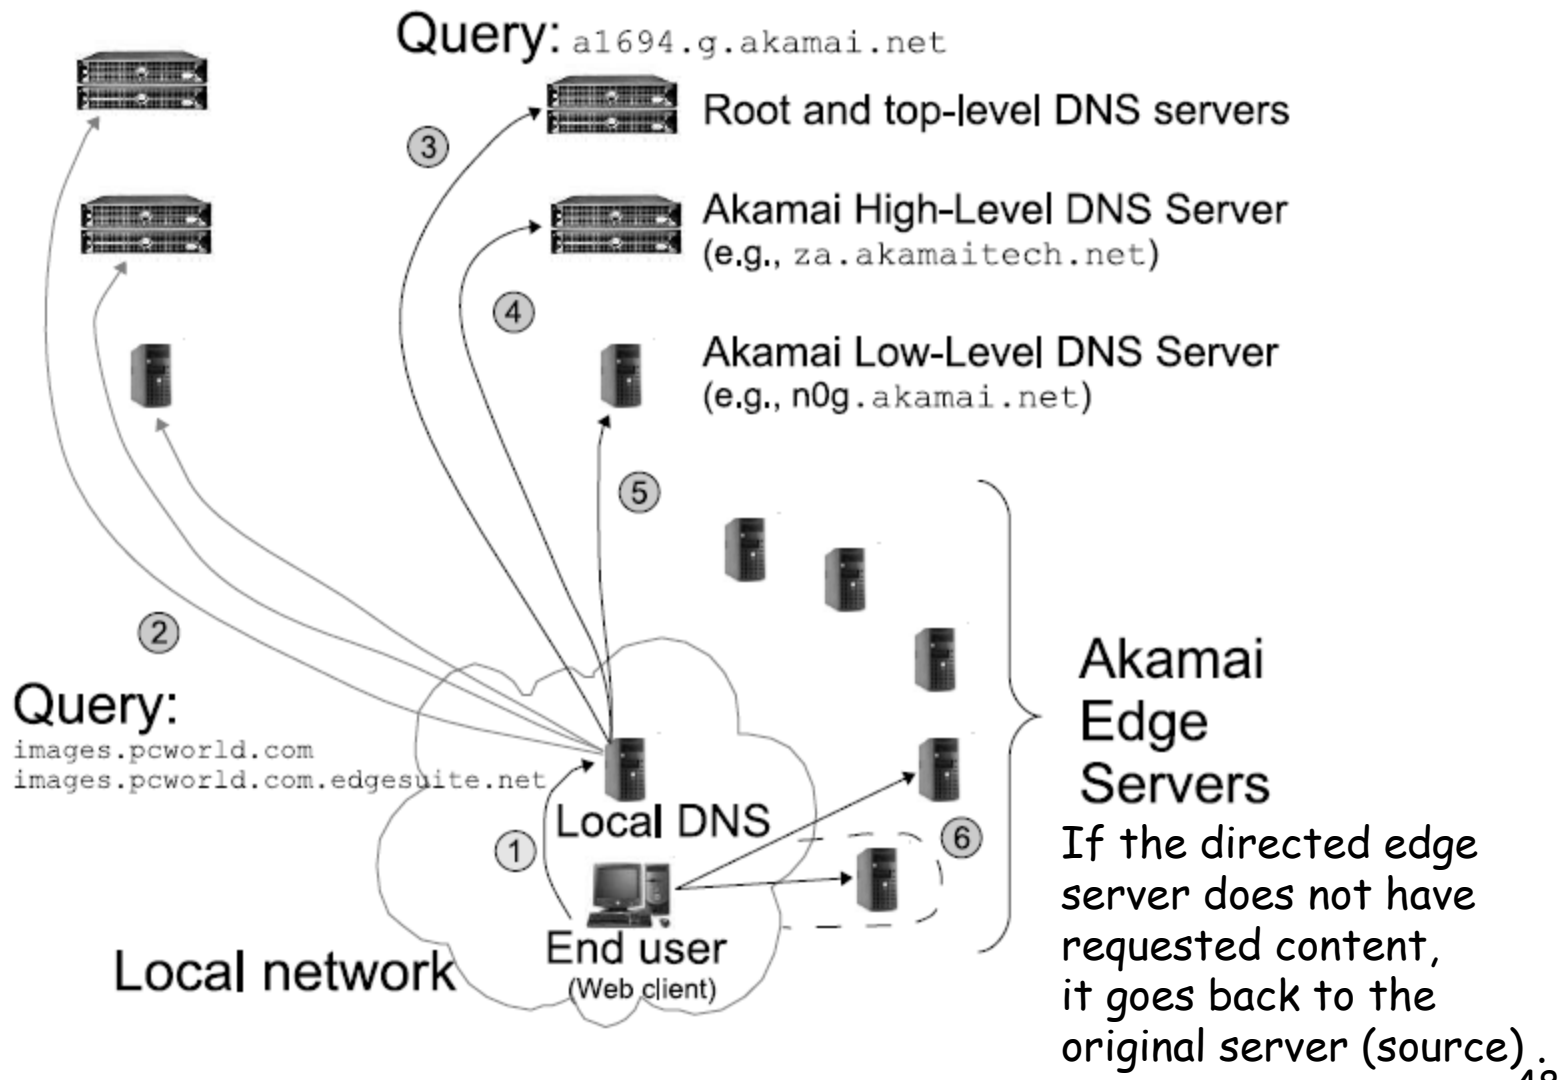
More details see “Global hosting system”: FT Leighton, DM Lewin – US Patent 6,108,703, 2000.

# Exercise: Zoo machine

---


- ❑ Check any web page of New York Times and find a page with an image
- ❑ Find the URL
- ❑ Use  
    `%dig +trace +recurse`  
    to see Akamai load direction

# Akamai Load Direction





# Two-Level Direction



proximity: high-level DNS determines  
client location; directs  
to low-level DNS, who manages  
a close-by cluster


# Local DNS Alg: Potential Input

- $p(m, e)$ : path properties (from a client site  $m$  to an edge server  $e$ )
  - Akamai might use a one-hop detour routing (see [akamai-detour.pdf](#))
- $a_m^k$ : request load from client site  $m$  to publisher  $k$
- $x_e$ : load on edge server  $e$
- caching state of a server  $e$

# Local DNS Alg

- ❑ Details of Akamai algorithms are proprietary
- ❑ A Bin-Packing algorithm (column 12 of Akamai Patent) every  $T$  second
  - Compute the load to each publisher  $k$  (called serial number)
  - Sort the publishers from increasing load
  - For each publisher, associate a list of random servers generated by a hash function
  - Assign the publisher to the first server that does not overload

# Hash Bin-Packing

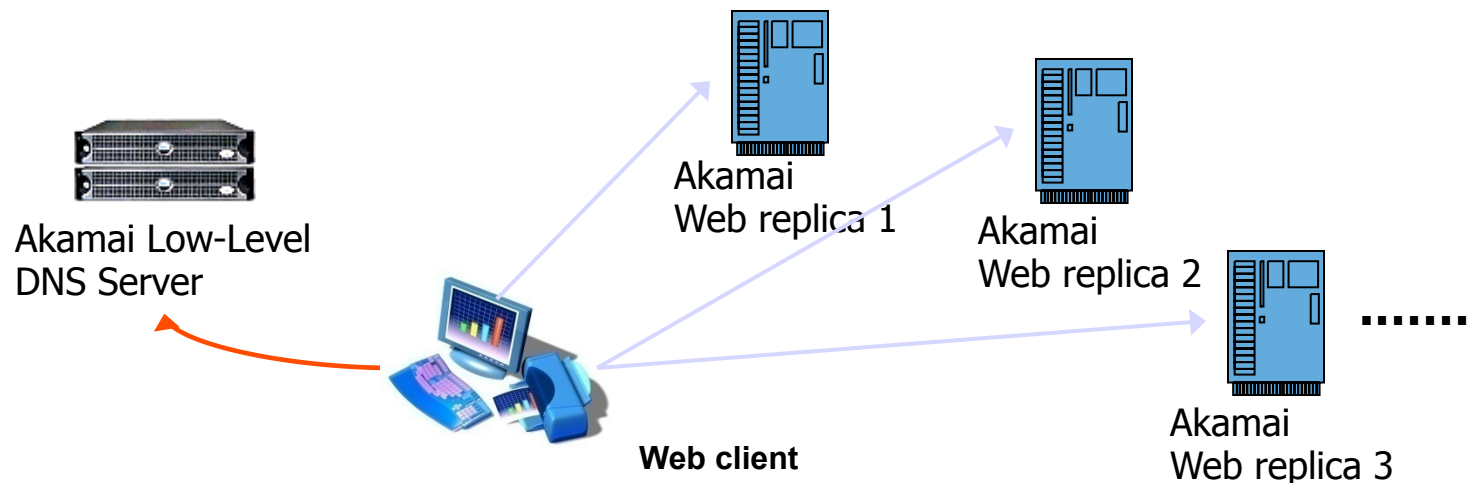


LB: maps request to individual  
machines inside cluster

# Experimental Study of Akamai Load Balancing

## ❑ Methodology


- 2-months long measurement
- 140 PlanetLab nodes (clients)
  - 50 US and Canada, 35 Europe, 18 Asia, 8 South America, the rest randomly scattered
- Every 20 sec, each client queries an appropriate CNAME for Yahoo, CNN, Fox News, NY Times, etc.




# Server Pool: to Yahoo

Target: a943.x.a.yimg.com (Yahoo)


Client 1: Berkeley




Client 2: Purdue



# Server Diversity for Yahoo




# Server Pool: Multiple Akamai Hosted Sites






# Load Balancing Dynamics




# Redirection Effectiveness: Measurement Methodology



# Do redirections reveal network conditions?

□ Rank =  $r1 + r2 - 1$

- 16 means perfect correlation



Percentage of time Akamai's selection is better or equal to rank

# (Offline Read) Facebook DNS

## Load Direction

---

- ❑ A system named Cartographer (written in Python) processes measurement data and configures the DNS maps of individual DNS servers (open source tinydns)

# Discussion

---

- Advantages and disadvantages of using DNS